

ECE411 Computer Organization and Design
Mid-Term Exam 1
10/11/2018

NetID:_____

Exam Guidelines:

1. This exam has **6 problems**. Make sure you have a complete exam before you begin **[20 pages + the cover page]**.
2. Write your name on every page in case pages become separated during grading.
3. You will have **3 hours** to complete this exam.
4. Write all of your answers on the exam itself. If you need more space to answer a given problem, continue on the back of the page, but clearly indicate that you have done so.
5. This exam is closed-book. You may use one sheet of notes. You may use a calculator.
6. **DO NOT** do anything that might be perceived as cheating. The minimum penalty will be a grade of zero.
7. Show all of your work on all problems. Correct answers that do not include work demonstrating how they were generated may not receive full credit, and answers that show no work cannot receive partial credit.
8. The exam is meant to test your understanding. Ample time has been provided. So be patient and read the questions/problems carefully before you answer.
9. **Good luck!**

Problem	Points
1. 24 pt	
2. 22 pt	
3. 30 pt	
4. 30 pt	
5. 40 pt	
6. 22 pt	
Total: 168 pt	

Question 1: Processor Pipelining

Given the following 5 stage pipeline with WB->EX & MEM->EX forwarding and register file bypass logic

Fetch stage combinational delay = 4ns

Decode stage combinational delay = 5ns

Execute stage combinational delay = 7ns

Memory stage combinational delay = 5ns

Writeback stage combinational delay = 4ns

Code:

```
ld    r1, DATA0 ; r1 <= M[DATA0]
andi  r3, r1, 0 ; r3 <= r1 & 0x0
ld    r2, DATA1 ; r2 <= M[DATA1]
add   r3, r3, r1; r3 <= r3 & r1
add   r3, r3, r2; r3 <= r3 & r2
```

```
ld    r4, DATA2
andi  r6, r4, 0
ld    r5, DATA3
add   r6, r6, r4
add   r6, r6, r5
```

Assume latch delay is 2ns. Assume the processor operates at maximum possible safe frequency. Assume 100% cache hit rates and branch prediction rates.

A. Calculate how long it will take (in ns) to run the given code. (6 points)

Name: _____

You now run the following code:

```
ld    r3, ADDR
add   r3, r1, r2
addi  r5, r3, 0x10
```

- B. You discover that this code produces the wrong value for r5 even after producing the correct result for the code at the beginning of the question. What might be a potential hardware bug in the pipeline that causes this error? (10 points)

- C. You are asked to evaluate the benefits of adding an additional WB->MEM forwarding path to the pipeline. Write an example piece of assembly code that could benefit from this new forwarding path. (8 points)

Name: _____

Question 2: Caches

Consider an architecture with a 32-bit address range and byte addressable memory. In order to exploit locality, you as the logic designer choose to implement a cache with 64-byte cache lines. 'Cache Type' refers to Direct Mapped, Set-Associative or Fully Associative in this question.

A. Assume each set has 128 bytes of data. There are 47 bits of bookkeeping (tag, valid, dirty and LRU) overhead per set. How many bits are allocated per tag, index and offset? (6 points)

B. Replacement policies are a major contributor to how well each way is utilized and the hit rate of the cache. Least Recently Used is a straightforward replacement policy with two common flavors; Pseudo LRU and True LRU.

How many bits are needed per set for a tree-based pseudo LRU scheme for a 16-way set associative cache? What is the bit storage overhead of using a queue to track True LRU for n-ways where $n > 2$? (6 points)

Name: _____

C. Prefetching is a scheme devised to predict what data will be useful in the near future and bring it into the caches before the CPU needs to access it. This allows potentially useful memory accesses to be overlapped with execution of non-memory operations.

Assume that we can build a highly accurate prefetcher that will only issue prefetch accesses to data that will be used at some point during program execution. Provide a scenario where performance degrades. (4 points)

D. The program below shows nested for loops loading and operating on some data from the `structure[]` array. A stride is used to determine the address difference between each subsequent load. Given the following program, answer the listed questions.

```
int structure[...]; // a very large array
uint32_t stride = 128;
for (int j = 0; j < 32; j++) {
    stride = stride * 2;

    for (uint32_t i = 0; i < 16; i++) {
        // Load an element from some structure
        int cur_element = structure[i*stride];

        // Operate on the cur_element
        ...
    }
}
```

Name:_____

You observe that the memory access times are increasing in the inner loop as the stride increases, which causes a decrease in performance. The memory access time saturates at very large strides. This behavior is caused by a near-zero hit rate in the data cache. What type of cache might be in the processor running the code above? Why? (6 points)

Name:_____

Question 3: Caches and Virtual Memory

Congratulations! you have accepted an offer from a new startup called "NoHope" as a HW architect in their processor design group. They are currently working on their cache design. The team leader explained their current preference - using virtually-indexed, virtually-tagged (VIVT) L1 instruction and data caches. After looking at their cache designs, you realized that they didn't handle any of the known issues associated with VIVT caches.

- A. After hearing you explain the issues, your team leader got convinced and asked you to propose 2 software solutions each for a minimum of two VIVT issues. What would be your solutions? (8 points)
- B. Unfortunately, the team leader worries about performance degradation and complexity typical of software solutions and asks you to suggest two alternatives (hardware only or hardware supported by software solution). What are you going to propose (caches still need to be VIVT)? (8 points)

Name: _____

- C. You also found out that they are using an 8 KB virtually indexed, physically tagged (VIPT) cache as the L2 cache. However, they would now like to double the size of the L2 cache (since new benchmarks have emerged) and decide on the associativity. With the information given in the table below what is the associativity that you recommend and why? (6 points)

Virtual address size	Page size	Size of each page table entry
32 bits	4 KB	8 bytes

- D. Frustrated by the poor design choices made by NoHope, you decide to interview with other processor design companies. During an interview at ALittleHope, you have been given the following information:

- A processor with a 32-bit virtual address space.
- A 64KB of physical memory space.
- The size of a page is 1KB.
- TLB is a direct-mapped cache with 4 entries.

You are then asked:

- a. How many bits are stored for each TLB tag entry? (4 points)

- b. How many bits are stored for each TLB data entry? (4 points)

Name: _____

Question 4: ISA

Consider the following ISAs and instruction latencies for some simple multicycle implementation (like your MP1 design) of each ISA (do not worry about cache misses and hits, assume AMAT is baked into the latency of all instructions that have memory accesses):

RegReg: a register-register ISA (assume 16 general purpose registers named r0-r15) (140MHz)

Assembly Instruction formats	Instruction latency (# of cycles)
load DR, label	8
store SR, label	8
add DR, SR1, SR2	6
sub DR, SR1, SR2	6
mul DR, SR1, SR2	10

Stack: a stack ISA (100MHz)

Assembly Instruction formats	Instruction latency (# of cycles)
push label	8
pop label	8
add	6
sub	6
mul	10

Name: _____

MemMem: a memory-memory ISA (140MHz) (label1 is always the destination)

Assembly Instruction formats	Instruction latency (# of cycles)
add label1, label2, label3	12
sub label1, label2, label3	12
mul label1, label2, label3	16
mov label1, label2	9

A) Write assembly code to execute the following C function in each ISA with no optimization: (12 pts)

```
/* global variables, no need to produce code for this part, just  
assume the labels are available to use */
```

```
int A;  
int B;  
int C;  
int tmp;  
int result;
```

```
void foo() {  
    result = result + (B*(C-A));  
}
```

As always, write any assumptions you make.

RegReg:

Stack:

Name: _____

MemMem:

B) Imagine you want to add an atomic integer addition instruction to one of these ISAs. An atomic operation should write its result to memory before any other thread has a chance to overwrite the memory operands of the instruction. Assume this is done by simply asserting a lock signal on the bus interface across multiple read/write operations (like the mem_read signal is asserted to do a read). This amendment would be easiest to implement with which ISA and why (6 pts)?

C) What do you think would be the latency of this new instruction in number of cycles and why (6 pts)?

D) Which ISA uses the instruction cache most efficiently, considering that labels take many bits to encode (much more than register operands) and that compilers must emit more instructions for some ISAs than others? Why? (6 pts)

Question 5: MPs

This question contains multiple parts and aims to assess your ability to understand a specification and implement desired functionality using good design practices.

Programmers that optimize their code for high performance often desire information about the utilization of the underlying hardware during the execution of their programs. In order to provide this information, modern processors include blocks which monitor the state of the machine and instructions which allow access to these state registers.

- a. You are tasked with writing a module that tracks the number of instructions executed by your processor as well as the number of cycles spent waiting on memory accesses. You should write this code in the provided skeleton below. You are told that the *enable* signal will be high for exactly one cycle during the execution of each unique instruction. The memory signals, *mem_read*, *mem_write*, and *mem_resp* behave identically to the class MPs. The *rst* signal is **active low** and should set your counters to zero. (4 pts)

```

module cpu_state_reg (
    input logic clk, rst, enable, mem_read, mem_write, mem_resp,
    output logic[31:0] instrs_exec, mem_cycles
);
// Declare internal signals here

// Combinational logic
always_comb begin

end
// Sequential logic
always_ff @(posedge clk) begin

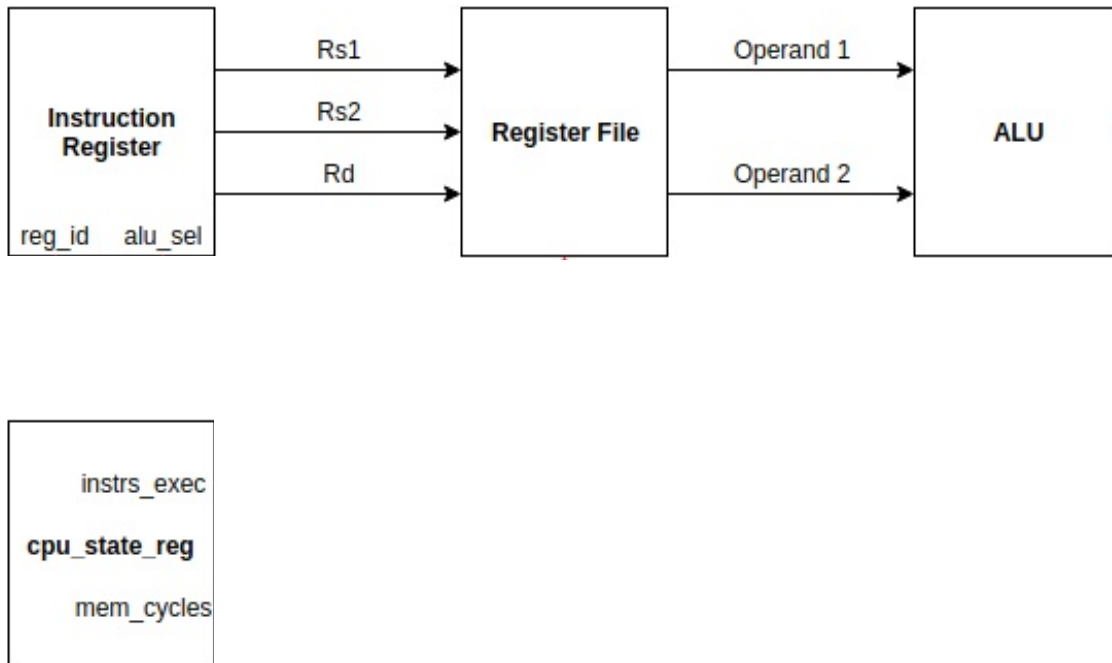
end
endmodule : cpu_state_reg

```

Name: _____

Now that your design can keep track of its state, you must add support to access these registers and allow a user to see the statistics of their program. The new instruction, *rfsr*, stands for “read from state register” and has the following format: *rfsr Rd, <reg_id>* . This instruction reads from the module you created above and writes the counter value into the destination register, *Rd*. For this problem, *<reg_id>* is a 1-bit signal and can be thought of as the “address” of the counter which should be read. Let *reg_id = 0* correspond to the instruction count and *reg_id = 1* correspond to the memory cycle count. In real processors, there are many different registers to choose from.


- b. Modify the partial datapath below by filling in the missing logic required to support the new *rfsr* instruction. Include labels where needed. Write a sentence or two describing each of the changes you’ve made to the datapath. If you utilize a block other than logic gates or muxes, you must explain the function of the block. *Alu_sel* is a 1-bit signal which is high when an ALU operation is being performed and low otherwise. (5 pts)



Name: _____

Finally, we must consider the impact of this instruction on the CPU controller. For this machine, there are three fetch states and a single decode state. These states operate identically to those implemented in MP0 and MP1. Following the decode state, you decide to add new state(s) to handle the *rfsr* instruction.

- c. Draw the new state(s) needed to support the the *rfsr* instruction. The decode state has been provided for you. For each state you draw, you should write a sentence or two to describe what is being accomplished. Feel free to use signals to describe the functionality of the state(s), just be sure the signal names match your datapath drawn above. (4 pts)



Name: _____

Question 6: Potpourri:

- Consider an ISA with three instructions (Inst1, Inst2, Inst3). Write an instruction sequence with 3 instructions that will have worse performance on a single cycle design compared to multi-cycle design? Why? Assume that inst1, inst2, and inst3 take 1, 2, and 3 cycle respectively on the multi-cycle design. Will this instruction sequence have higher performance on a 3-stage pipeline (vs single-cycle design)? Why? (8 points)
- If D is dynamic power of a CMOS processor design, S is static power of the design, draw the D/S curve showing how it has changed over last 20 years. Make sure to mark 1 on y-axis. Make sure to mark 2018 on x-axis. (6 points)
- Here's a comparison between STT-RAM, a memory technology based on magnetic spins and SRAM:

	Area (mm ²)	Read Energy (nJ)	Write Energy (nJ)	Leakage Power at (mW)	Read Latency (ns)	Write latency (ns)	Read @ 2 GHz (cycles)	Write @2 GHz (cycles)
1 MB SRAM	2.61	0.578	0.578	4542	1.012	1.012	2	2
4MB STT- RAM	3.00	1.035		2524	0.998		2	

Where would you put STT-RAM in the memory hierarchy? Why? (6 points)

Name:_____

- d. If McDonalds is “I am loving it”, Skittles is “Taste the Rainbow”, Google is “Don’t be Evil”, ECE411 is

Appendix A:

	15	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0
ADD ⁺	0001				DR			SR1			0	00		SR2		
ADD ⁺	0001				DR			SR1			1	imm5				
AND ⁺	0101				DR			SR1			0	00		SR2		
AND ⁺	0101				DR			SR1			1	imm5				
BR	0000				n	z	p	PCOffset9								
JMP	1100				000			BaseR			000000					
JSR	0100				1	PCOffset11										
JSRR	0100				0	00		BaseR			000000					
LDB ⁺	0010				DR			BaseR			offset6					
LDI ⁺	1010				DR			BaseR			offset6					
LDR ⁺	0110				DR			BaseR			offset6					
LEA ⁺	1110				DR			PCOffset9								
NOT ⁺	1001				DR			SR			111111					
RET	1100				000			111			000000					
RTI	1000				000000000000											
SHF ⁺	1101				DR			SR			A	D	imm4			
STB	0011				SR			BaseR			offset6					
STI	1011				SR			BaseR			offset6					
STR	0111				SR			BaseR			offset6					
TRAP	1111				0000			trapvect8								

Figure 1.2: LC-3b Instruction Formats. NOTE: + indicates instructions that modify condition codes.

Appendix B:

31	27	26	25	24	20	19	15	14	12	11	7	6	0	
funct7				rs2		rs1		funct3		rd		opcode		R-type
imm[11:0]						rs1		funct3		rd		opcode		I-type
imm[11:5]				rs2		rs1		funct3		imm[4:0]		opcode		S-type
imm[12 10:5]				rs2		rs1		funct3		imm[4:1 11]		opcode		B-type
imm[31:12]										rd		opcode		U-type
imm[20 10:1 11 19:12]										rd		opcode		J-type

RV32I Base Instruction Set

imm[31:12]										rd		0110111		LUI
imm[31:12]										rd		0010111		AUIPC
imm[20:10:1 11 19:12]										rd		1101111		JAL
imm[11:0]				rs1		000		rd		1100111		JALR		
imm[12:10:5]				rs2		rs1		000		imm[4:1 11]		1100011		BEQ
imm[12:10:5]				rs2		rs1		001		imm[4:1 11]		1100011		BNE
imm[12:10:5]				rs2		rs1		100		imm[4:1 11]		1100011		BLT
imm[12:10:5]				rs2		rs1		101		imm[4:1 11]		1100011		BGE
imm[12:10:5]				rs2		rs1		110		imm[4:1 11]		1100011		BLTU
imm[12:10:5]				rs2		rs1		111		imm[4:1 11]		1100011		BGEU
imm[11:0]				rs1		000		rd		0000011		LB		
imm[11:0]				rs1		001		rd		0000011		LH		
imm[11:0]				rs1		010		rd		0000011		LW		
imm[11:0]				rs1		100		rd		0000011		LBU		
imm[11:0]				rs1		101		rd		0000011		LHU		
imm[11:5]				rs2		rs1		000		imm[4:0]		0100011		SB
imm[11:5]				rs2		rs1		001		imm[4:0]		0100011		SH
imm[11:5]				rs2		rs1		010		imm[4:0]		0100011		SW
imm[11:0]				rs1		000		rd		0010011		ADDI		
imm[11:0]				rs1		010		rd		0010011		SLTI		
imm[11:0]				rs1		011		rd		0010011		SLTIU		
imm[11:0]				rs1		100		rd		0010011		XORI		
imm[11:0]				rs1		110		rd		0010011		ORI		
imm[11:0]				rs1		111		rd		0010011		ANDI		
0000000				shamt		rs1		001		rd		0010011		SLLI
0000000				shamt		rs1		101		rd		0010011		SRLI
0100000				shamt		rs1		101		rd		0010011		SRAI
0000000				rs2		rs1		000		rd		0110011		ADD
0100000				rs2		rs1		000		rd		0110011		SUB
0000000				rs2		rs1		001		rd		0110011		SLL
0000000				rs2		rs1		010		rd		0110011		SLT
0000000				rs2		rs1		011		rd		0110011		SLTU
0000000				rs2		rs1		100		rd		0110011		XOR
0000000				rs2		rs1		101		rd		0110011		SRL
0100000				rs2		rs1		101		rd		0110011		SRA
0000000				rs2		rs1		110		rd		0110011		OR
0000000				rs2		rs1		111		rd		0110011		AND