

# ECE411: Computer Organization and Design

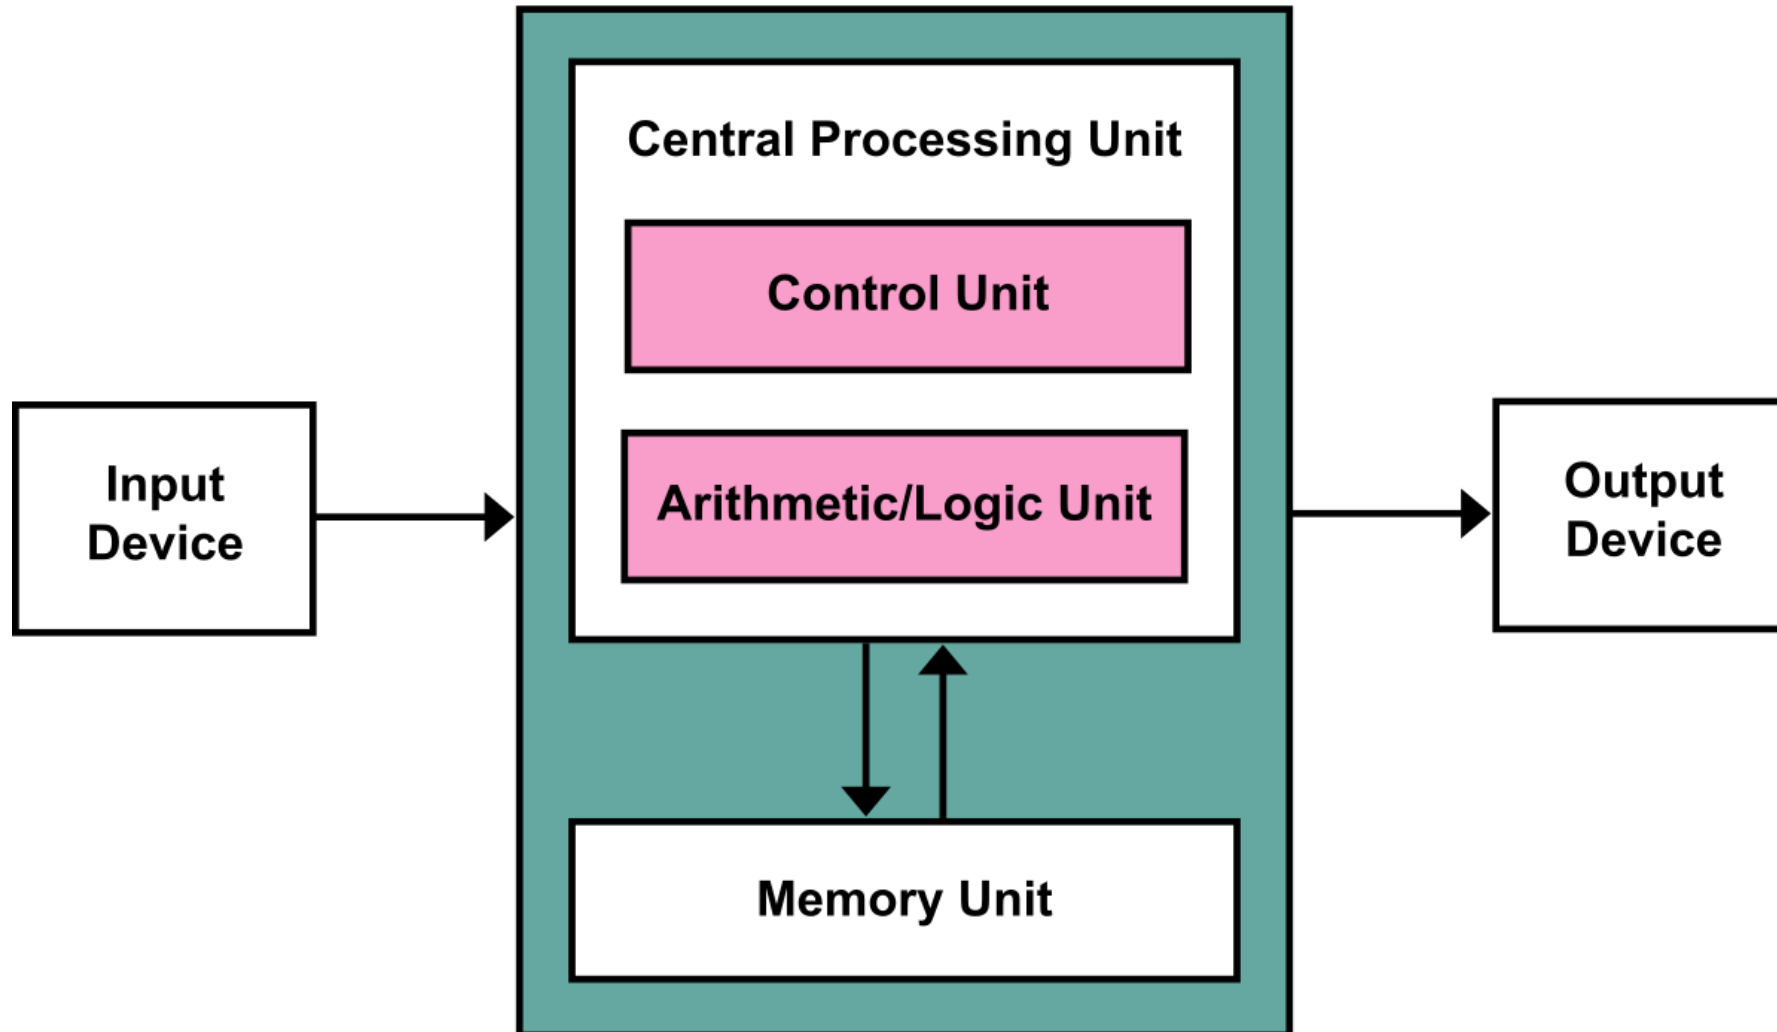
## Lecture 2: Instruction Set Architecture

Rakesh Kumar

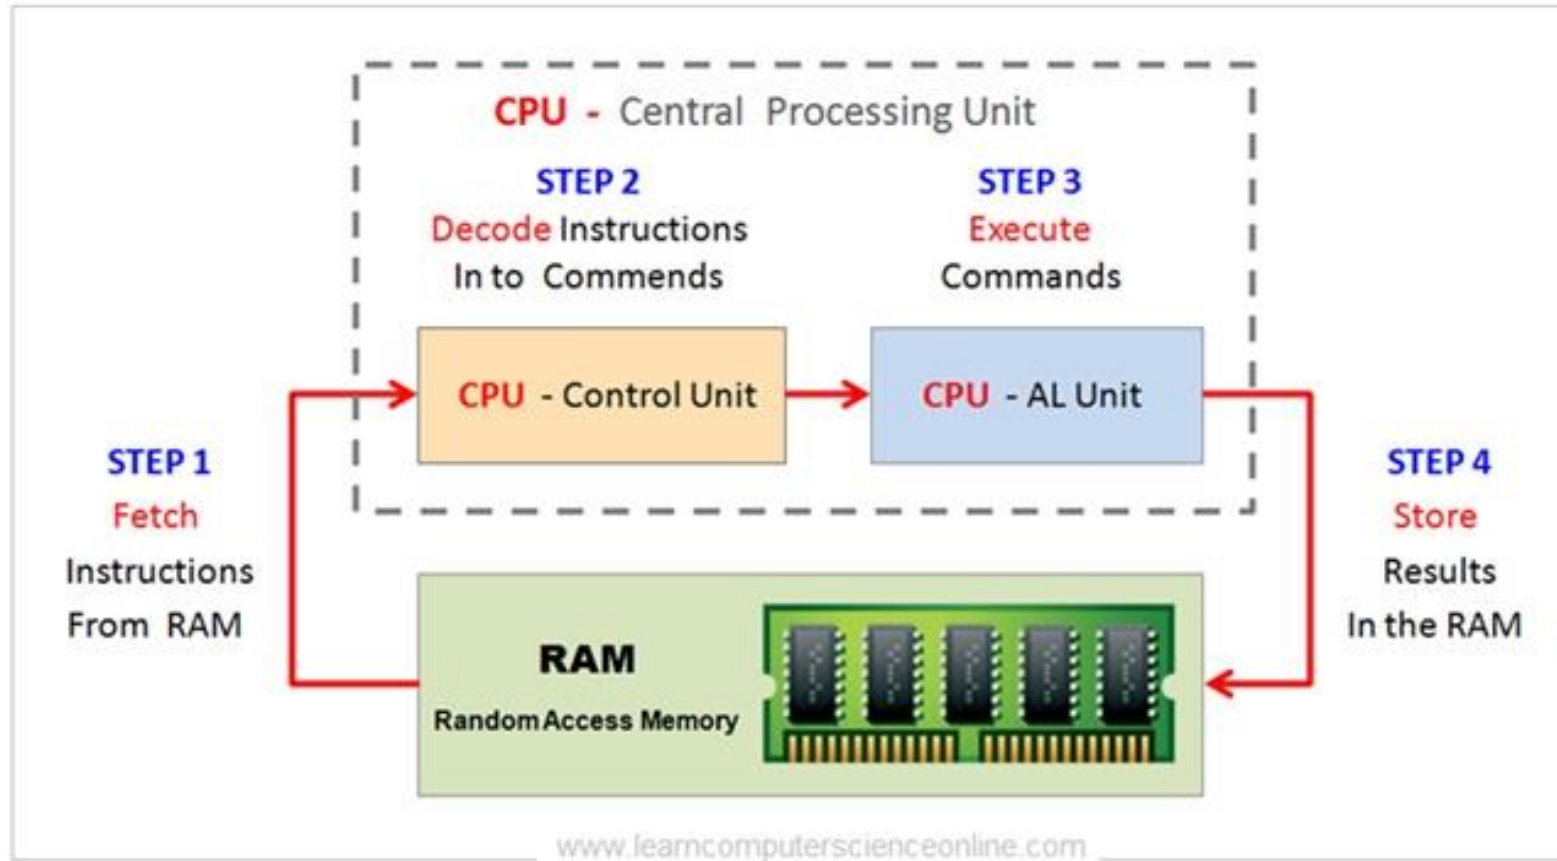


# **Von-Neumann Computer Architecture**

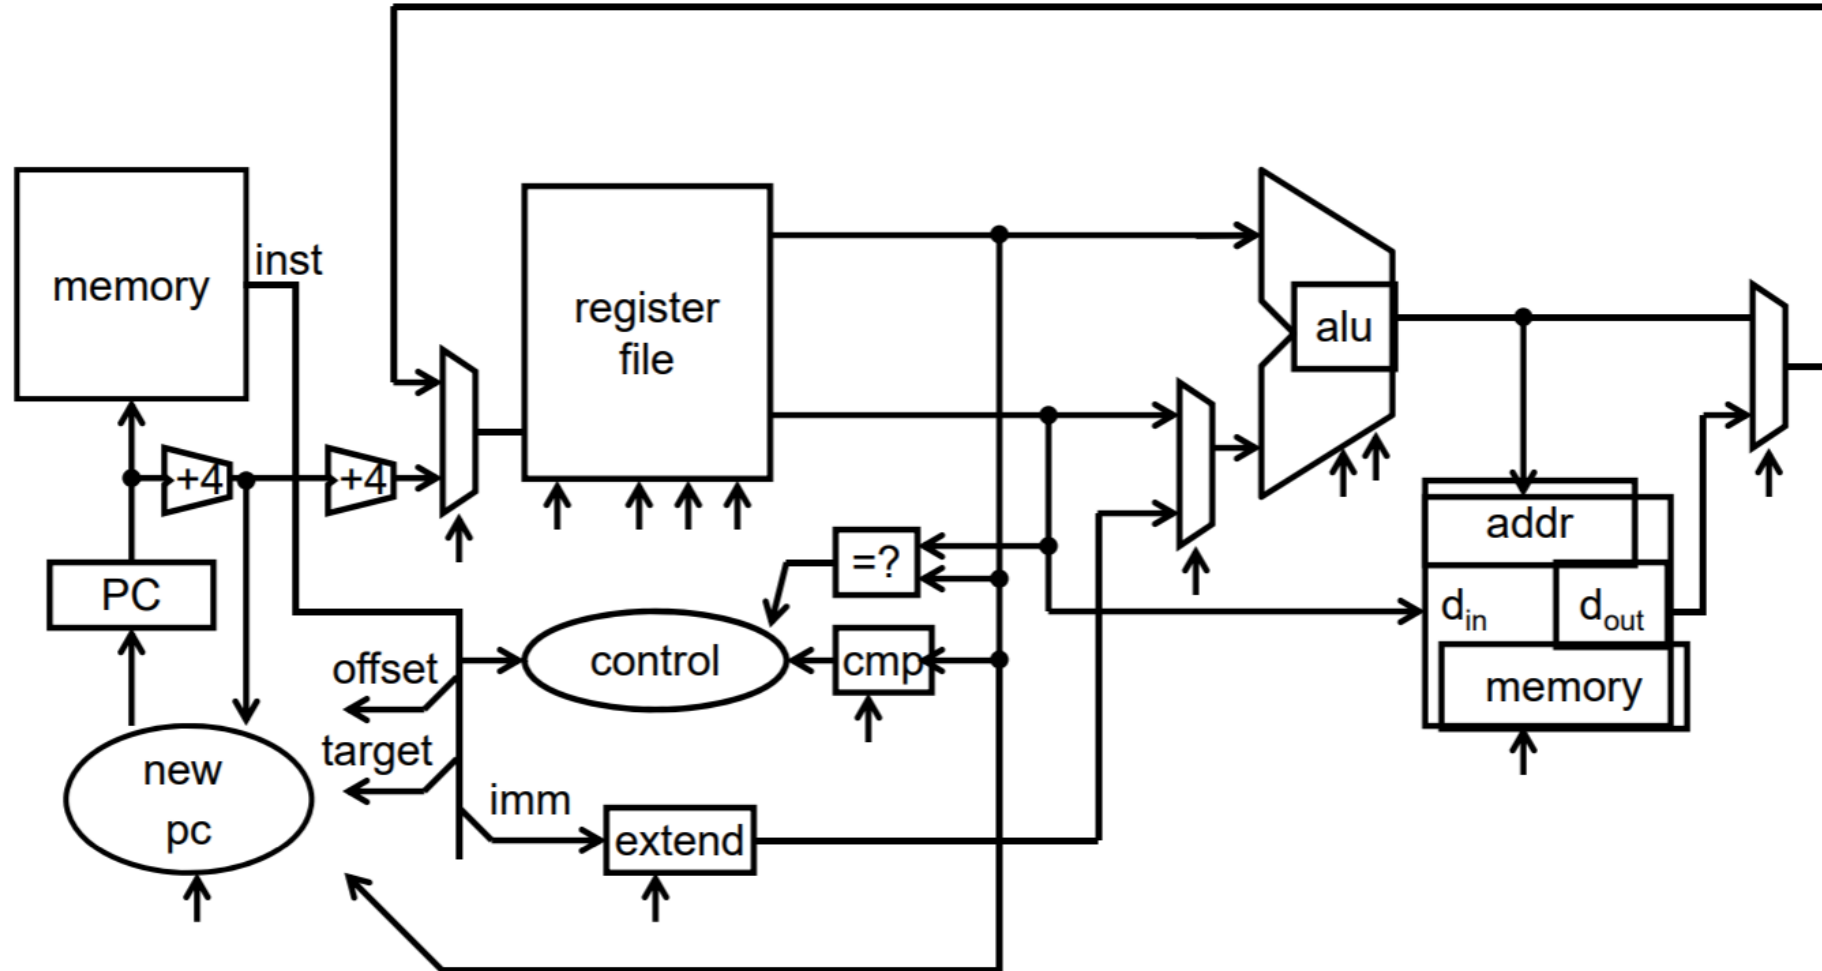
# Von-Neumann Computer Organization



# Instruction Execution Cycles



# A Simple Processor



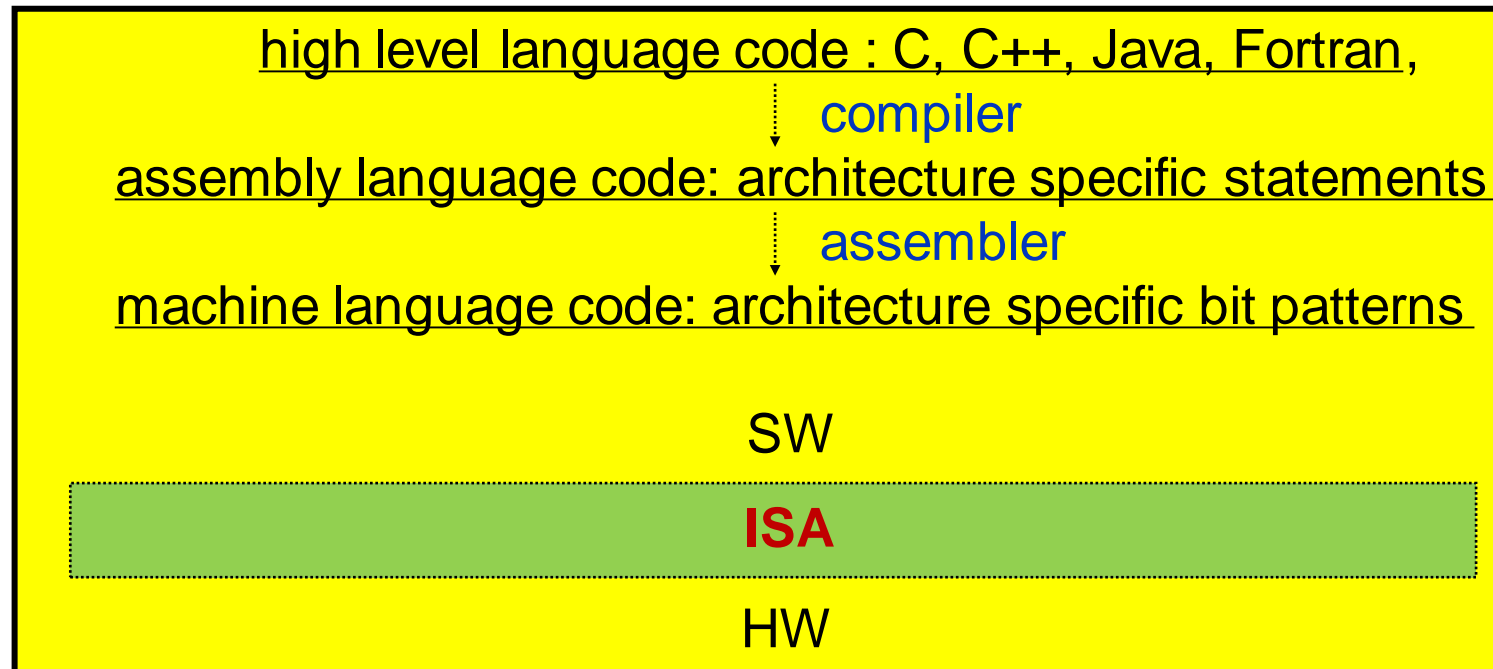
# A Simple Processor (Detailed Explanation)

- executes instructions, which are represented as bit patterns with specific fields interpreted differently
- contains a register file, which holds data that will be referenced by instructions
- has a program counter that holds the address of the instruction being executed
- can access memory
  - ⦿ use address to select the location we want to access
  - ⦿ random-access: time to access a piece of data is independent of which address the data is stored in
- I/O
  - ⦿ keyboard, mouse, display, network
  - ⦿ long-term data storage: hard disk, CD-ROM, SSD, etc.

# Instruction Set Architecture (ISA)

Q: What is the function of ISA?


- serves as an **interface** between software and hardware
- provides a mechanism by which the software tells the hardware what should be done



# Interface Design

Q: What does interface mean in computer systems/architecture design?

- a good interface:
  - ⊙ lasts through many implementations (portability, compatibility)
  - ⊙ is used in different ways (generality)
  - ⊙ provides convenient functionality to higher levels
  - ⊙ permits an efficient implementation at lower levels







# Today's ISAs



MCU Unit Shipments (2013), by Architecture





# Case Study: ARM in server space

- 1. There's no desktop development environment (where apps can be developed/debugged before deploying on the cloud).
- 2. Very few players, so no reason to develop a software ecosystem.
- 3. Very public failures, discourages new efforts.
- 4. Poor performance of existing ARM servers
- 5. Geopolitics may discourage ARM further
- 6. Resurgence of AMD means that need for ARM is weaker.

# Case Study: RISC-V


ISA Comparison				
	Custom ISA	Licensed ISA	Licensed ISA with architecture license	Open-Source ISA
<b>Design Flexibility</b>	High	Low	Moderate to high	High
<b>License Fees</b>	None	\$0 to millions	\$0 to millions	None
<b>Royalty Fees</b>	None	0 to a few %	0 to a few %	0 to a few %
<b>Available Software</b>	None to little	Moderate to extensive	Moderate to extensive	Little to moderate
<b>Processor Cores</b>	Custom	Standard	Standard to custom	Custom or standard
<b>Hardware Engineering Costs</b>	High	Low	Low to moderate	High
<b>Software Engineering Costs</b>	High	Low to moderate	Low to moderate	Moderate to high
<b>Development Time</b>	Long	Short	Short to moderate	Long
<b>Manufacturability Risk</b>	High	Low	Low to moderate	Moderate to high
<b>Time-to-Market</b>	Long	Short	Short to moderate	Moderate to high

<https://www.forbes.com/sites/tiriasresearch/2018/04/05/what-you-need-to-know-about-processor-architectures/#72fec5344f57>

# Interface Design

Q: What does interface mean in computer systems/architecture design?

- a good interface:
  - ⊙ lasts through many implementations (portability, compatibility)
  - ⊙ is used in different ways (generality)
  - ⊙ provides convenient functionality to higher levels
  - ⊙ permits an efficient implementation at lower levels




# Instruction Set Design Issues

Q: As a hardware architect, what you should consider in ISA design?

- instruction set design issues include:
  - ⊙ where are operands stored?
    - registers, memory, stack, accumulator
  - ⊙ how many explicit operands are there?
    - 0, 1, 2, or 3
  - ⊙ how is the operand location specified?
    - register, immediate, indirect, . . .
  - ⊙ what type & size of operands are supported?
    - byte, int, float, double, string, vector. . .
  - ⊙ what operations are supported?
    - add, sub, mul, move, compare . . .


# Instruction Length

variable:



**x86 – Instructions vary from 1 to 17 Bytes long**

fixed:



**MIPS, PowerPC, and most other RISC's:  
all instruction are 4 Bytes long**



# Instruction Length

- variable-length instructions (x86):
  - ⊙ require multi-step, complex fetch and decode (-)
  - ⊙ allow smaller binary programs that require less disk storage, less DRAM at runtime, less memory, bandwidth and better cache efficiency (+)
- fixed-length instructions (RISC's)
  - ⊙ allow easy fetch and decode (+)
  - ⊙ simplify pipelining and parallelism (+)
  - ⊙ result in larger binary programs that require more disk storage, more DRAM at runtime, more memory bandwidth and lower cache efficiency (-)

# Case Study: ARM

How can we overcome the limitation of fixed-length instructions?

- ARM (Advanced RISC Machine)
  - ⊙ started with fixed, 32-bit instruction length
  - ⊙ **added thumb instructions**
    - a subset of the 32-bit instructions
    - all encoded in 16 bits
    - all translated into equivalent 32-bit instructions within the processor pipeline at runtime
    - can access only 8 general purpose registers
  - ⊙ motivated by many resource constrained embedded applications that require less disk storage, less DRAM at runtime, less memory bandwidth and better cache efficiency

Q: How about RISC-V?

# How many registers?

- most computers have a small set of registers
  - ⊙ on-chip memory to hold values that will be used soon
  - ⊙ a typical instruction use 2 or 3 register values
- advantages of a small number of registers:
  - ⊙ it requires fewer instruction bits to specify which one
  - ⊙ less hardware
  - ⊙ faster access (shorter wires, fewer gates)
  - ⊙ faster context switch (when all registers need saving)
- advantages of a larger number:
  - ⊙ fewer loads and stores needed
  - ⊙ easier to express several operations in parallel

# Where do operands reside?

When the ALU needs them?

- stack machine:
  - ⊙ **push** loads memory into 1st register (“top of stack”), moves other regs down
  - ⊙ **pop** does the reverse
  - ⊙ **add** combines contents of first two registers, moves rest up
- accumulator machine:
  - ⊙ only 1 register (called the “accumulator”)
  - ⊙ instruction includes **store** and  $ACC \leftarrow ACC + MEM$
- register-memory machine :
  - ⊙ arithmetic instructions can use data in registers and/or memory
- load-store machine (aka register-register machine):
  - ⊙ arithmetic instructions can only use data in registers

# Classifying ISAs

## accumulator (before 1960, e.g., 68HC11):

1-address	add A	$\text{acc} \leftarrow \text{acc} + \text{mem}[A]$
-----------	-------	--

## stack (1960s to 1970s):

0-address	add	$\text{tos} \leftarrow \text{tos} + \text{next}$
-----------	-----	--

## memory-memory (1970s to 1980s):

2-address	add A, B	$\text{mem}[A] \leftarrow \text{mem}[A] + \text{mem}[B]$
-----------	----------	--

3-address	add A, B, C	$\text{mem}[A] \leftarrow \text{mem}[B] + \text{mem}[C]$
-----------	-------------	--

## register-memory (1970s to present, e.g., 80x86):

2-address	add R1, A	$R1 \leftarrow R1 + \text{mem}[A]$
-----------	-----------	------------------------------------

	load R1, A	$R1 \leftarrow \text{mem}[A]$
--	------------	-------------------------------

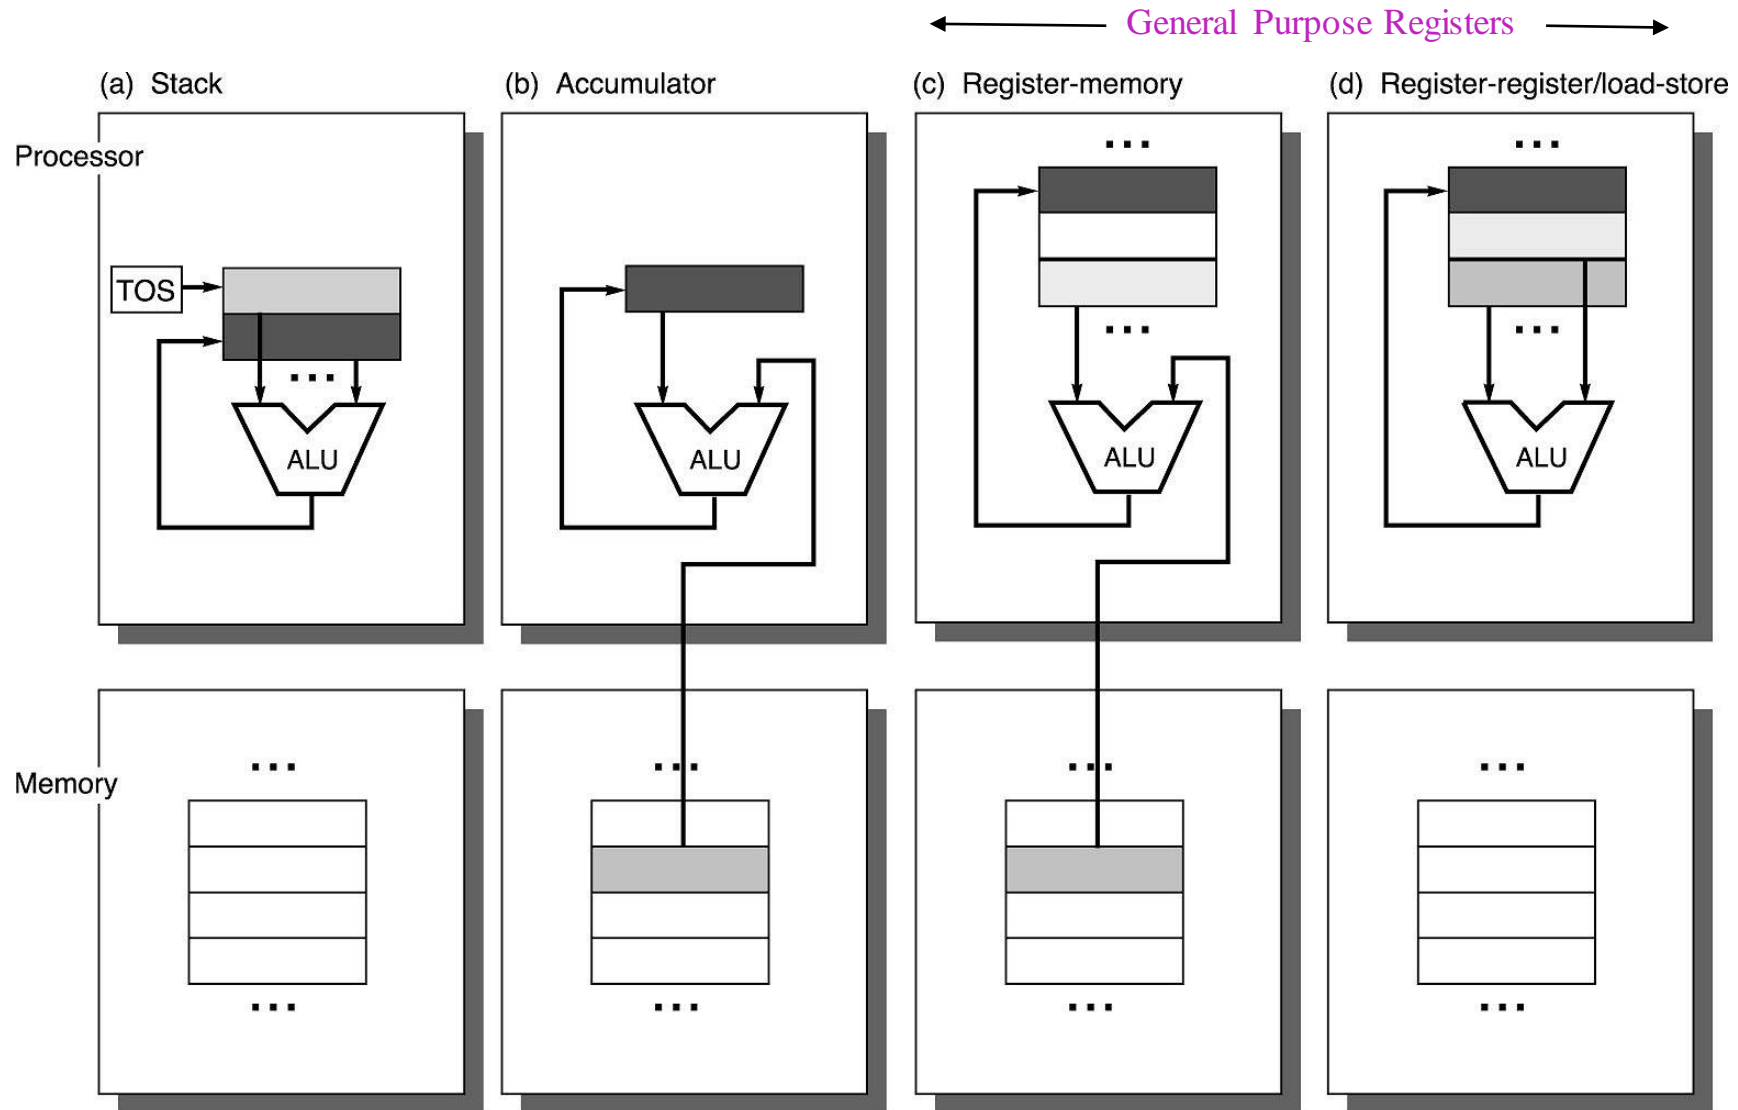
## register-register (load/store) (1960s to present, e.g., MIPS):

3-address	add R1, R2, R3	$R1 \leftarrow R2 + R3$
-----------	----------------	-------------------------

	load R1, R2	$R1 \leftarrow \text{mem}[R2]$
--	-------------	--------------------------------

	store R1, R2	$\text{mem}[R1] \leftarrow R2$
--	--------------	--------------------------------


# Operand Locations in Four ISA Classes



# Four Instruction Sets


- code sequence  $C = A + B$

stack	accumulator	Register (register-memory)	register (load-store)
push A push B add pop C	load A add B store C	load R1, A add R1, B store C, R1	load R1, A load R2, B add R3, R1, R2 store C, R3



# More About General Purpose Registers

- why do almost all new architectures use GPRs?
  - ⊙ registers are much faster than memory (even cache)
    - register values are available immediately
    - when memory isn't ready, processor must wait ("stall")
  - ⊙ registers are convenient for variable storage
    - compiler assigns some variables just to registers
    - more compact code since small fields specify registers (compared to memory addresses)






# Latency Numbers Every Programmer Should Know

- Registers 0.5 – 2 ns
- L1 cache reference 0.5 ns
- Branch mispredict 5 ns
- L2 cache reference 7 ns 14x L1 cache
- Mutex lock/unlock 25 ns
- Main memory reference 100 ns 20x L2 cache, 200x L1 cache
- Compress 1K bytes with Zippy 3,000 ns 3 us
- Send 1K bytes over 1 Gbps network 10,000 ns 10 us
- Read 4K randomly from SSD\* 150,000 ns 150 us ~1GB/sec SSD
- Read 1 MB sequentially from memory 250,000 ns 250 us
- Round trip within same datacenter 500,000 ns 500 us
- Read 1 MB sequentially from SSD\* 1,000,000 ns 1,000 us 1 ms ~1GB/sec SSD, 4X memory
- Disk seek 10,000,000 ns 10,000 us 10 ms 20x datacenter roundtrip
- Read 1 MB sequentially from disk 20,000,000 ns 20,000 us 20 ms 80x memory, 20X SSD
- Send packet CA->Netherlands->CA 150,000,000 ns 150,000 us 150 ms

*Credits: measured by Google (numbers vary on different machines)*


# Stack Architecture

- instruction set:
  - ⊙ add, sub, mult, div, . . .
  - ⊙ push A, pop A



- example:  $A*B - (A+C*B)$

- ⊙ push A
- ⊙ push B
- ⊙ mul
- ⊙ push A
- ⊙ push C
- ⊙ push B
- ⊙ mul
- ⊙ add
- ⊙ sub




# Stacks: Pros and Cons

- pros
  - ⊙ good code density (implicit top of stack)
  - ⊙ low hardware requirements
  - ⊙ easy to write a simpler compiler for stack architectures
- cons
  - ⊙ stack becomes the bottleneck
  - ⊙ little ability for parallelism or pipelining
  - ⊙ data is not always at the top of stack when need, so additional instructions like **swap** are needed
  - ⊙ difficult to write an optimizing compiler for stack architectures

# Accumulator Architectures

- instruction set:


- ⊙ add A, sub A, mult A, div A, . . .
- ⊙ load A, store A



$acc = acc +,-,*,/ mem[A]$

- example:  $A*B - (A+C*B)$

- ⊙ load B
- ⊙ mul C
- ⊙ add A
- ⊙ store D
- ⊙ load A
- ⊙ mul B
- ⊙ sub D



# Accumulators: Pros and Cons

- pros
  - ⊙ very low hardware requirements
  - ⊙ easy to design and understand
- cons
  - ⊙ accumulator becomes the bottleneck
  - ⊙ little ability for parallelism or pipelining
  - ⊙ high memory traffic

# Memory-Memory Architecture

- instruction set:

(3 operands)	add A, B, C	sub A, B, C	mul A, B, C
(2 operands)	add A, B	sub A, B	mul A, B

- example:  $A*B - (A+C*B)$

(3 operands)

mul D, A, B  
mul E, C, B  
add E, A, E  
sub E, D, E

(2 operands)

mov D, A  
mul D, B  
mov E, C  
mul E, B  
add E, A  
sub E, D

# Memory-Memory: Pros and Cons

- pros
  - ⊙ requires fewer instructions (especially if 3 operands)
  - ⊙ easy to write compilers for (especially if 3 operands)
- cons
  - ⊙ very high memory traffic (especially if 3 operands)
  - ⊙ variable number of clocks per instruction
  - ⊙ with two operands, more data movements are required


# Register-Memory Architecture

- instruction set:

add R1, A   sub R1, A        mul R1, B  
load R1, A   store R1, A

- example:  $A * B - (A + C * B)$

load R1, A			
mul R1, B	/*	$A * B$	*/
store R1, D			
load R2, C			
mul R2, B	/*	$C * B$	*/
add R2, A	/*	$A + CB$	*/
sub R2, D	/*	$AB - (A + C * B)$	*/



$R1 = R1 +, -, *, / \text{ mem}[B]$




# Register-Memory: Pros and Cons

- pros
  - ⊙ some data can be accessed w/o loading first
  - ⊙ instruction format easy to encode
  - ⊙ good code density
- cons
  - ⊙ operands are not equivalent (execution time is determined by the slowest part)

# Load-Store (Register-Register) Architecture

- instruction set:  
add R1, R2, R3      sub R1, R2, R3    mul R1, R2, R3  
load R1, A            store R1, A      move R1, R2
- example:  $A * B - (A + C * B)$   
load R1, A  
load R2, B  
load R3, C  
mul R7, R3, R2      /\*      C\*B      \*/  
add R8, R7, R1      /\*      A + C\*B      \*/  
mul R9, R1, R2      /\*      A\*B      \*/  
sub R10, R9, R8      /\*      A\*B - (A+C\*B)      \*/



$R3 = R1 +, -, *, / R2$

# Register-Register: Pros and Cons

- pros
  - ⊙ simple, fixed length instruction encodings
  - ⊙ instructions take similar number of cycles
  - ⊙ relatively easy to pipeline and make superscalar
- cons
  - ⊙ higher instruction count
  - ⊙ not all instructions need three operands
  - ⊙ dependent on good compiler

# Registers: Advantages and Disadvantages (Review)

- advantages

- ⊙ faster than cache or main memory (no addressing mode or tags)
- ⊙ deterministic (no misses)
- ⊙ can replicate (multiple read ports)
- ⊙ short identifier (typically 3 to 8 bits)
- ⊙ reduce memory traffic

- disadvantages

- ⊙ need to save and restore on procedure calls and context switch
- ⊙ can't take the address of a register (for pointers)
- ⊙ fixed size (can't store strings or structures efficiently)
- ⊙ compiler must manage
- ⊙ limited number

Every ISA designed after 1980 uses a load-store ISA (e.g., RISC, to simplify CPU design).

# Instruction Format

32-bit RISC-V Instruction Formats

Instruction Formats	31	30	29	28	27	26	25	24	23	22	21	20	19	18	17	16	15	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0
Register/register	funct7							rs2					rs1					funct3			rd					opcode						
Immediate	imm[11:0]												rs1					funct3			rd					opcode						
Upper Immediate	imm[31:12]																				rd					opcode						
Store	imm[11:5]							rs2					rs1					funct3			imm[4:0]					opcode						
Branch	[12]	imm[10:5]							rs2					rs1					funct3			imm[4:1]			[11]	opcode						
Jump	[20]	imm[10:1]										[11]	imm[19:12]							rd					opcode							

- **opcode (7 bit):** partially specifies which of the 6 types of *instruction formats*
- **funct7 + funct3 (10 bit):** combined with **opcode**, these two fields describe what operation to perform
- **rs1 (5 bit):** specifies register containing first operand
- **rs2 (5 bit):** specifies second register operand
- **rd (5 bit):** Destination register specifies register which will receive result of computation

# Types of Operations

- arithmetic and logic: AND, ADD
- data transfer: MOVE, LOAD, STORE
- control: BRANCH, JUMP, CALL
- system: OS CALL, VM
- floating point: ADDF, MULF, DIVF
- decimal: ADDD, CONVERT
- string: MOVE, COMPARE
- graphics: (DE)COMPRESS

# Conditional branch

- how do you specify the **destination** of a branch/jump?
  - ⊙ theoretically, the destination is a full address
    - 16 bits for LC3b
    - 32 bits for MIPS, RISC-V
- studies show that almost all conditional branches go **short distances** from the current program counter (loops, if-then-else)
  - ⊙ we can specify a relative address in much fewer bits than an absolute address
  - ⊙ e.g., `beq $1, $2, 100` => if ( $\$1 == \$2$ )  $PC = PC + 100 * 4$

# Jumps

- need to be able to jump to an absolute address sometimes
  - ⦿ jump -- j 10000 => PC = 10000
- need to be able to do procedure calls and returns
  - ⦿ jump and link--jal 100000 => \$31 = PC + 4; PC = 10000
    - used for procedure calls
  - ⦿ jump register -- jr \$31 => PC = \$31
    - used for returns, but can be useful for lots of other things

OP	target
----	--------



# ISA Tradeoffs (An Example of MIPS)

	6 bits	5 bits	5 bits	5 bits	5 bits	6 bits
r format	OP	rs	rt	rd	sa	funct
i format	OP	rs	rt	immediate		
j format	OP	target				

- what if?
  - ⊙ 64 registers
  - ⊙ 20-bit immediate
  - ⊙ 4 operand instruction (e.g.,  $Y = AX + B$ )


# Byte Ordering

- how should bytes within multi-byte word be ordered in memory?
- conventions
  - ⊙ Sun's, Mac's are "big endian" machines
    - least significant byte has highest address
  - ⊙ Alphas, PC's are "little endian" machines
    - least significant byte has lowest address


# Byte Ordering Example

- big endian
  - ⊙ least significant byte has highest address
- little endian
  - ⊙ least significant byte has lowest address
- example
  - ⊙ variable x has 4-byte representation 0x01234567
  - ⊙ address given by &x is 0x100

big endian



little endian



# Reading Byte-Reversed Listings

- Disassembly
  - ⦿ text representation of binary machine code
  - ⦿ generated by program that reads the machine code
- Example fragment

address	instruction code	assembly rendition
8048365:	5b	pop %ebx
8048366:	81 c3 ab 12 00 00	add \$0x12ab,%ebx
804836c:	83 bb 28 00 00 00 00	cmpl \$0x0,0x28(%ebx)

- Deciphering numbers
  - ⦿ value:
  - ⦿ pad to 4 bytes:
  - ⦿ split into bytes:
  - ⦿ reverse:

0x12ab  
0x000012ab  
00 00 12 ab  
ab 12 00 00

# Takeaways

Function of ISA

ISA Design Issues

Classical ISAs: Their Pros and Cons

RISC-V Instruction Format

Practice RISC-V in your MPs

Byte ordering: big-endian vs. little-endian

# Announcement

- Next lecture
  - ⦿ performance, energy, and power metric: [HP2] Chapter 1.5, 1.8
- MP assignment
  - ⦿ MP1 deadline: 9/1/2022