




# Superscalar Complexity

---

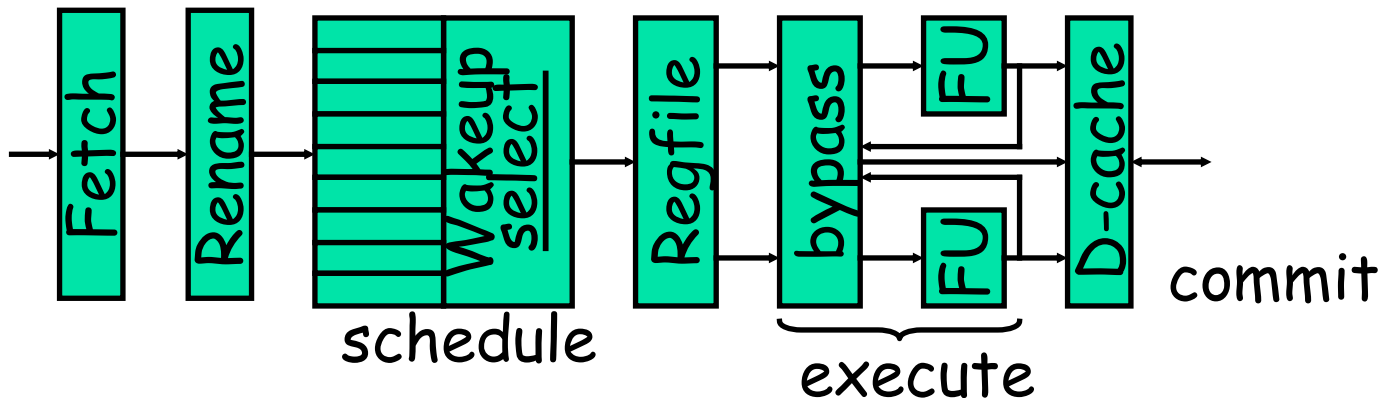
# View of Superscalar Execution



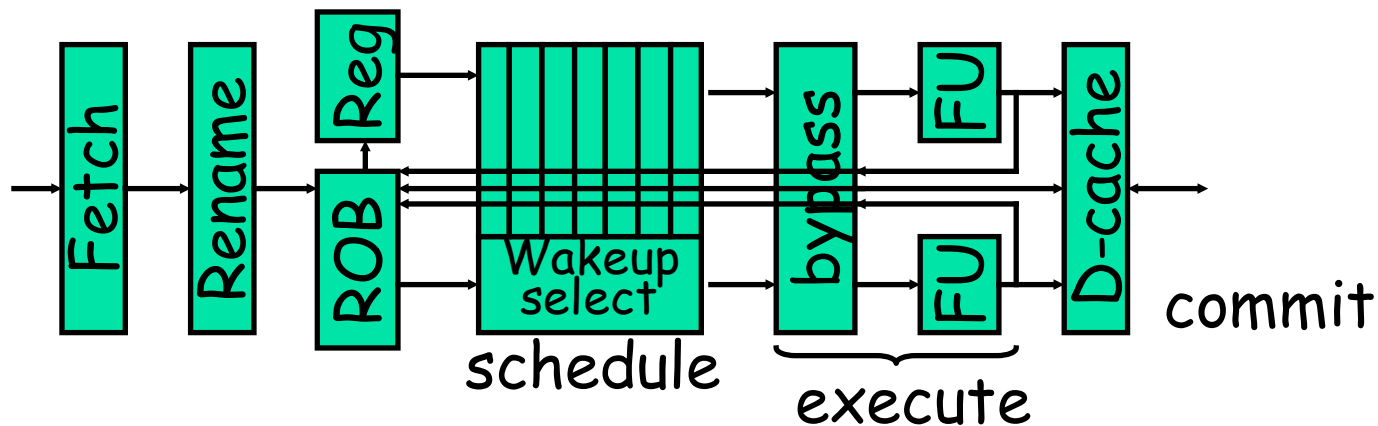
# Generic Superscalar Processor Models



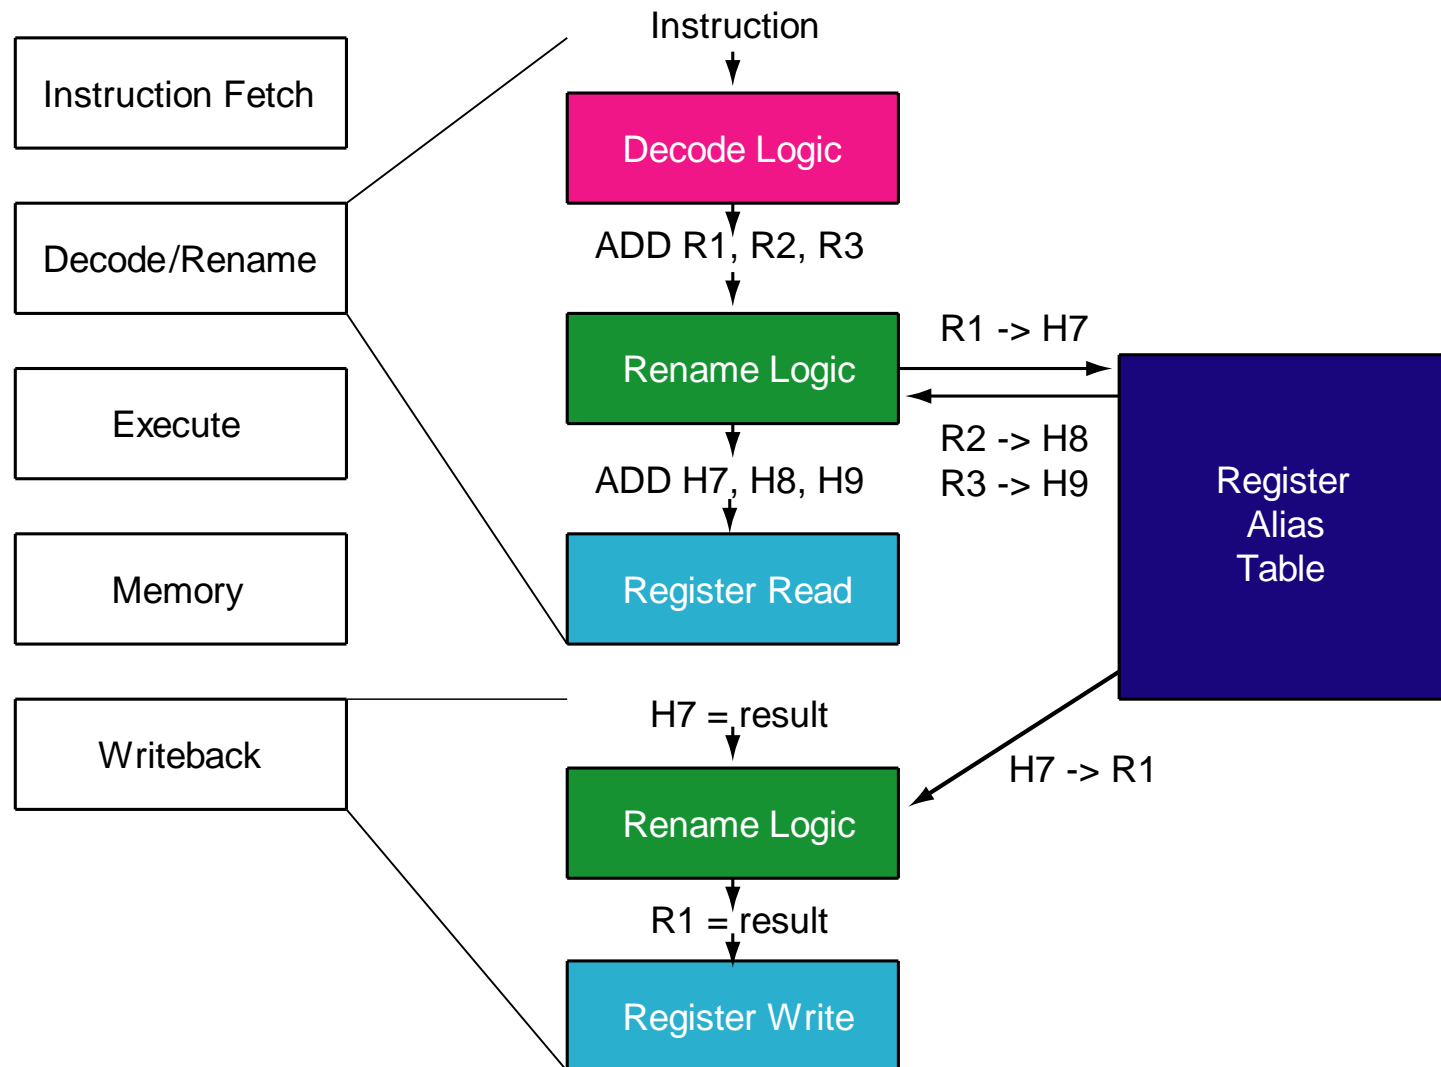

Issue queue based




Reservation based



# Register Renaming



# Register Renaming



- Alpha 21264+, MIPS R10K+, Pentium 4 use explicit register renaming
  - Registers are not read until instruction dispatches (begins execution)
  - Register renaming ensures no conflicts

```


DIV R5, R4, R2
ADD R7, R5, R1
SUB R5, R3, R2
LW  R7, 1000(R5)

```


RAT

R1	PR23
R2	PR2
R3	PR17
R4	PR45
R5	PR13
R6	PR20
R7	PR30

...



# Register Renaming



- Assume that **PR37** is free and allocated to DIV's destination register R5

```
DIV R5, R4, R2
ADD R7, R5, R1
SUB R5, R3, R2
LW  R7, 1000(R5)
```


```
DIV PR37, PR45, PR2
```

RAT


R1	PR23
R2	PR2
R3	PR17
R4	PR45
R5	<del>PR13</del>
R6	PR20
R7	PR30

**PR37**

...



# Register Renaming



- Assume that **PR37** is free and allocated to DIV's destination register R5


```
DIV R5, R4, R2
ADD R7, R5, R1
SUB R5, R3, R2
LW  R7, 1000(R5)
```

```
DIV PR37, PR45, PR2
```


RAT

R1	PR23
R2	PR2
R3	PR17
R4	PR45
R5	<b>PR37</b>
R6	PR20
R7	PR30

...



# Register Renaming



- Assume that **PR4** is free and allocated to ADD's destination register R7

```
DIV R5, R4, R2
ADD R7, R5, R1
SUB R5, R3, R2
LW  R7, 1000(R5)
```


```
DIV PR37, PR45, PR2
ADD PR4 , PR37, PR23
```

RAT

R1	PR23
R2	PR2
R3	PR17
R4	PR45
R5	<b>PR37</b>
R6	PR20
R7	<b>PR30</b>


**PR4**

...





# Register Renaming



```
DIV R5, R4, R2
ADD R7, R5, R1
SUB R5, R3, R2
LW  R7, 1000(R5)
```


```
DIV PR37, PR45, PR2
ADD PR4 , PR37, PR23
SUB PR42, PR17, PR2
```

RAT


R1	PR23
R2	PR2
R3	PR17
R4	PR45
R5	<del>PR37</del>
R6	PR20
R7	<del>PR4</del>

PR42

...



# Register Renaming



```
DIV R5, R4, R2
ADD R7, R5, R1
SUB R5, R3, R2
LW  R7, 1000(R5)
```


```
DIV PR37, PR45, PR2
ADD PR4 , PR37, PR23
SUB PR42, PR17, PR2
LW  PR19, 1000(PR42)
```

RAT

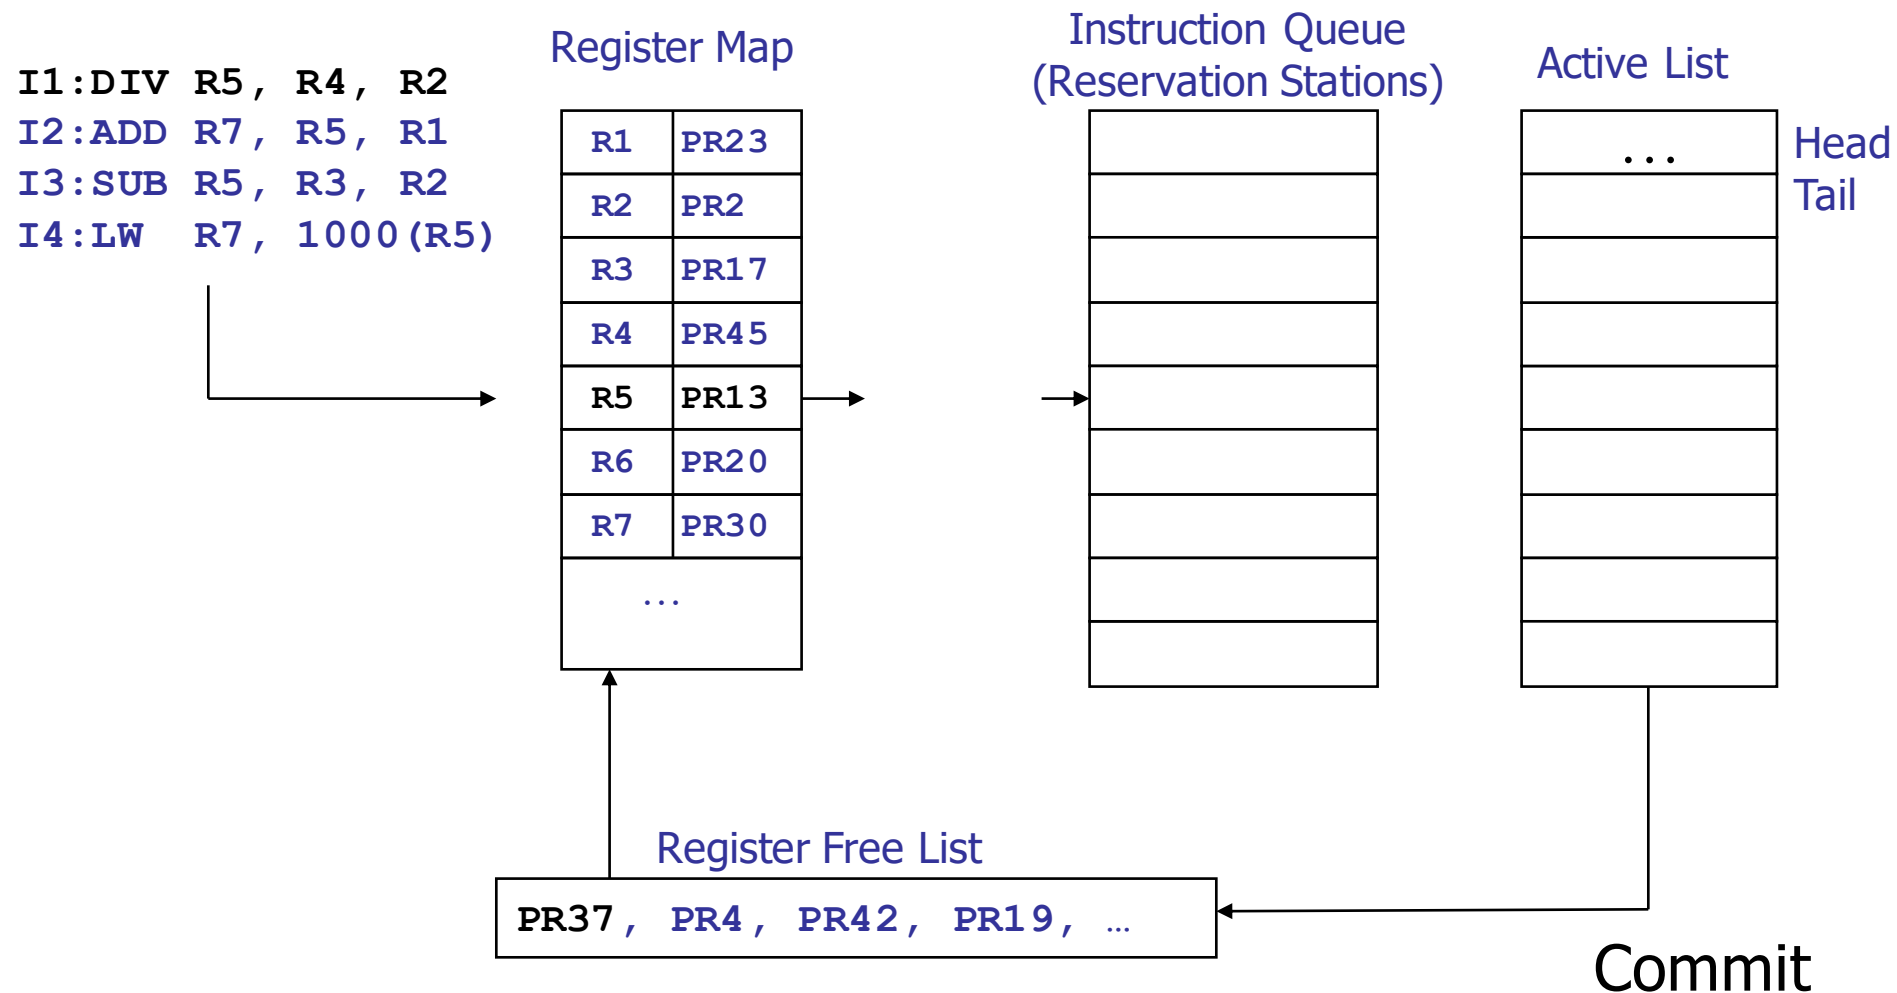

R1	PR23
R2	PR2
R3	PR17
R4	PR45
R5	PR42
R6	PR20
R7	<del>PR4</del>

PR19

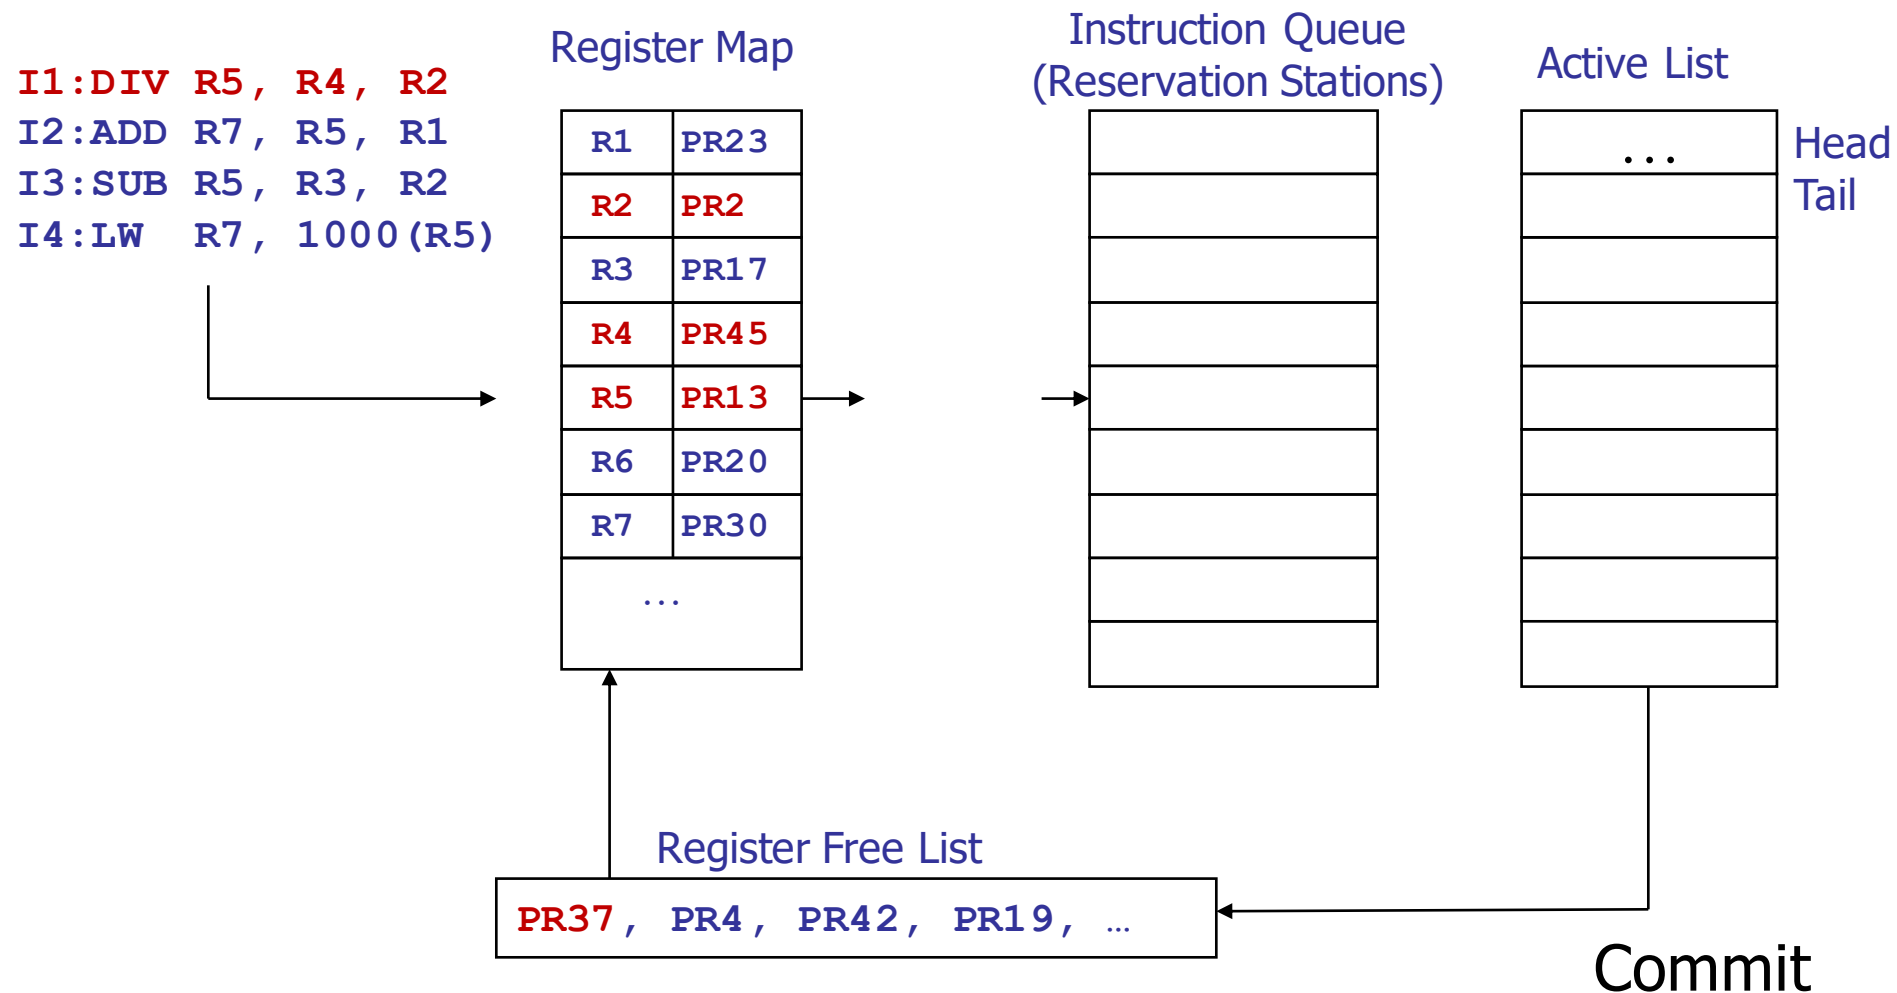

...



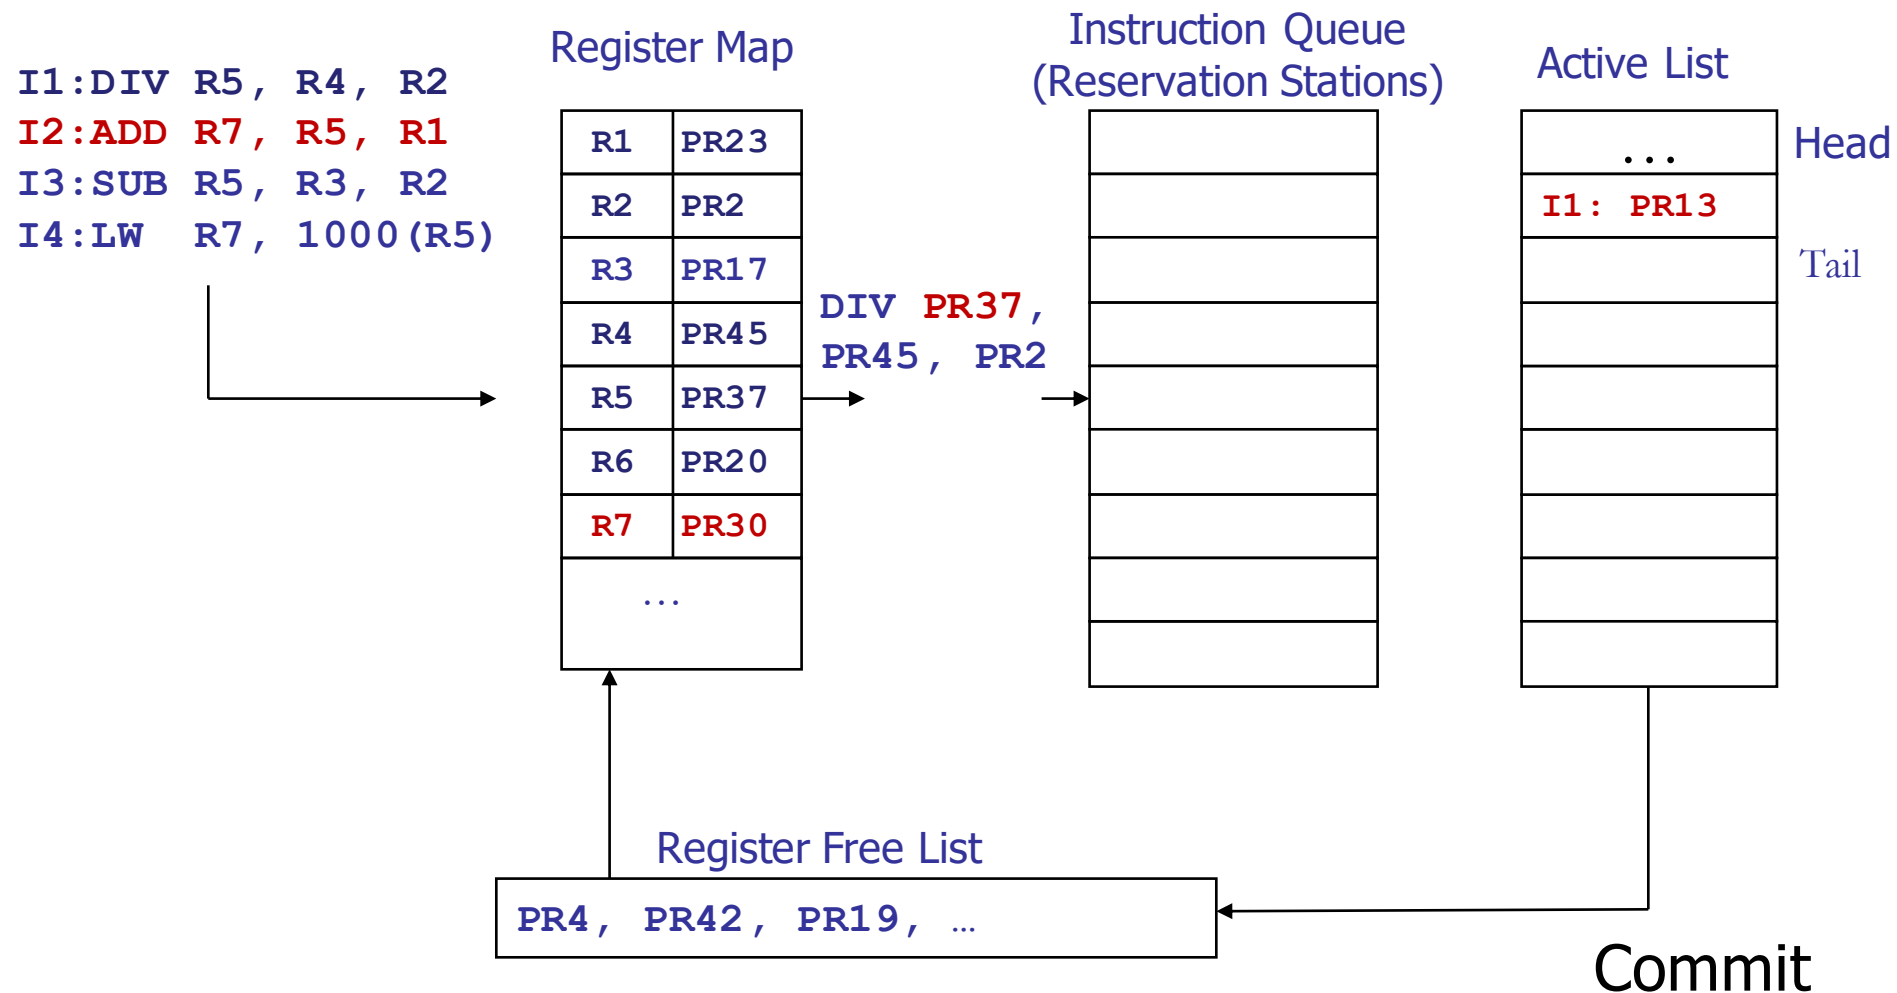

# MIPS R10000, Recycling Physical Registers



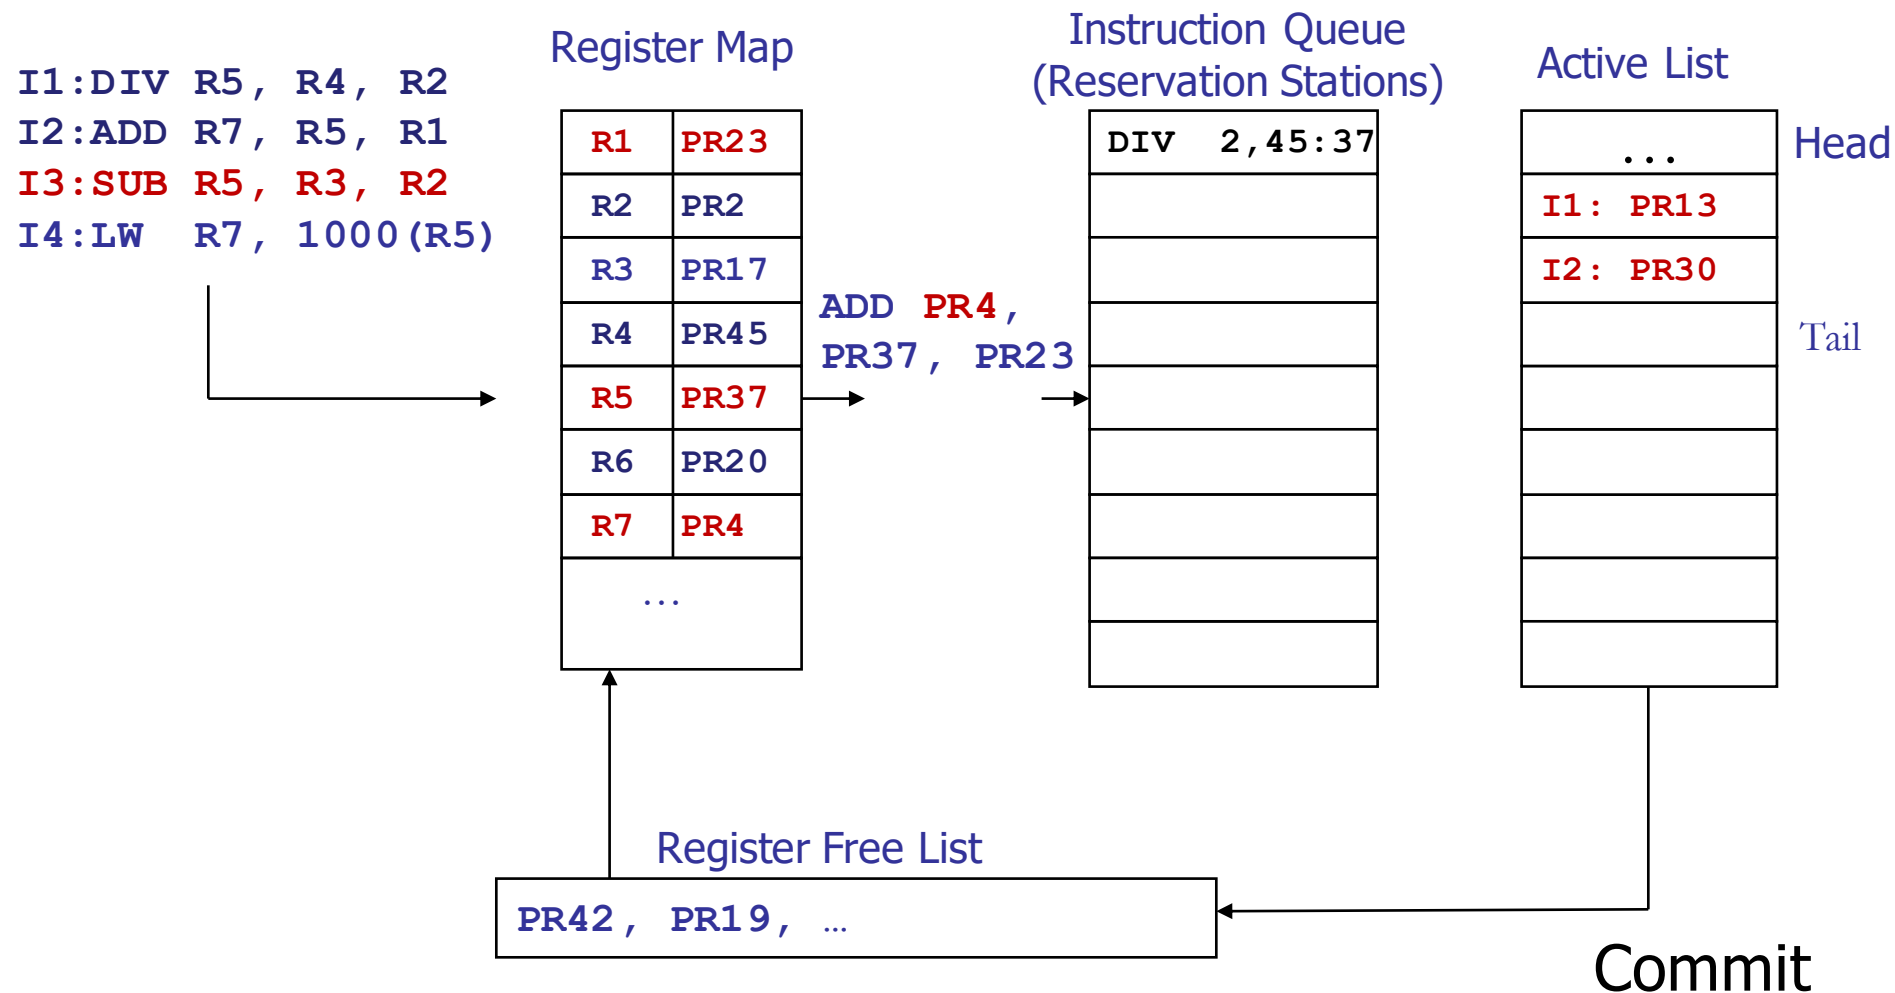

# MIPS R10000, Recycling Physical Registers



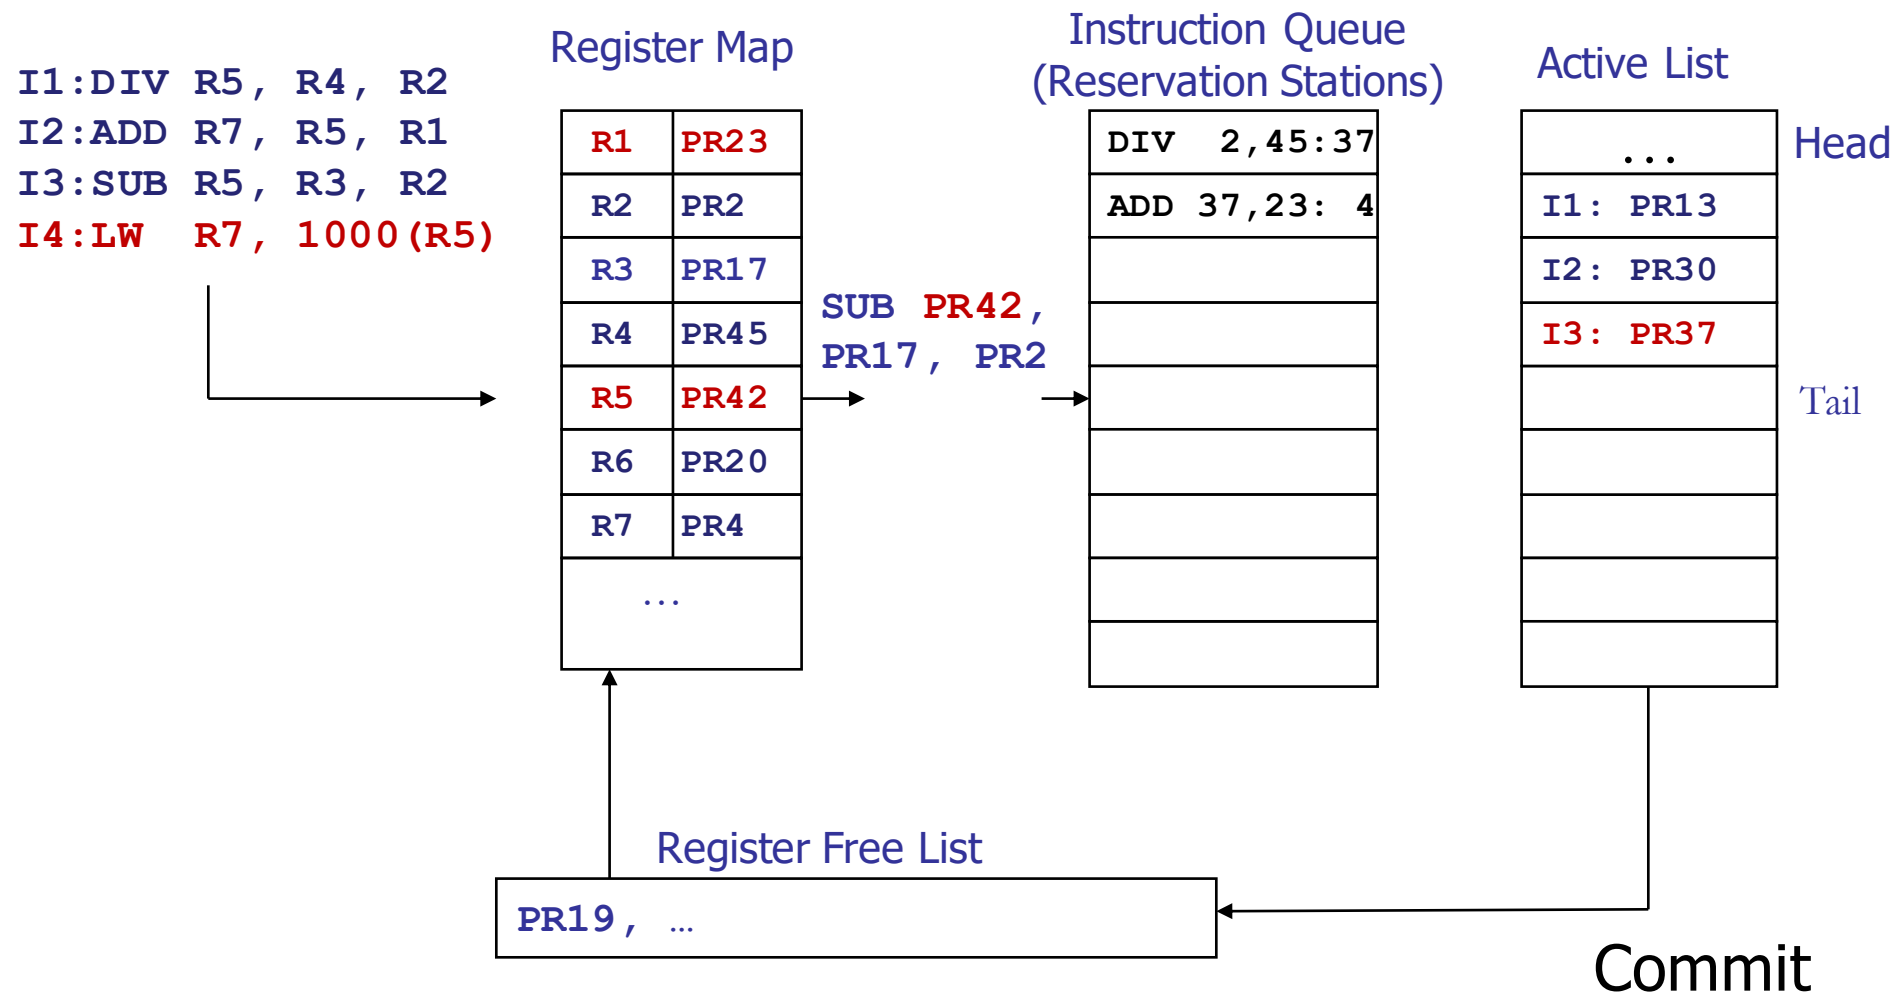

# MIPS R10000, Recycling Physical Registers



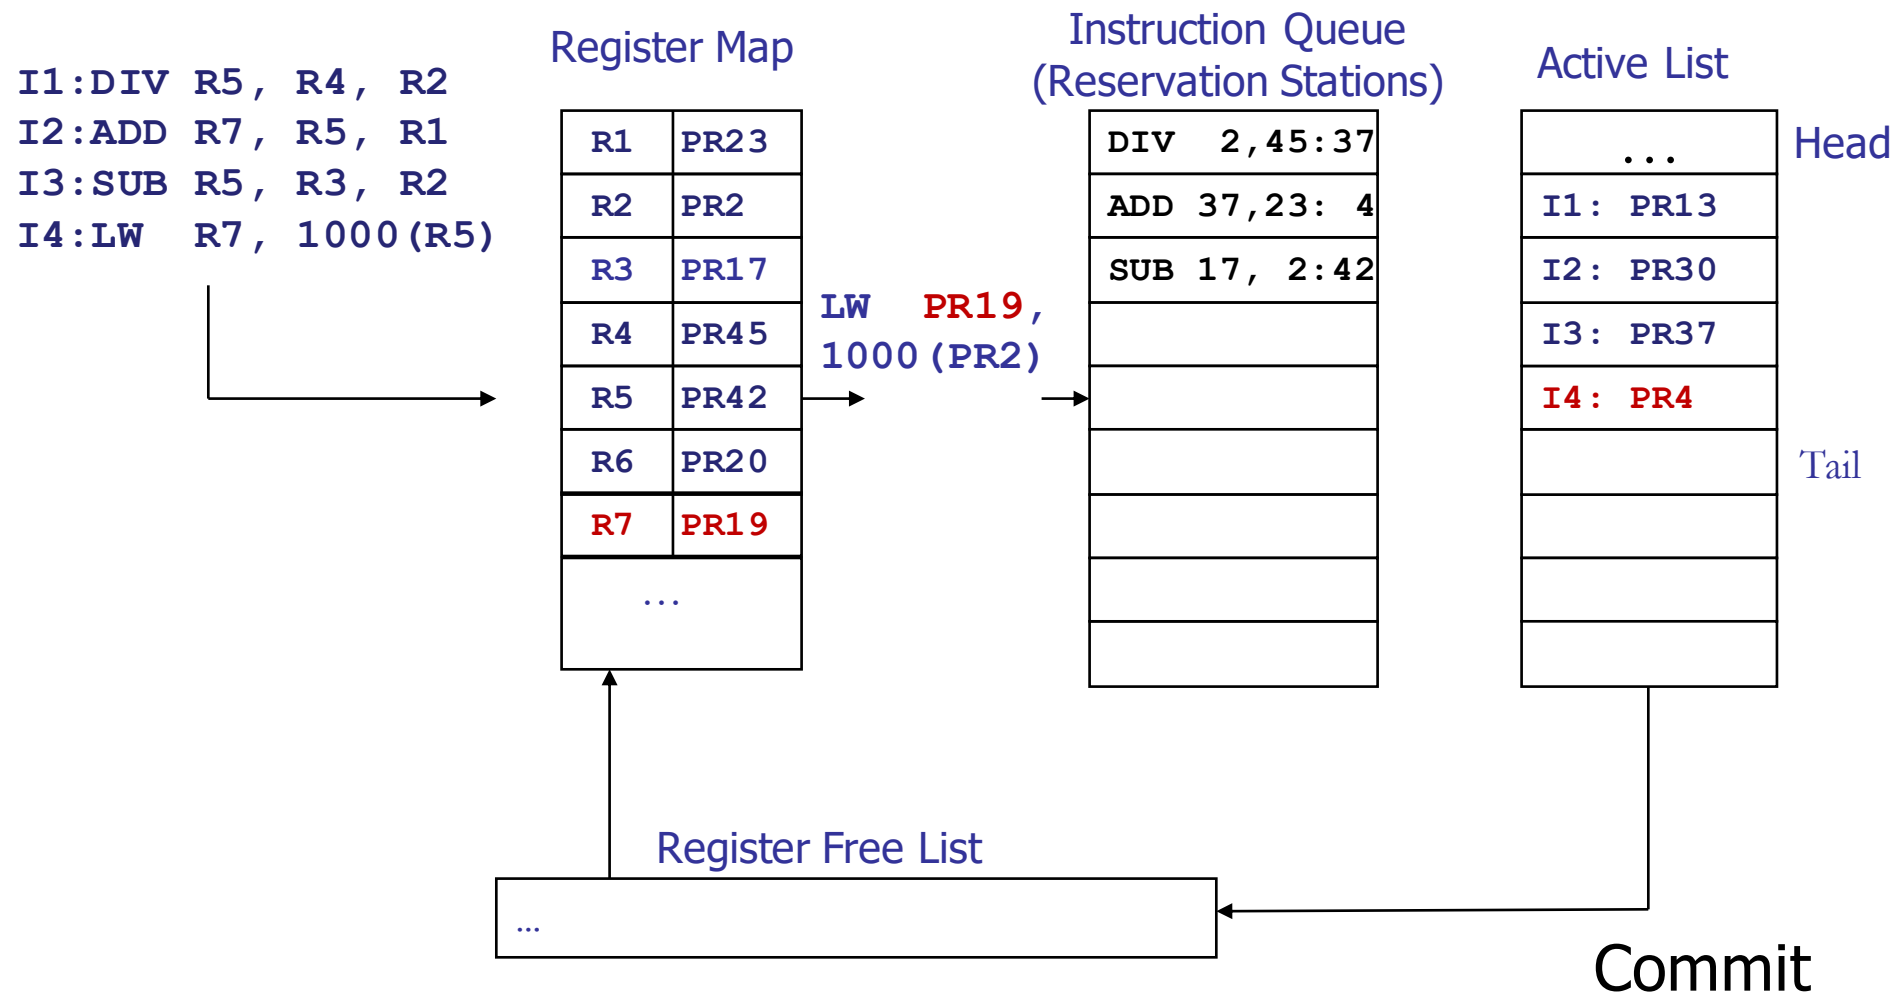

# MIPS R10000, Recycling Physical Registers



# MIPS R10000, Recycling Physical Registers

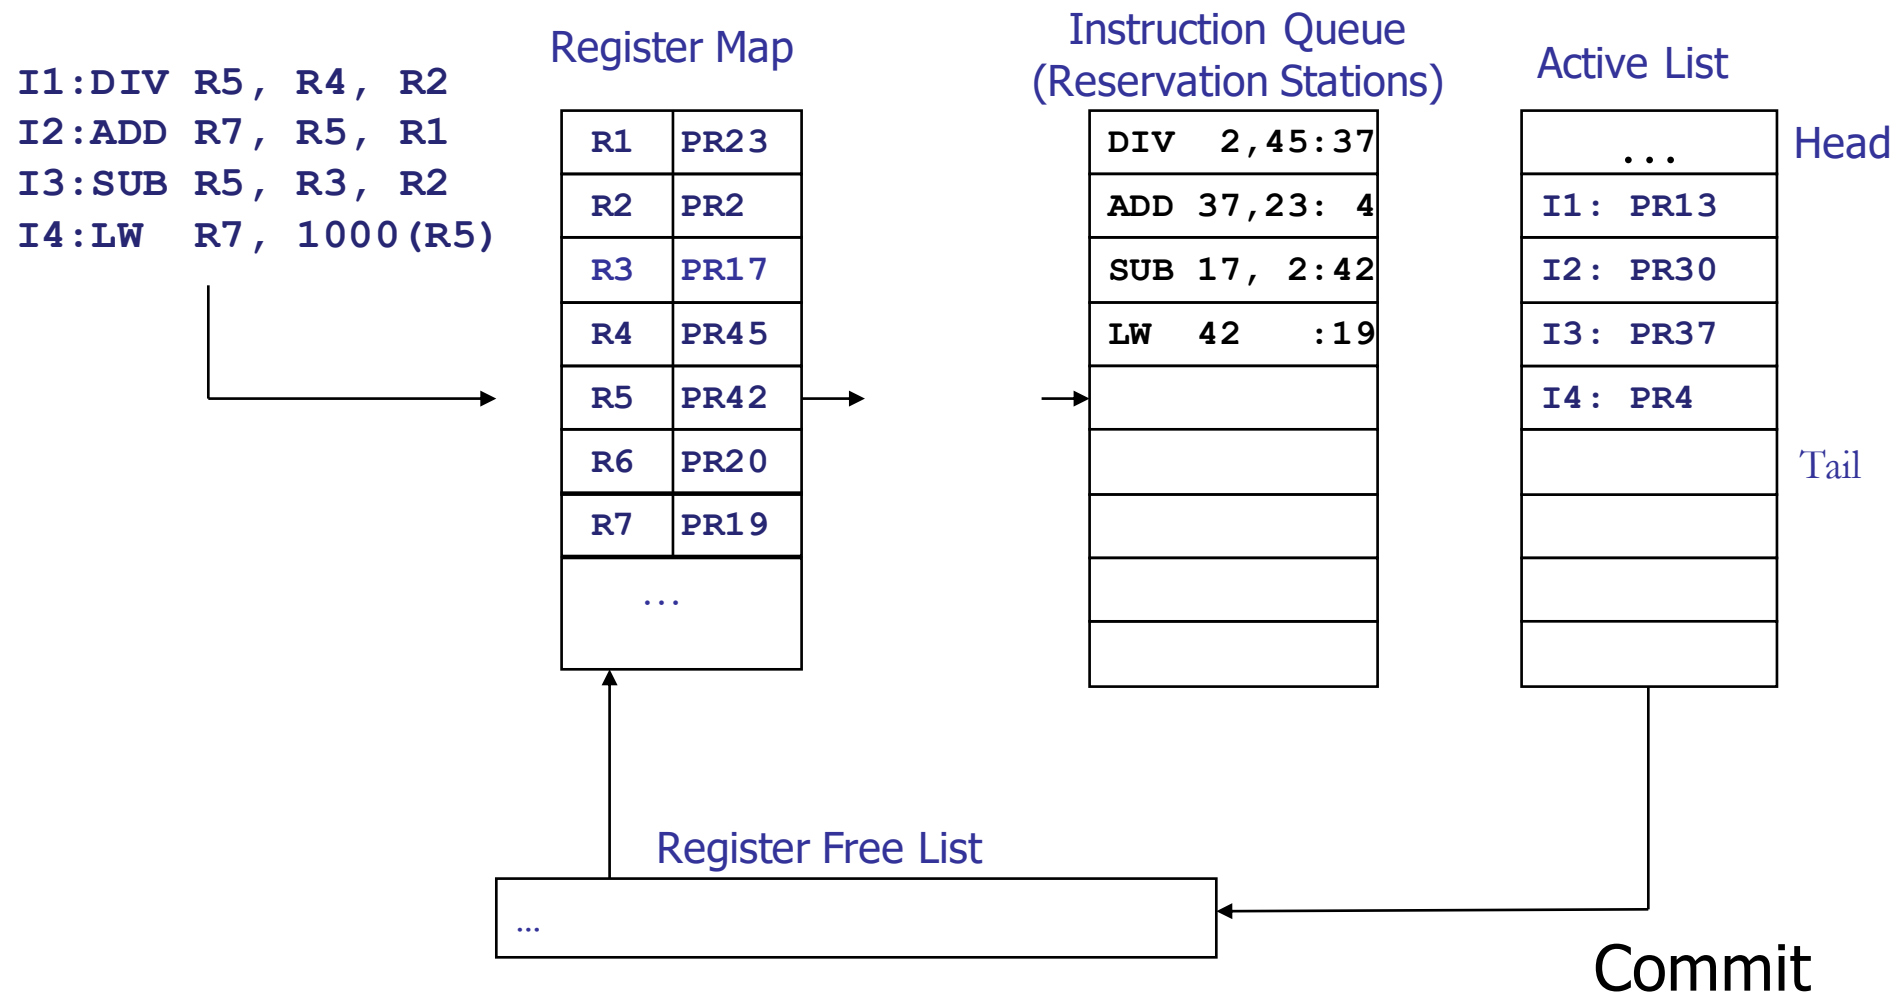



# MIPS R10000, Recycling Physical Registers

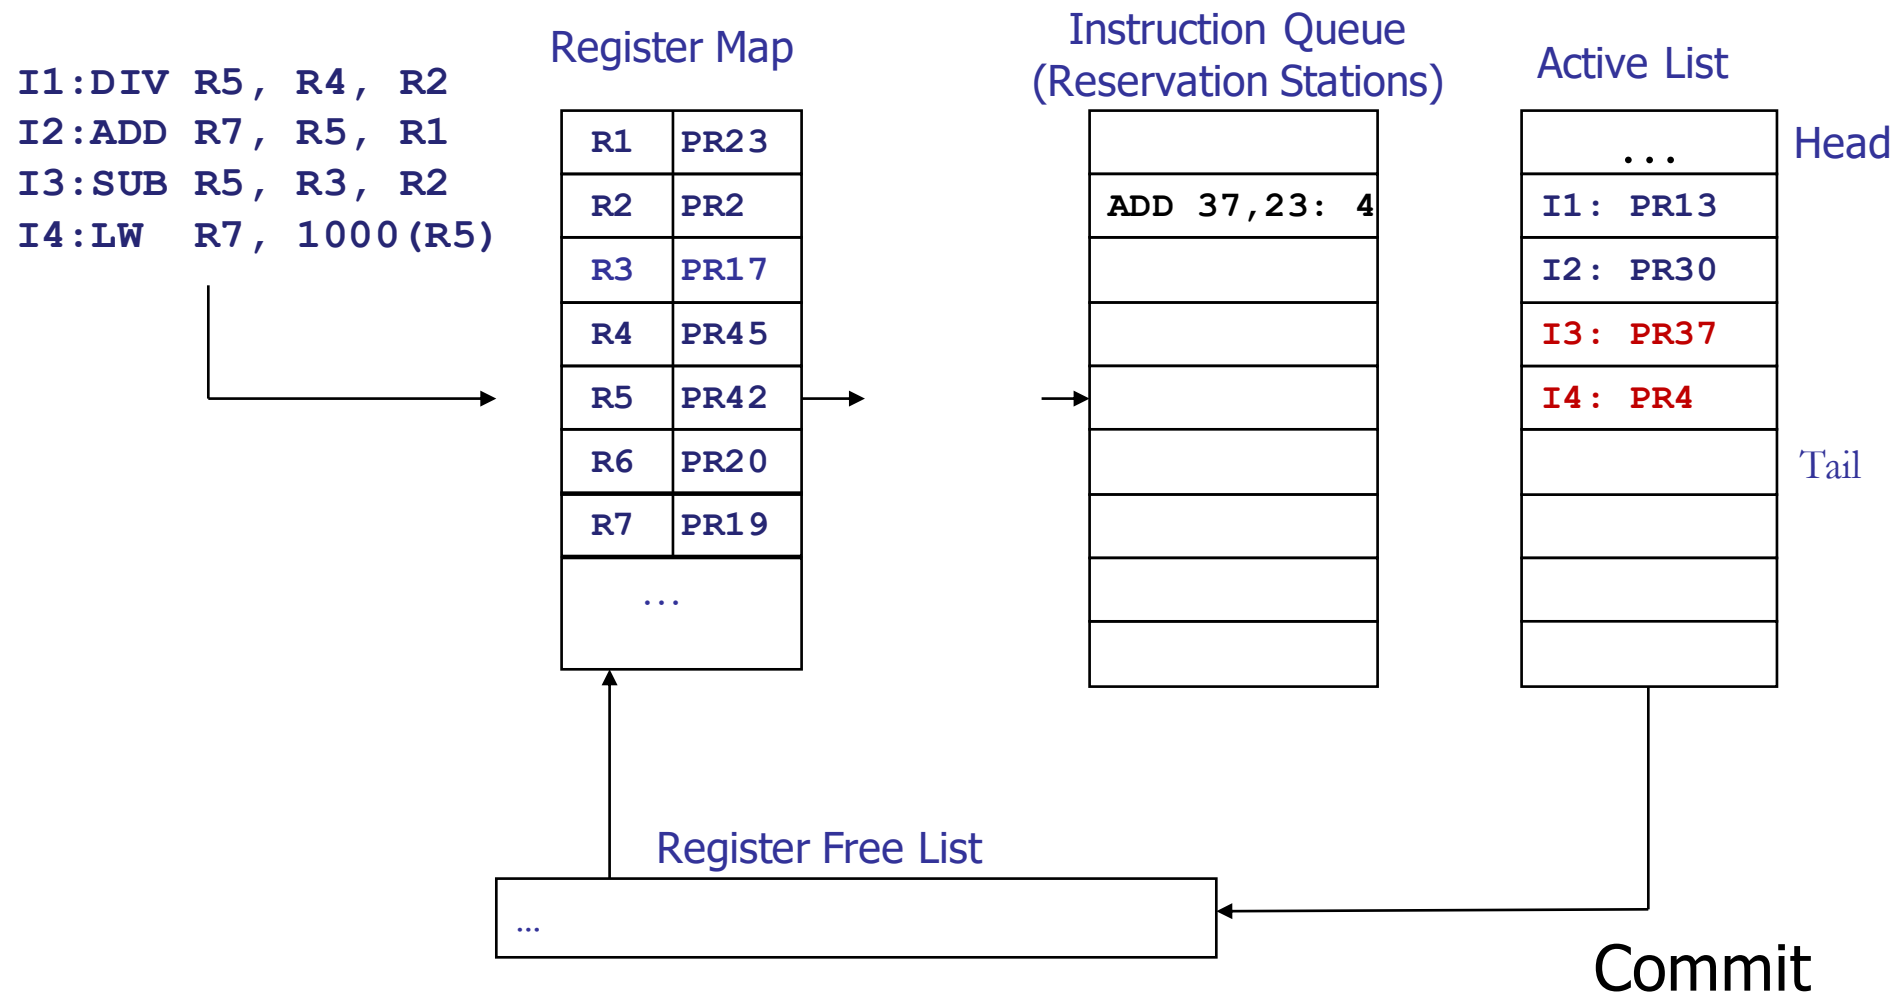





# MIPS R10000, Recycling Physical Registers

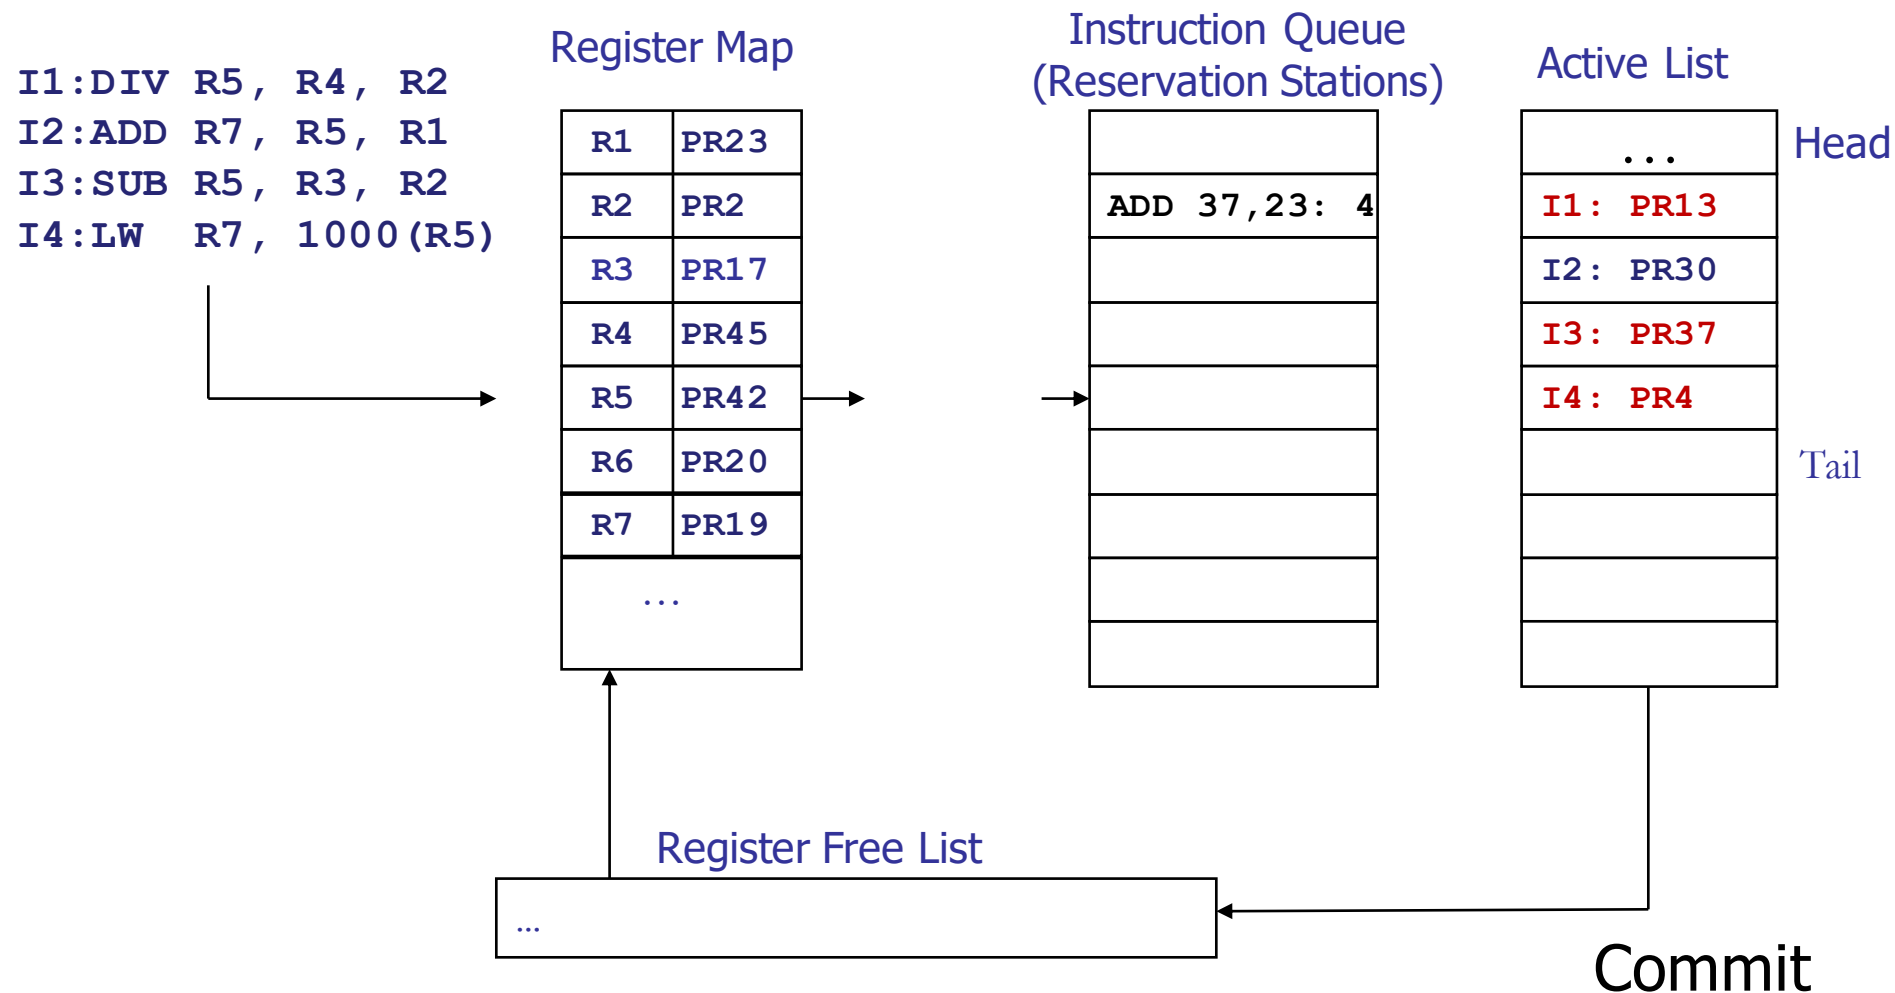



# MIPS R10000, Recycling Physical Registers



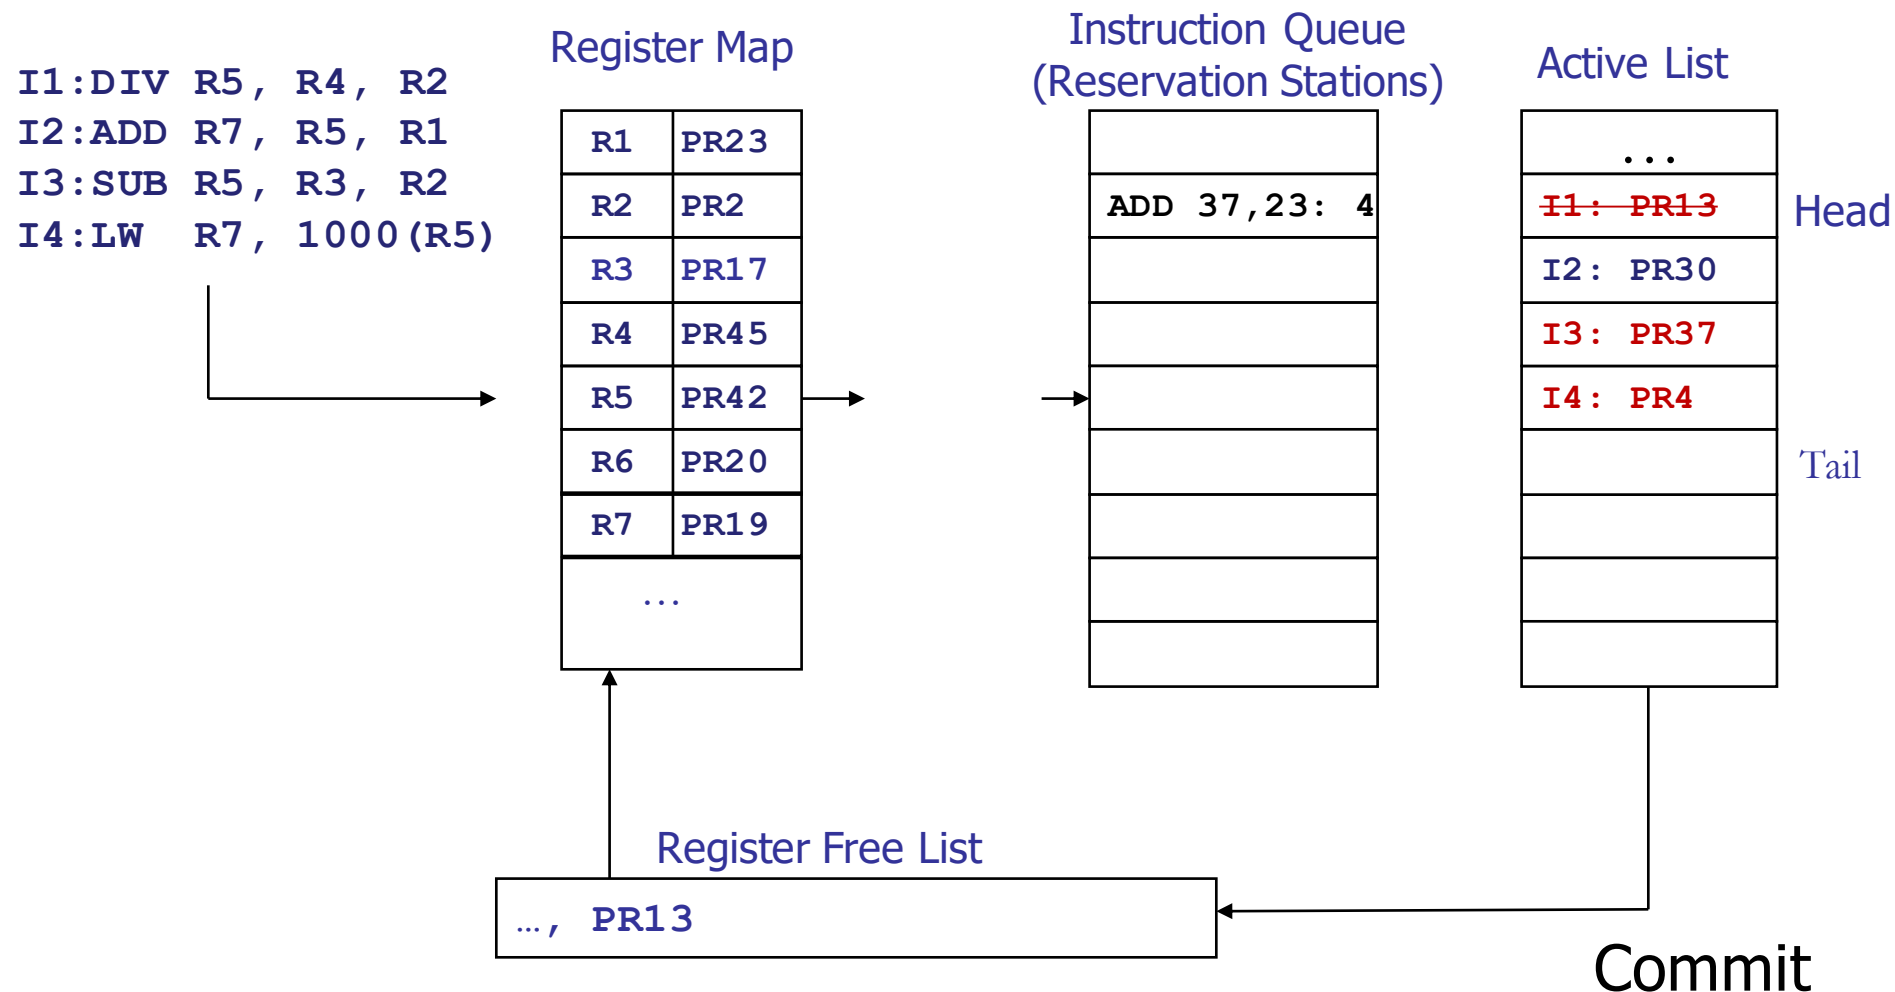

**I4**, producing register **19**, completes, broadcasts a completion signal to IQ

# MIPS R10000, Recycling Physical Registers




I1, producing register 37, completes, broadcasts a completion signal to IQ

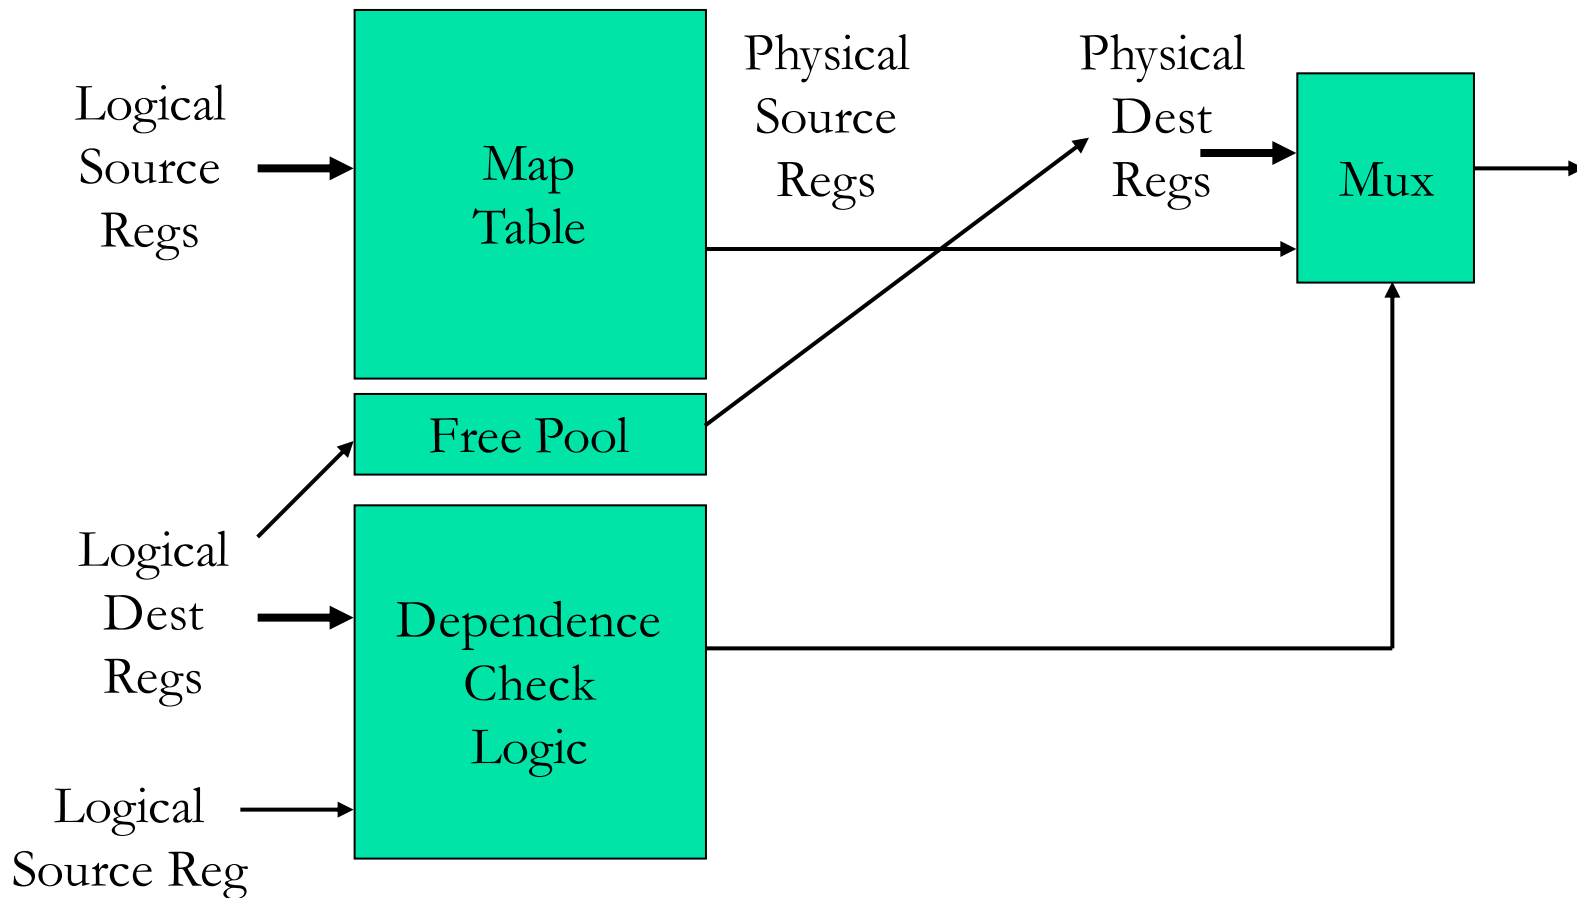
# MIPS R10000, Recycling Physical Registers




I2, producing register 4, completes, broadcasts a completion signal to IQ  
 I1, commits




# Register Rename Logic




# Instr Scheduling: Wakeup & Select



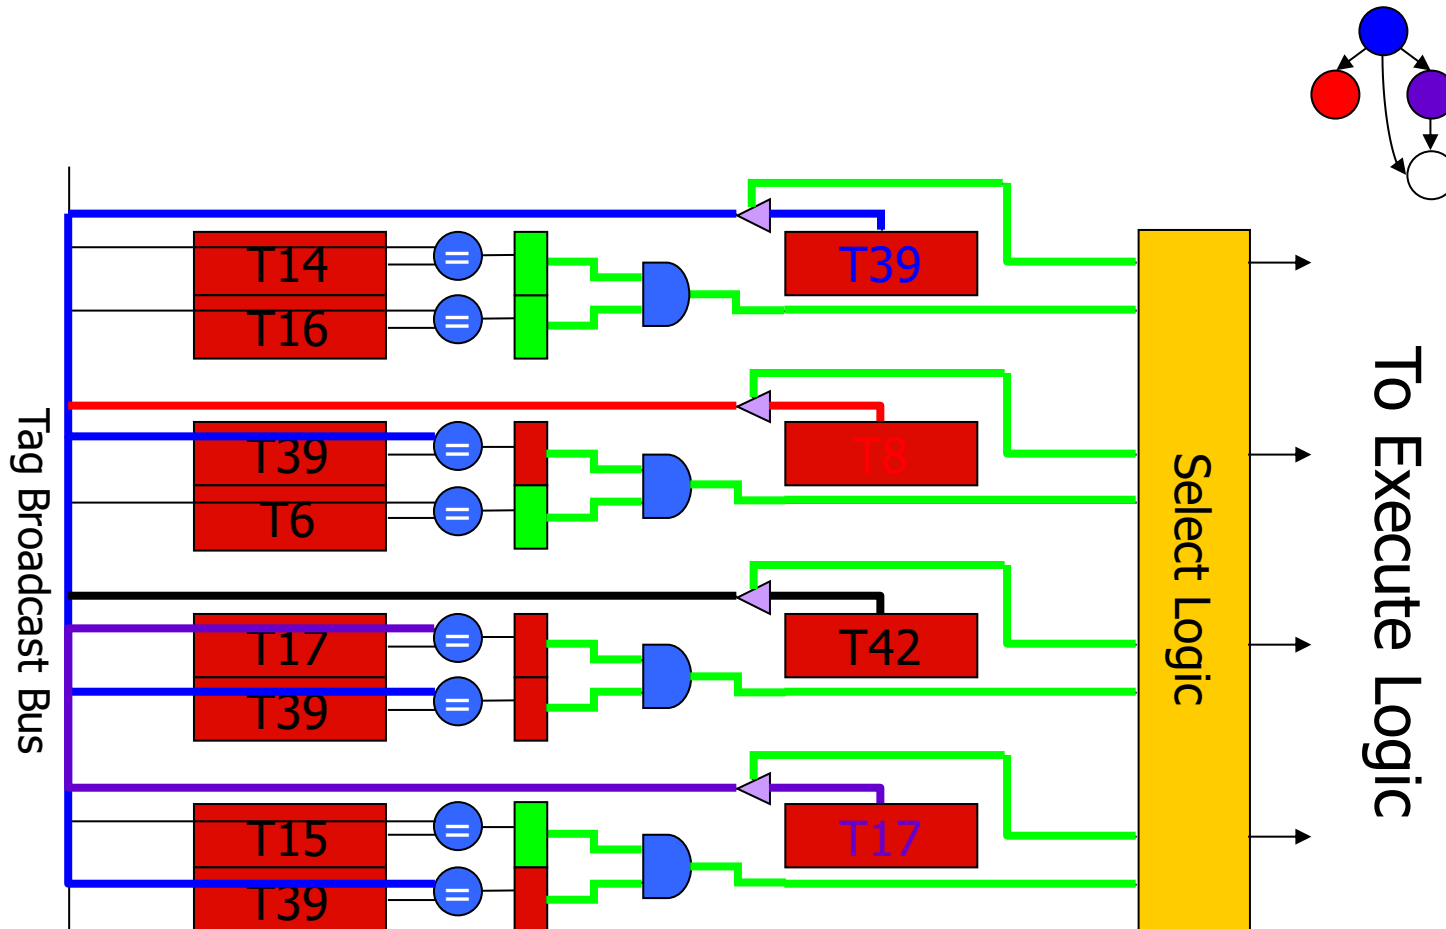

- Wakeup Logic
  - To notify the resolution of data dependency of input operands
  - Wake up instructions with zero input dependency
- Select Logic
  - Choose and fire ready instructions
  - Deal with structure hazard
- Wakeup-select is likely on the critical path
  - Associative match



# Wakeup Logic




# Scalar Scheduler (Issue Width = 1)

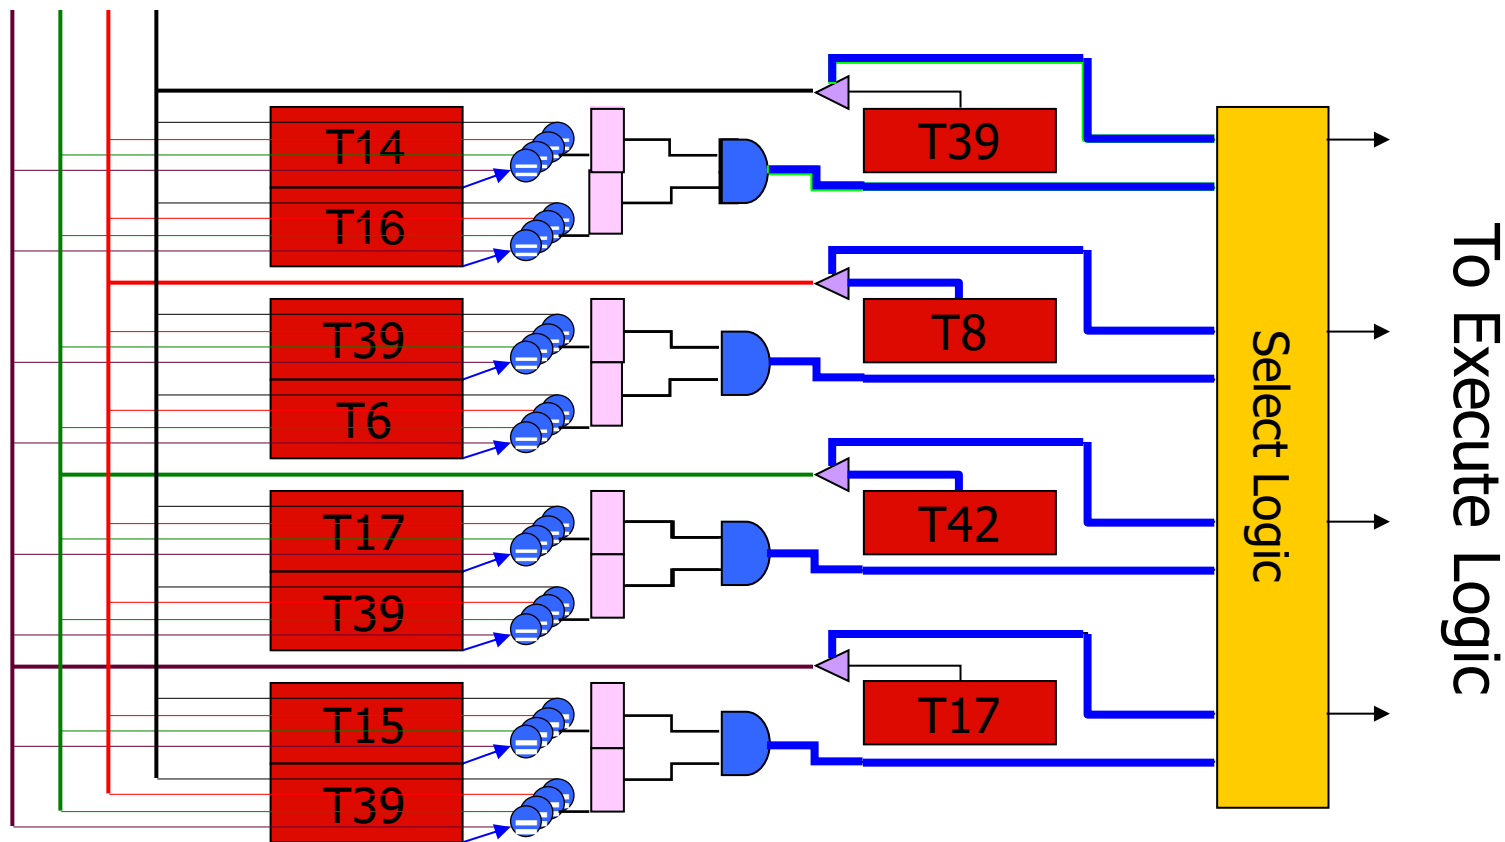





# Superscalar Scheduler (Width = 4)



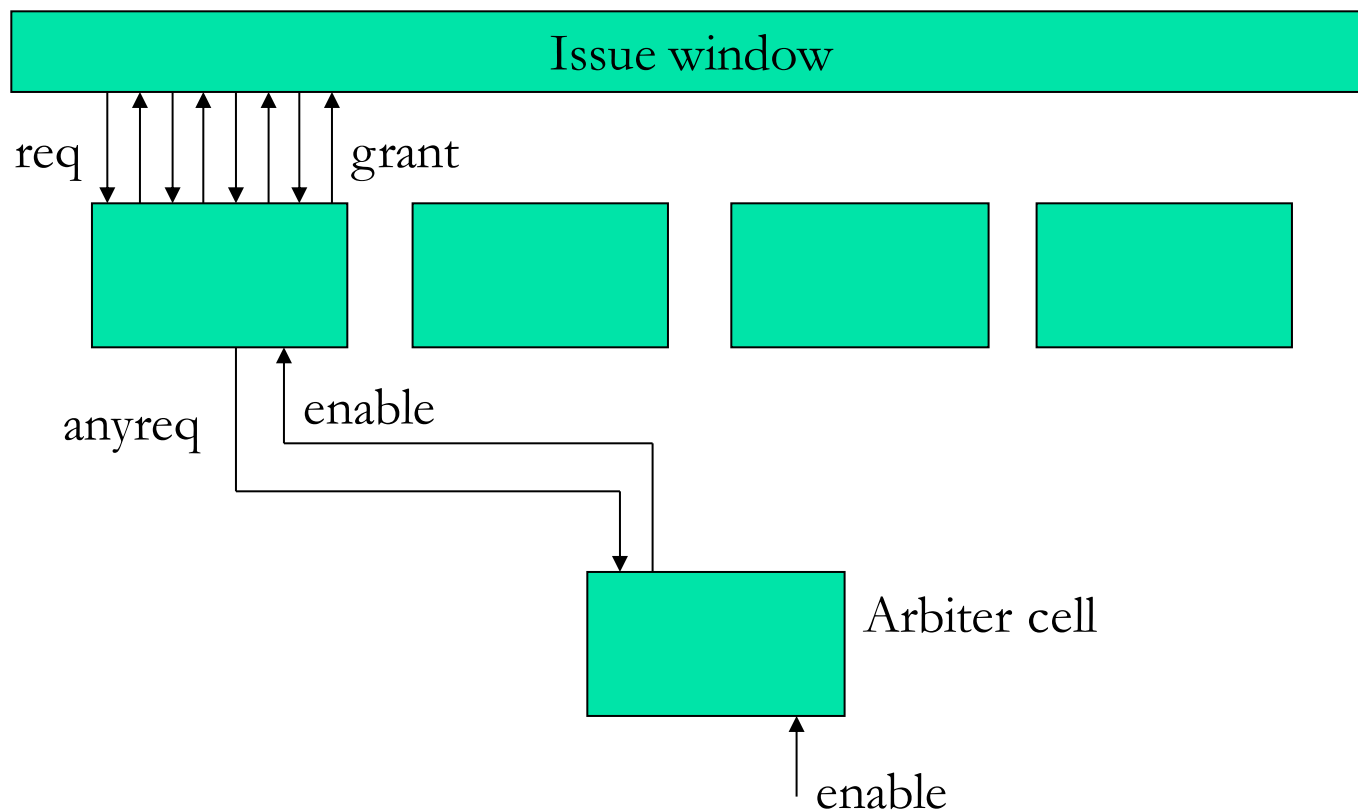
Tag Broadcast Bus [3..0]



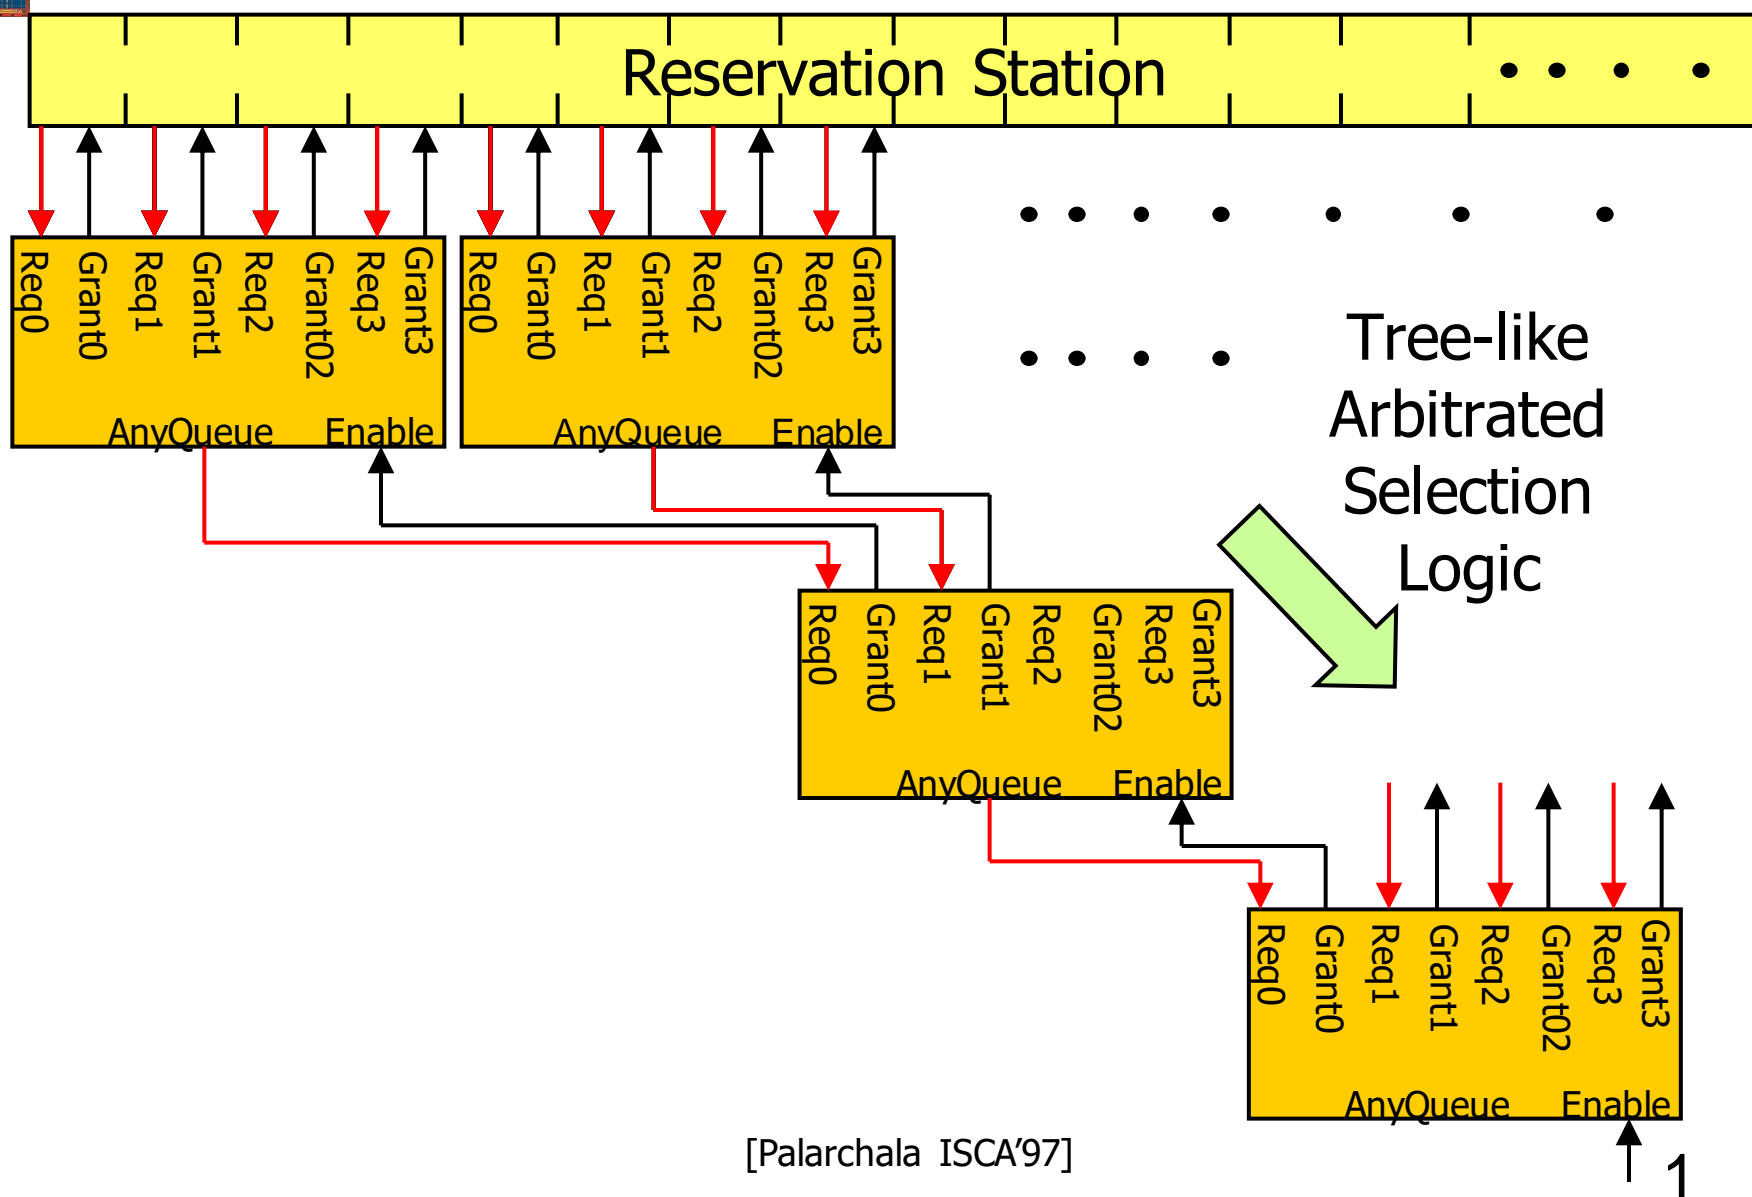
Snapshot of RS (only 4 entries shown)



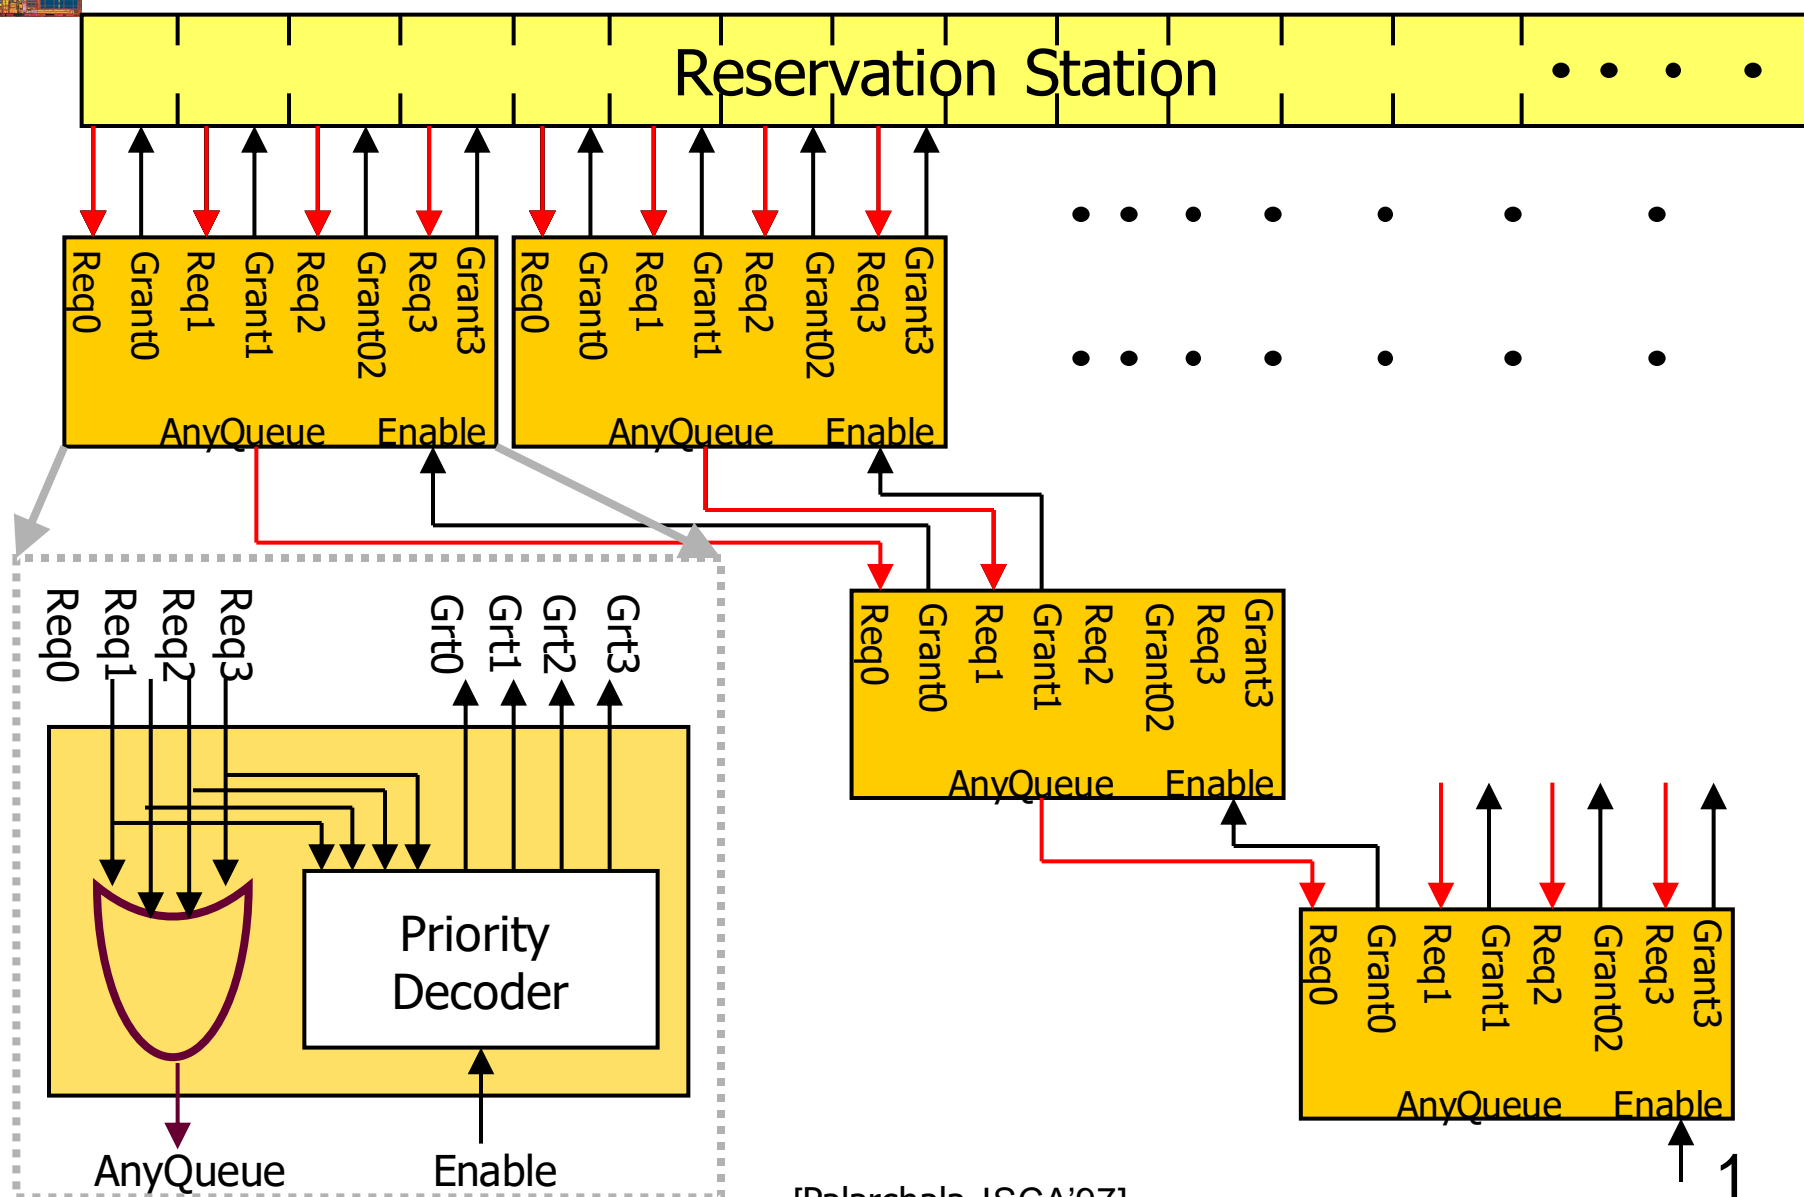

# Selection Logic



# Simple (?) Select Logic Implementation

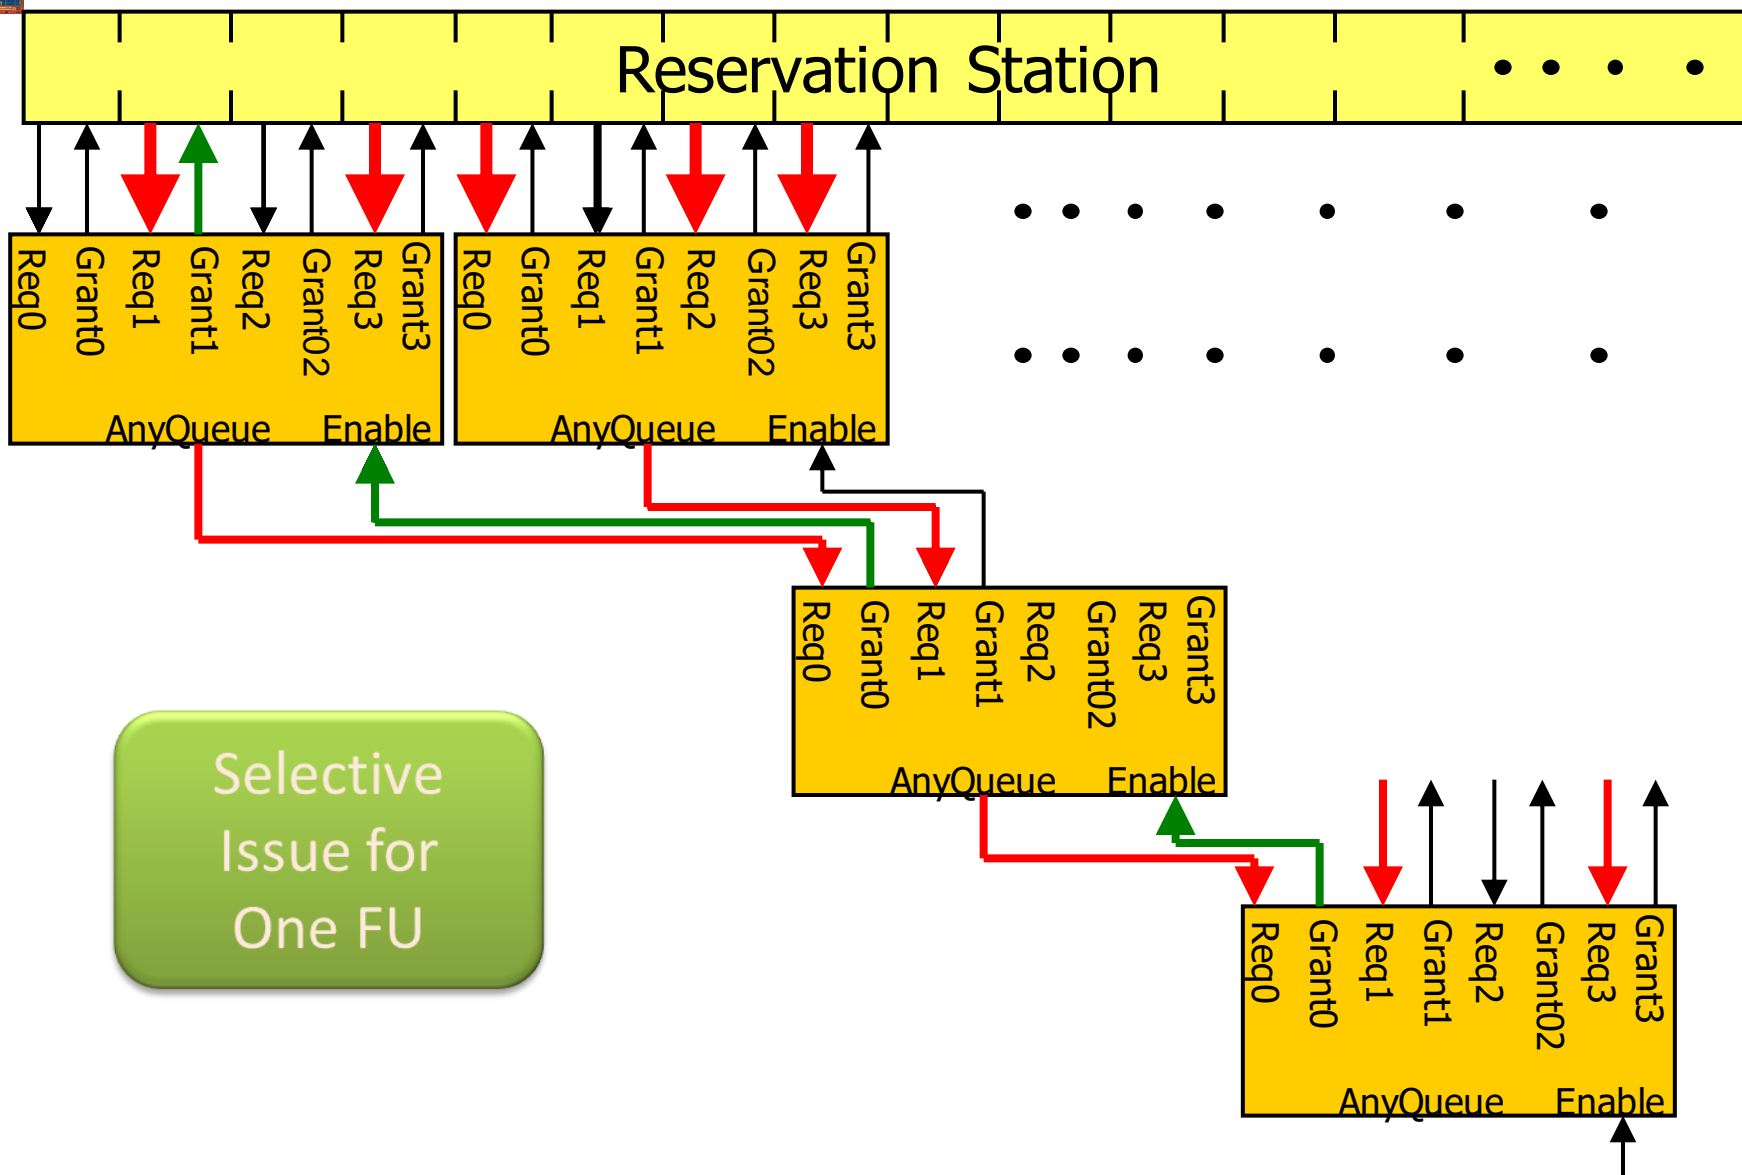


# Simple Select Logic Implementation






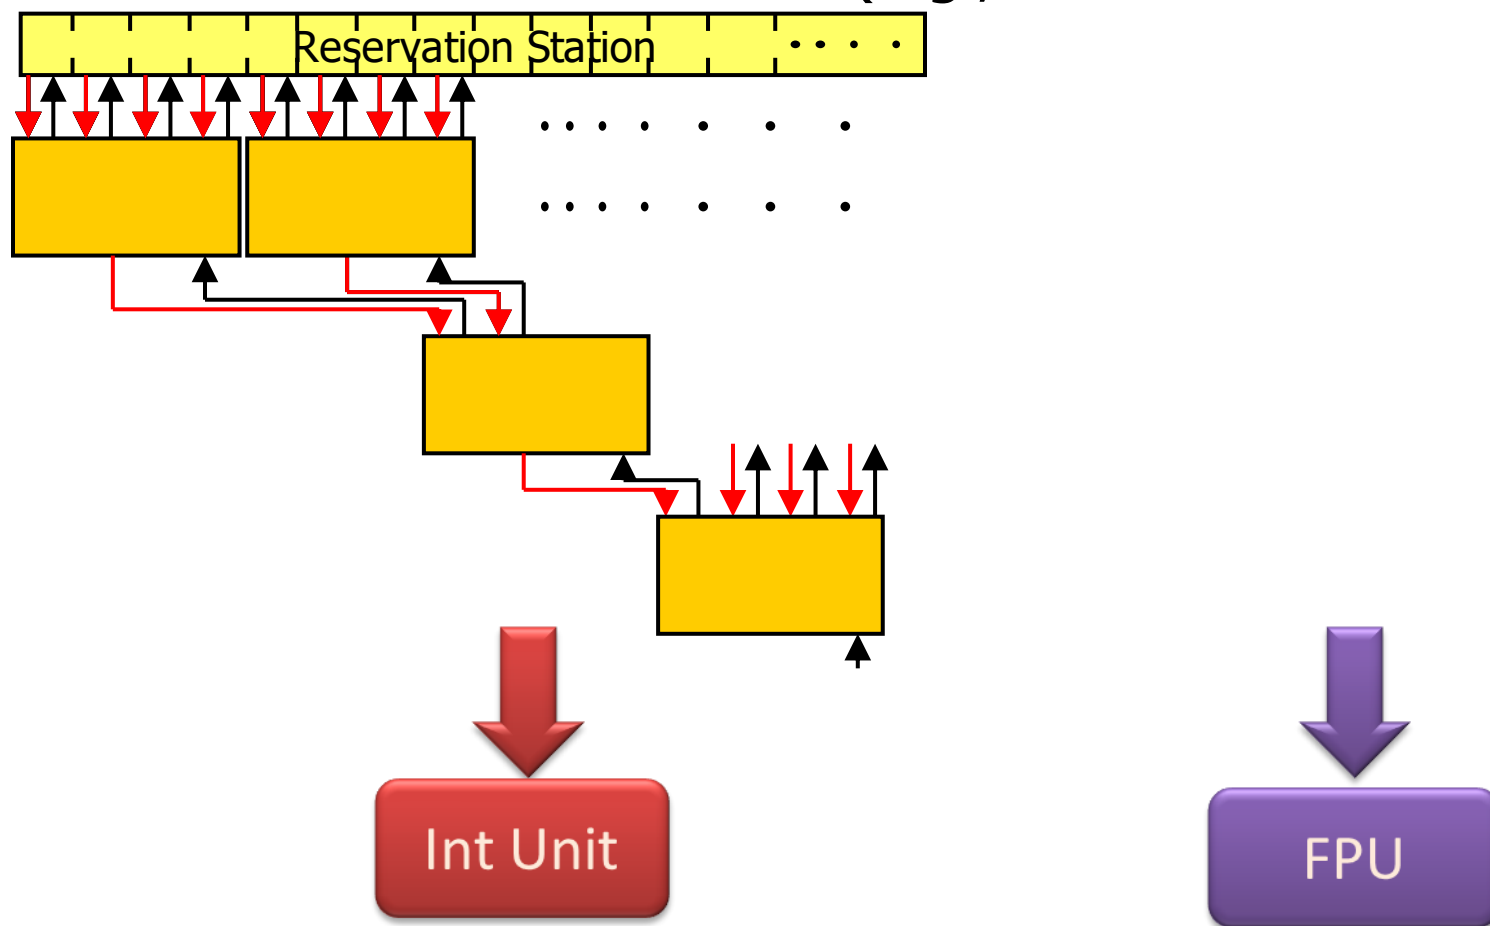
# Simple Select Logic Implementation



# Issues to Distinctive Functional Units

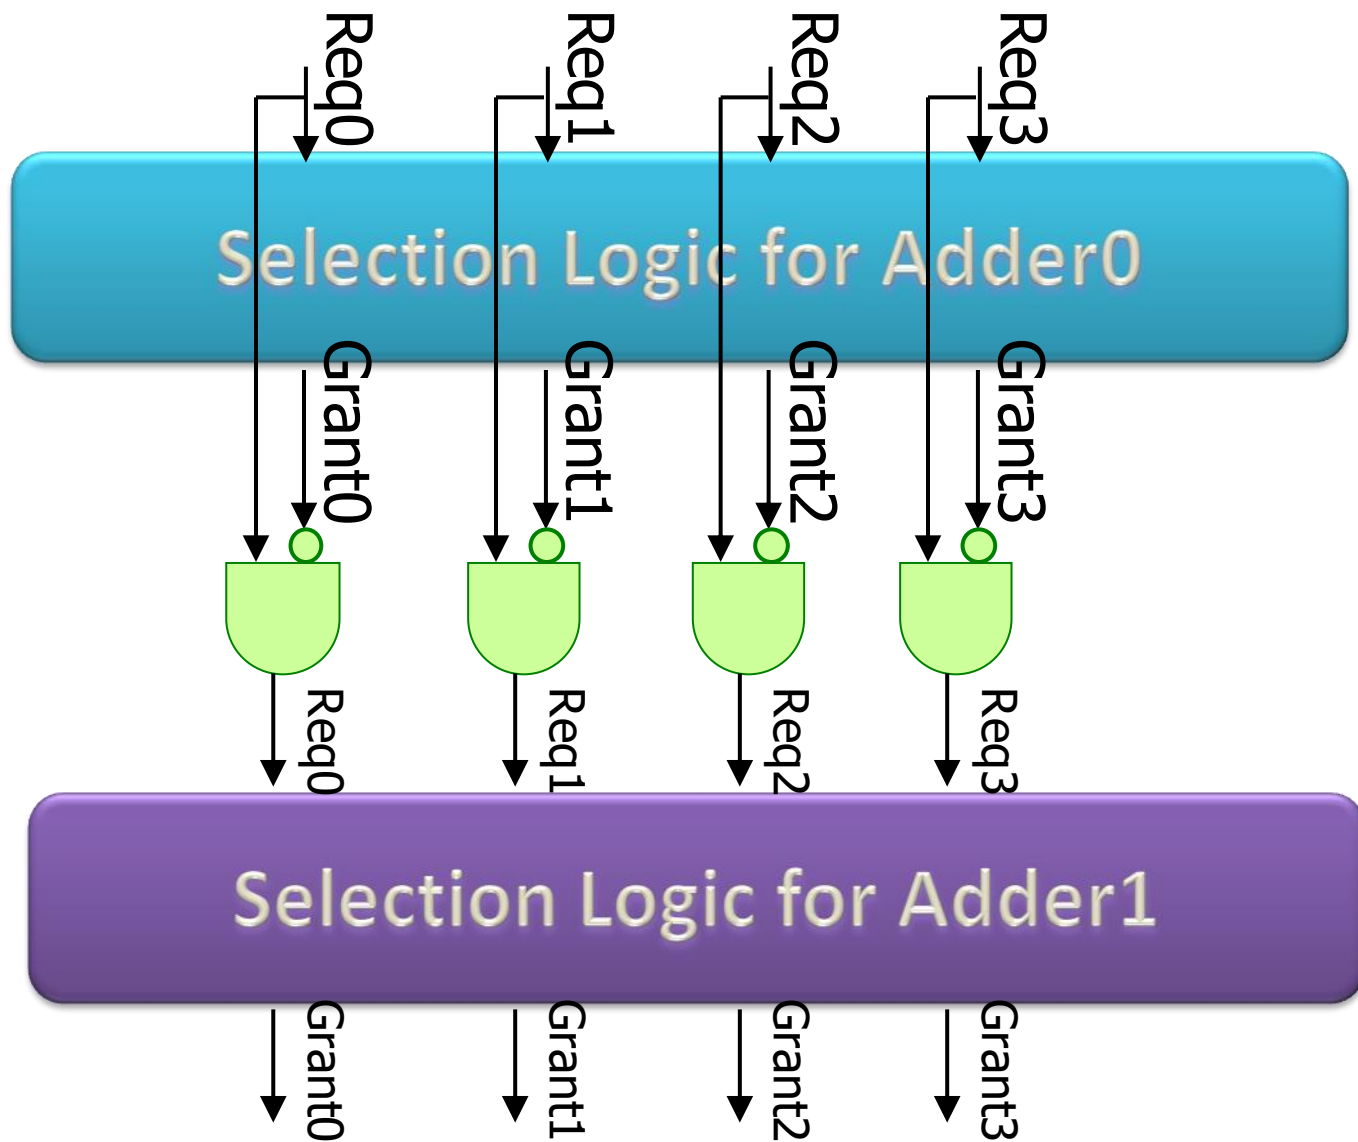



Distributed Instruction Windows (e.g., MIPS R1000 or Alpha 21264)




Faster to have separate instruction schedulers for different instruction

# Dual Issues to Multiple Units








## Bypass Delay

---

- The number of bypass paths equals  $2 \times IW^2 \times S$   
(S is the number of pipeline stages)
- Wire length  $\sim IW$ , hence, delay  $\sim IW^2$
- The layout and pipeline depth (capacitive load) also matter


# Other structures

---




- Registers
- Caches

# Limitations of ILP




- Assume an ideal processor model
  - Register renaming
    - Infinite number of virtual registers
    - All WAW and WAR hazards can be removed
  - Branch prediction
    - Perfect branch prediction
  - Jump prediction
    - All jumps are perfectly predicted
  - Memory address alias analysis
    - All memory addresses are known exactly
  - Perfect caches
    - All memory accesses take 1 clock cycle



# Impact of Window Size



Benchmarks





# Impact of Branch Prediction



2K window  
Max. 64 instruction issues per  
cycle

# Impact of Finite Registers




2K window

Max. 64 instruction issues per cycle

8K entry tournament predictors


2K jump and return predictors

# Limits of Multiple-Issue Processors: Revisited



- Doubling the issue rate above the current 3-6 issue, i.e. 6-12 issue requires
  - Issue 3-4 data memory accesses per cycle
  - Resolve two or three branches per cycle
  - Rename and access more than 20 registers per cycle
  - Fetch 12-24 instructions per cycle
  - ➔ The complexity of implementing these capabilities would sacrifice the maximum clock rate
- Another issue is power!
  - Modern microprocessors are primarily power limited.
    - *Static power grows proportionally to the transistor count*
    - *Dynamic power is proportional to the product of the number of switching transistors and the switching rate*
  - Microprocessors trying to achieve both a high IPC and a

# Limits of Multiple-Issue Processors

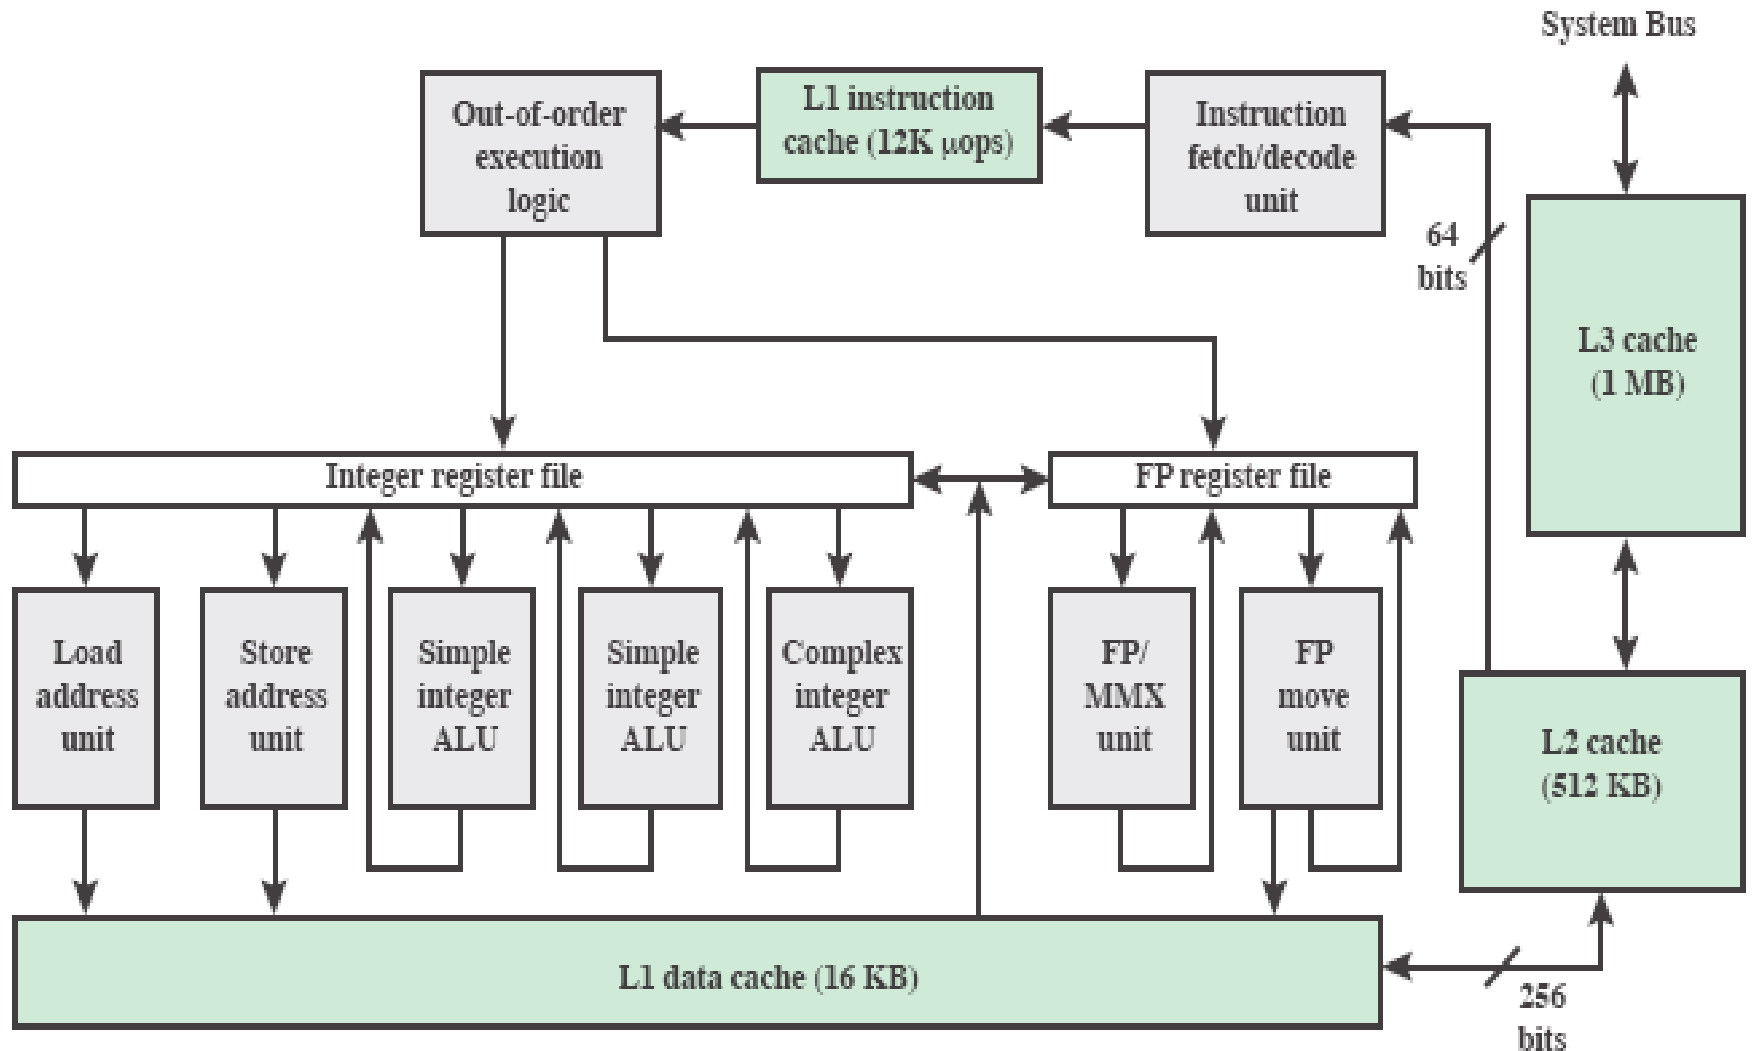



- Energy inefficiency of multiple-issue processors
  - Multiple instruction issue incurs overhead in logic that grows faster than the issue rate increase
    - Dependence checking, register renaming, wakeup and select, etc.
    - Without voltage reduction, higher IPC will lead to lower performance/watt!
  - Growing gap between peak issue rate and sustained performance
    - The number of transistor switching is proportional to the peak issue rate
    - The performance is proportional to the sustained rate
    - Thus, growing gap translates to increasing energy per unit performance!
- For example, speculation is inherently inefficient
  - It can never be perfect, thus there is inherently waste in executing computations before we know that whether the path is taken or not

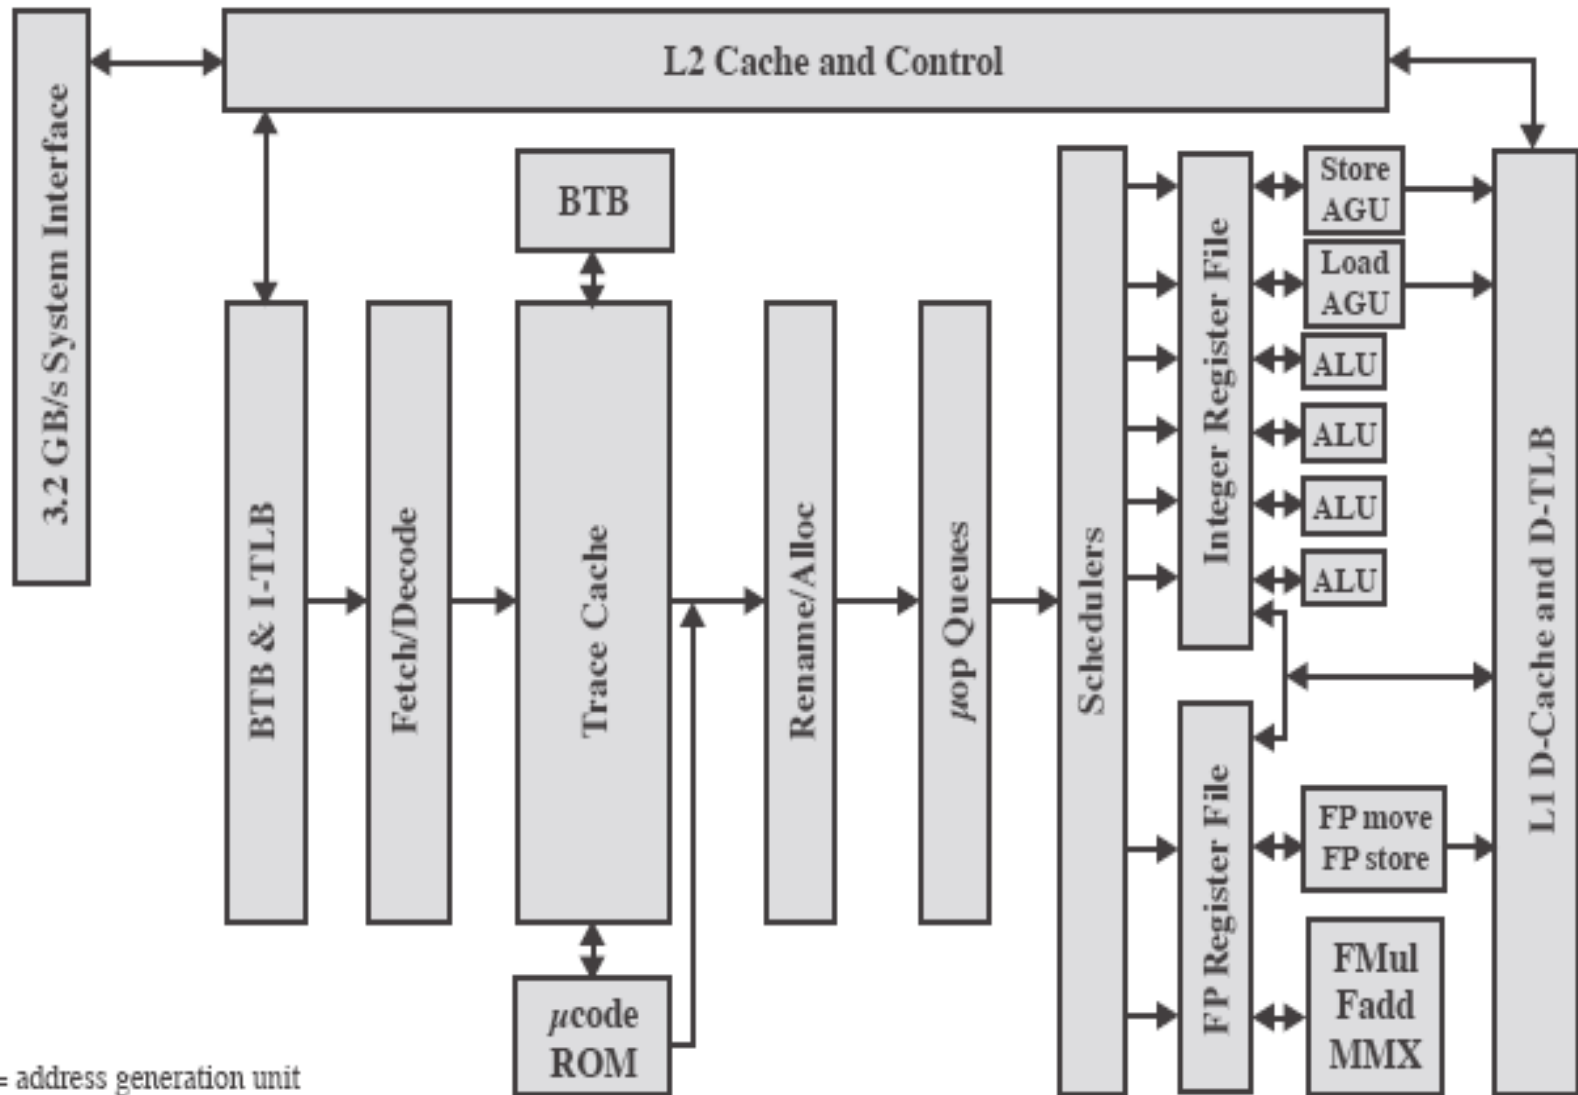


# Example: Pentium 4

## A Superscalar CISC Machine



# Pentium 4 alternate view





AGU = address generation unit

BTB = branch target buffer

D-TLB= data translation lookaside buffer

I-TLB = instruction translation lookaside buffer

# Pentium 4 pipeline



TC Next IP = trace cache next instruction pointer

TC Fetch = trace cache fetch

Alloc = allocate

Rename = register renaming

Que = micro-op queuing

Sch = micro-op scheduling

Disp = Dispatch

RF = register file

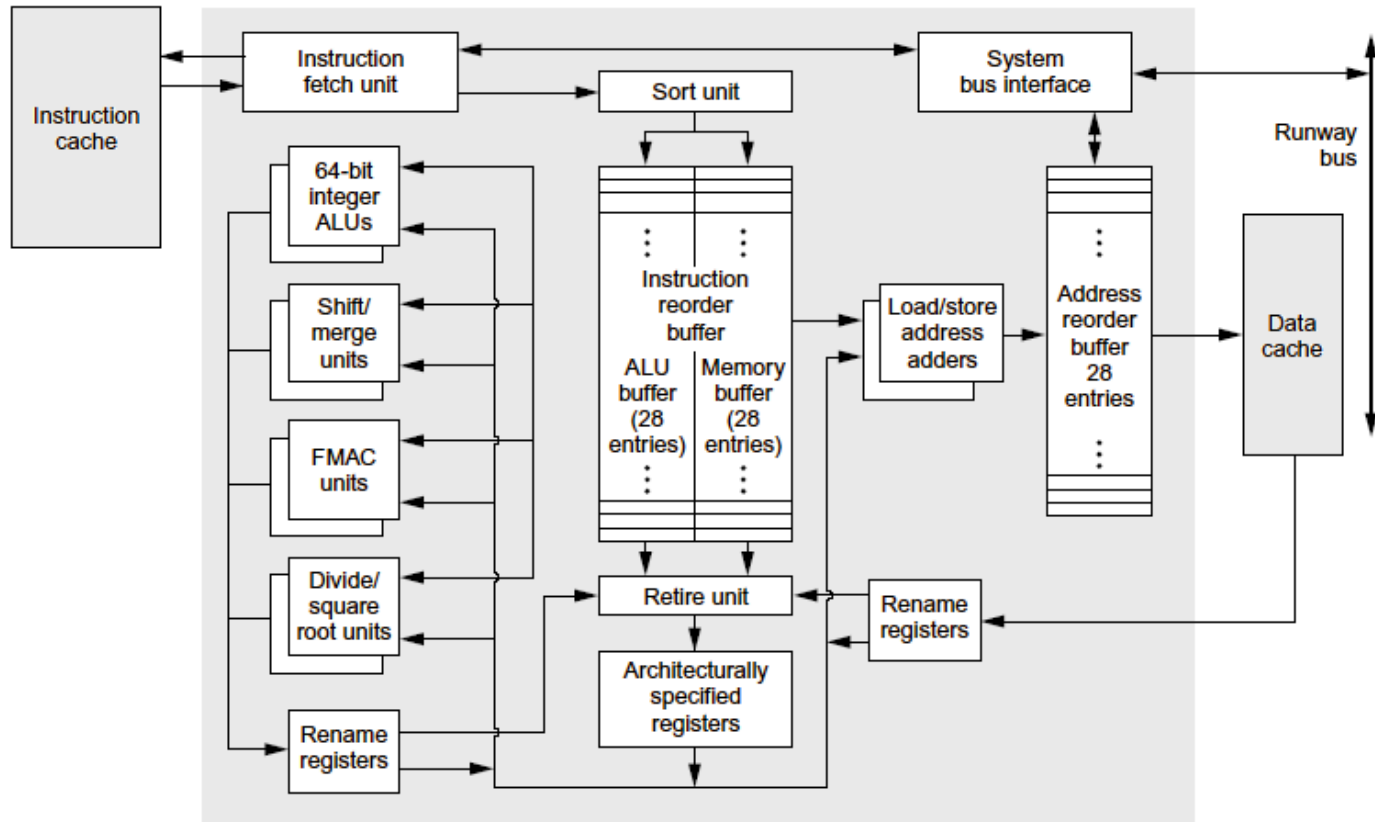

Ex = execute

Flgs = flags



Br Ck = branch check

## 20 stages !



# HP PA-RISC 8000



# Processor Performance with Time



# Price being paid



# MIPS-X (1988)

#### 4.58. hsc - Halt and Spontaneously Combust

[illegible]

## Assembler

### hsc

## Operation

$$\text{Reg}(31) \Leftarrow \text{PC}$$

The processor stops fetching instructions and self destructs.

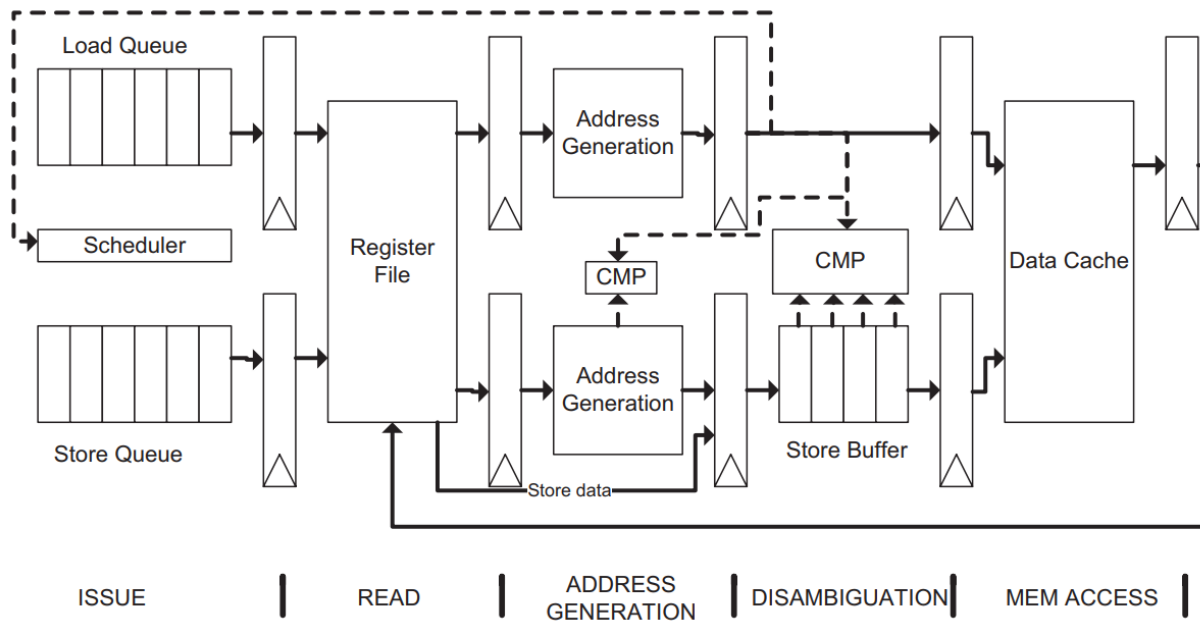
Note that the contents of `Reg(31)` are actually lost.

### Description

This is executed by the processor when a protection violation is detected. It is a privileged instruction available only on the *-NSA* versions of the processor.

# Load Ordering and Store Ordering on AMD K6

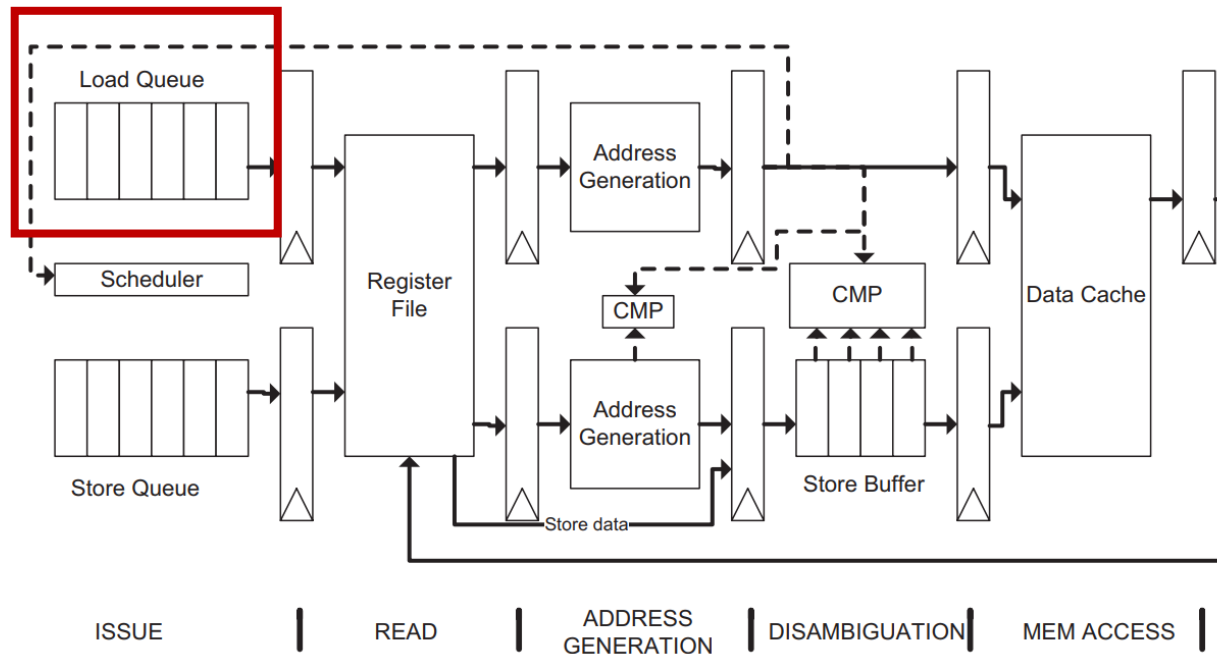
- two separate pipelines for load and store operations w/ some level of communication
- instructions flow in strict order inside each pipeline





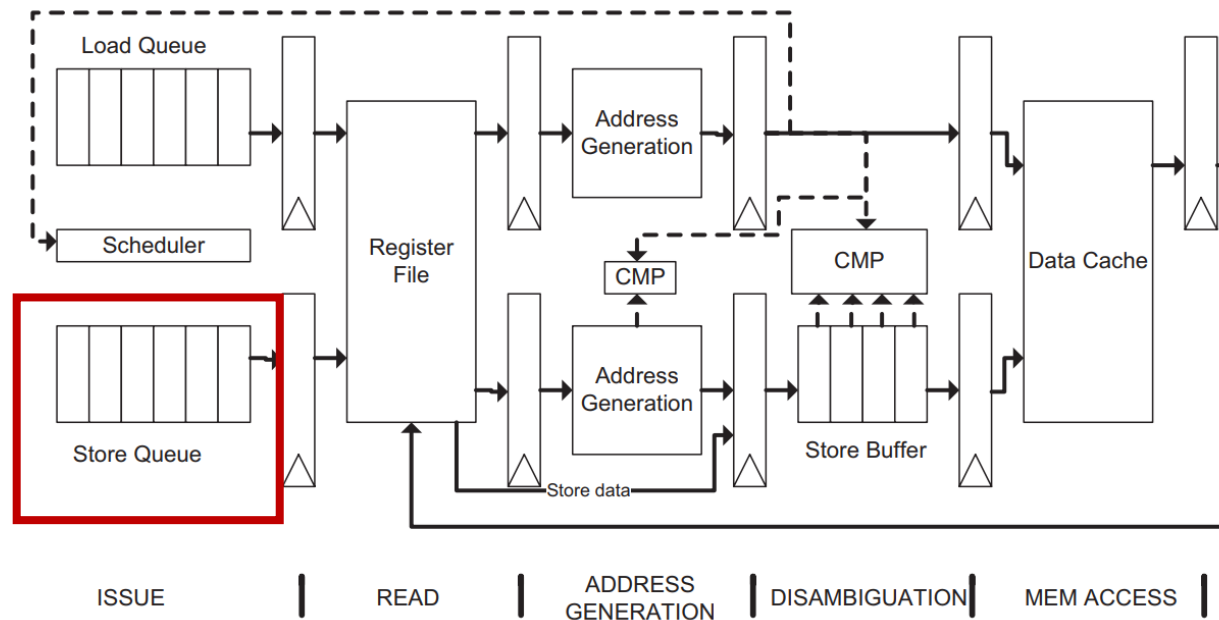
# Load Ordering and Store Ordering on AMD K6

- Load queue: stores LD operations in order
  - LD operations are inserted in this queue after renaming and reside there until they become oldest on queue and their source operands are ready.



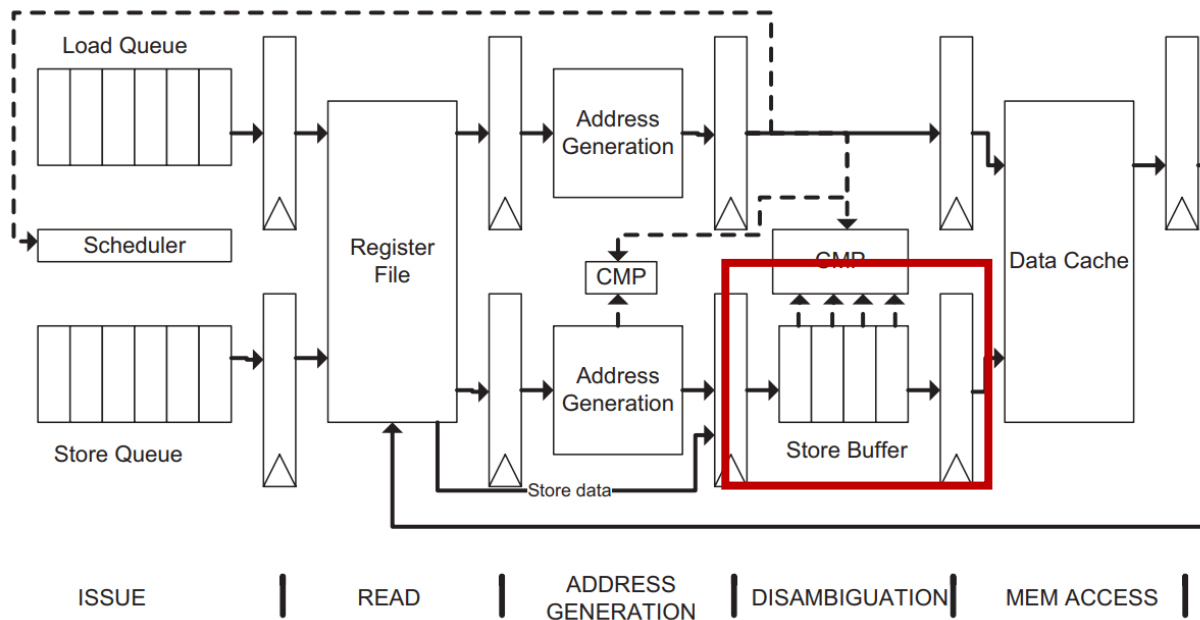
# Load Ordering and Store Ordering on AMD K6

- Store queue: stores the ST operations in order
  - ST operations reside here since they have been renamed until they become oldest instruction on queue; source operands they need to compute the address are available



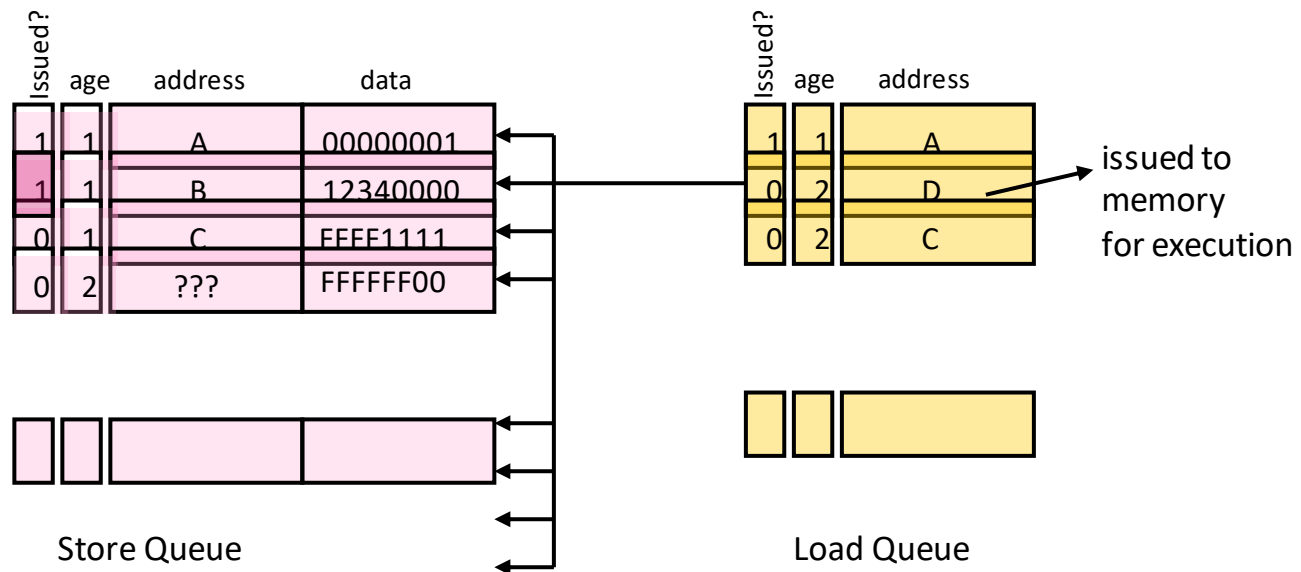
# Load Ordering and Store Ordering on AMD K6

- Store buffer: keeps the ST operations in order until they become oldest in-flight instruction in the processor, and then they proceed to update the memory



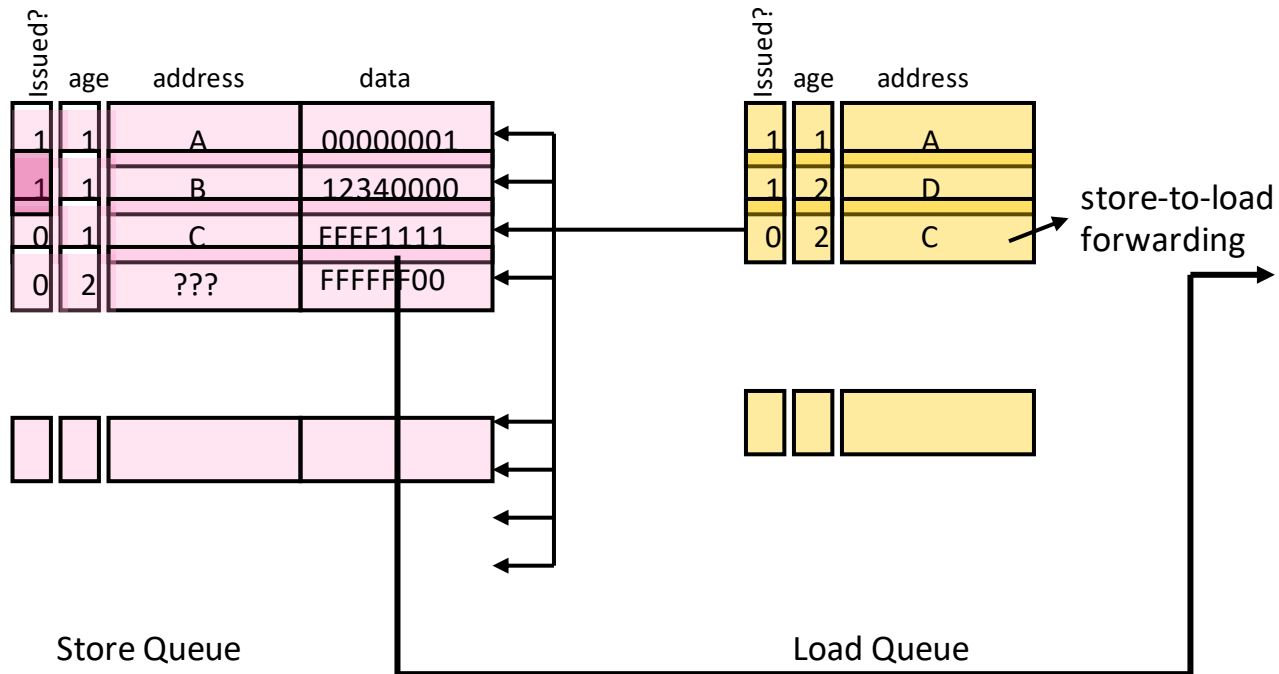
# LSQ

- load will compare its memory address with addresses of stores in store buffer **that are older than it**
- each load checks against older stores
  - associative search → performance scalability issue



# LSQ

- each load checks against older stores
  - associative search → performance scalability issue



# Amdahl's Law


- law of diminishing returns

- ⊙ make the common case faster


- ⊙  $\text{speedup} = \text{Perf}_{\text{new}} / \text{Perf}_{\text{old}} = T_{\text{old}} / T_{\text{new}} = \frac{1}{(1 - f) + \frac{f}{P}}$

- Example

- ⊙ Suppose floating point instructions occupy 10% of a program, they are improved to run 2X faster, what's the speedup?

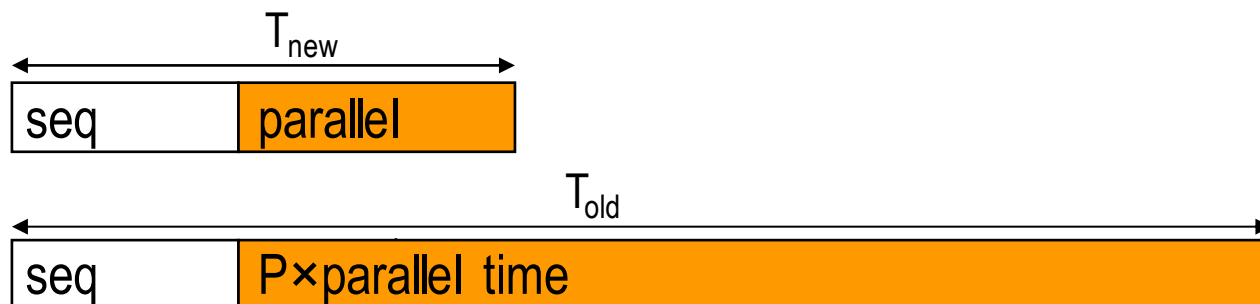


# Parallelism vs. Speedup



# Gustafson's Law

- Amdahl's Law killed massive parallel processing (MPP)
- Gustafson came to rescue
  - ✓ more workloads




- ✓ assume:  $\text{seq} + \text{parallel} = 1$  ( $T_{\text{new}}$ )
- ✓  $\text{speedup} = \text{seq} + p \times (1 - \text{seq})$  where  $p$  = parallel factor
- ✓ if seq diminishes w/ increased problem size,  $\text{speedup} \rightarrow p$



# New Breed of Metrics

- performance/Watt
  - ⊙ performance achievable at the same cooling capacity
- performance/Joule (energy)
  - ⊙ achievable performance at the lifetime of the same energy source (i.e., battery = energy)
  - ⊙ equivalent to reciprocal of energy-delay product (the product of E and D)


# (Dynamic) Power Dissipation



$$\text{energy/transition} = C_L \times V_{DD}^2 \times P_{0/1 \rightarrow 1/0}$$

$$\text{power} = C_L \times V_{DD}^2 \times f$$

# (Short Circuit) Power Dissipation



$$\text{Energy/transition} = t_{sc} \times V_{DD} \times I_{peak} \times P_{0/1 \rightarrow 1/0}$$

$$\text{Power} = t_{sc} \times V_{DD} \times I_{peak} \times f$$

# CMOS Energy & Power Equations

$$E = C_L V_{DD}^2 P_{0 \rightarrow 1} + t_{sc} V_{DD} I_{peak} P_{0 \rightarrow 1} + V_{DD} I_{leakage}$$

$$f_{0 \rightarrow 1} = P_{0 \rightarrow 1} * f_{clock}$$

$$P = C_L V_{DD}^2 f_{0 \rightarrow 1} + t_{sc} V_{DD} I_{peak} f_{0 \rightarrow 1} + V_{DD} I_{leakage}$$

dynamic power  
( $\approx 40 - 70\%$  today and  
decreasing relatively)

short-circuit power  
( $\approx 10\%$  today and  
decreasing absolutely)

leakage power  
( $\approx 20 - 50\%$  today  
and increasing)



# Power and Energy Metrics

- power in Watts
  - ⊙ **Watt = Energy / Time**
  - ⊙ poses constraints:
    - i.e., processor can only work fast enough to max out the power delivery or cooling solution
  - ⊙ Peak power
    - determines power ground wiring designs
    - sets packaging/cooling limits
    - impacts signal noise margin and reliability analysis

# Power and Energy Metrics

- Energy in Joules
  - ⊙ ultimate metric, i.e., the true “cost” of performing a fixed task
  - ⊙ energy = power × delay
    - Joules = Watts × seconds
    - Lower energy means less power to perform a computation at the same frequency
- Example:
  - ⊙ if processor A consumes 1.2× power of processor B, but finishes the task in 30% less time, which processor is more energy efficient and by how much?


# Power versus Energy



# More Energy Metrics

EDP = Energy Delay Product =  $E \times D$


EDDP = Energy Delay<sup>2</sup> Product =  $ED^2$  (more emphasis on performance)






# Research on Power/Energy

- What happens when a CPU heatsink is removed
  - ◎ Check this [YouTube link](#)



Google Data Center Cooling System



Microsoft Experimental Datacenter