


ECE411: Computer Organization and Design

Lecture 16: Multi-core and Multi-threading


Rakesh Kumar




Uniprocessor (single-core) Architecture



Intel 386



Intel 486




Intel Pentium

Multi-core Architecture

- Performance demand from applications
- Faster computation



Intel Nehalem – First true multi-core processor

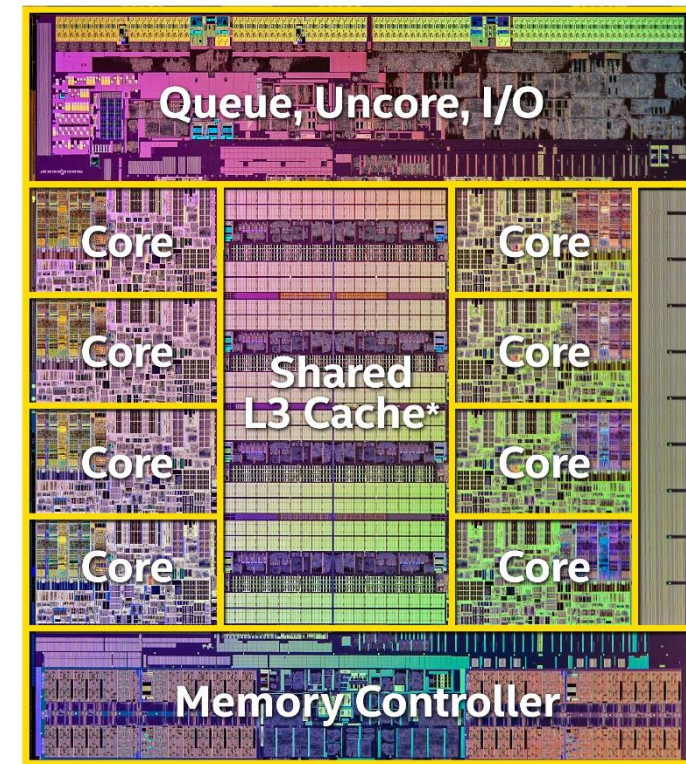


Multi-core Architecture

- Performance demand from applications
- Faster computation



New 8-Core Intel® Core™ i7 Processor Extreme Edition




Intel® Core™ i7-5960X Processor Extreme Edition
Transistor count: 2.6 Billion
Die size: 17.6mm x 20.2mm



* 20MB of cache is shared across all 8 cores



Why Multi-core Architecture?

- Difficult to make single-core clock frequencies even higher
 - ⦿ hard to extract more ILP b/c larger instruction windows are very expensive, slow, power-hungry circuits
- Deeply pipelined circuits:
 - ⦿ Heat problems
 - ⦿ Difficult design and verification
 - ⦿ Large design teams necessary
 - ⦿ Server farms need expensive air-conditioning




<https://dl.acm.org/doi/pdf/10.5555/774861.774897>

Power vs. Frequency




Power and Thermal Limit



How to reduce the power consumption

- Multicore
 - ⊙ One core with frequency of 2 GHz
 - ⊙ Two cores with 1GHz frequency (each)
 - Same throughput
 - Two 1GHz cores require less power/energy
- New Challenges – Performance
 - ⊙ How to utilize the cores
 - ⊙ It is difficult to find parallelism in programs to keep all these cores busy

Power and Thermal Limit




Utilization Wall: Dark Silicon's Effect on Multicore Scaling


Spectrum of tradeoffs
between # of cores and
frequency

Example:
65 nm → 32 nm ($S = 2$)

4 cores @ 1.8 GHz




65 nm



2x4 cores @ 1.8 GHz
(8 cores dark, 8 dim)

(Industry's Choice)



4 cores @ 2x1.8 GHz
(12 cores dark)

75% dark after 2 generations;
93% dark after 4 generations

32 nm

Figure 1: Multicore scaling leads to large amounts of Dark Silicon. (From[7].)

Why Multi-core Architecture?

- Parallel processors now play a major role
 - ⊙ Logical way to improve performance
 - Connect multiple microprocessors
 - ⊙ Not much left with ILP exploitation
 - ⊙ Data center servers have parallelism
- Multi-core processor architectures will become increasingly attractive
 - ⊙ Due to slowdown in advances of uniprocessors

Optional Reading: <https://arxiv.org/pdf/1911.11313v1.pdf>

https://education.dellemc.com/content/dam/dell-emc/documents/en-us/2019KS_Yellin-Saving_The_Future_of_Moores_Law.pdf

ILP versus TLP

- instruction-level parallelism (ILP)


- ⊙ parallelism at the machine-instruction level
- ⊙ the processor can re-order, pipeline instructions, split them into microinstructions, do aggressive branch prediction, etc.
- ⊙ instruction-level parallelism enabled rapid increases in processor speeds over the last 15 years

- thread-level parallelism (TLP)

- ⊙ parallelism on a more coarser scale
- ⊙ server can serve each client in a separate thread (web server, database server)
- ⊙ a computer game can do AI, graphics, and physics in three separate threads
- ⊙ single-core superscalar processors cannot fully exploit TLP
- ⊙ multi-core architectures are the next step in processor evolution: explicitly exploiting TLP

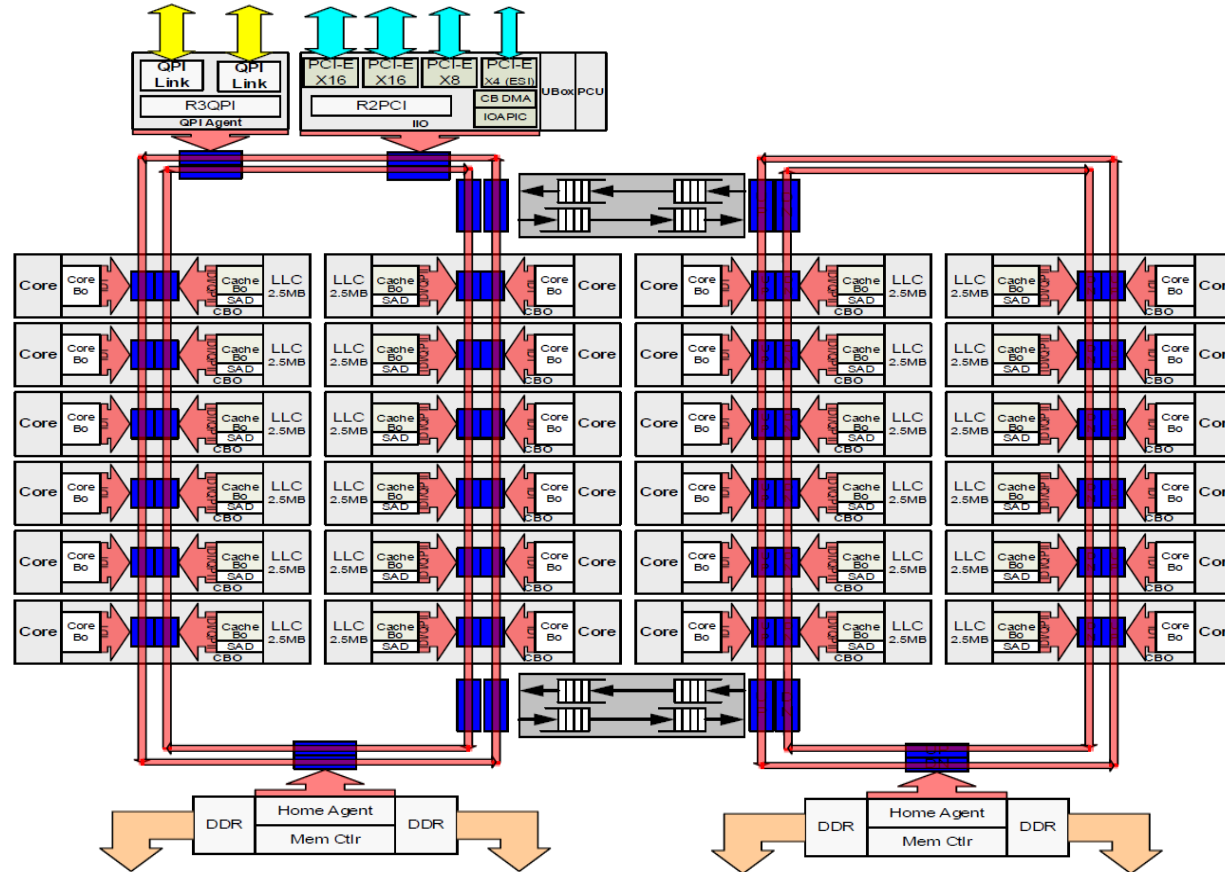
Multi-core CPU Chip

- Multiple cores fit on a single processor socket
 - ⦿ also called CMP (Chip Multi-Processor)



Multi-core Architectures

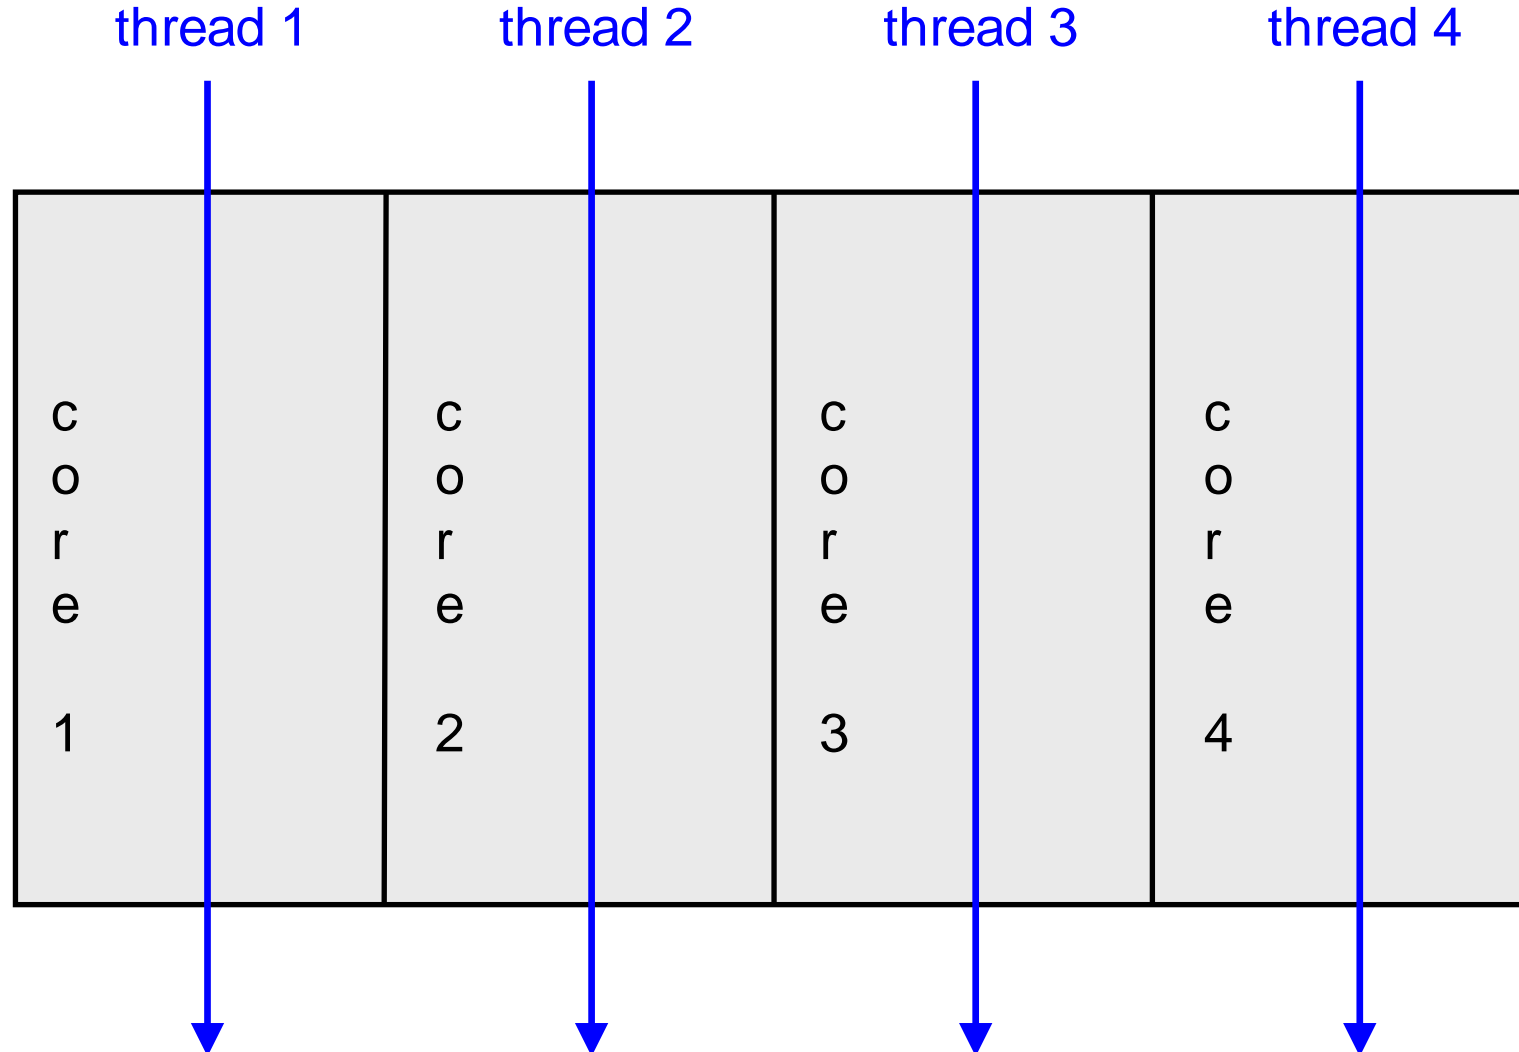
Intel® Xeon® Processor E5 v4 Product Family HCC



multi-core CPU chip


Cores Run in Parallel

- Exploiting thread-level parallelism



Cores Run in Parallel

- Threads can be time-sliced (just like on a uniprocessor) in each core



Interaction with OS

- OS perceives each core as a separate processor
- OS scheduler maps threads/processes to different cores
- Most major OS support multi-core today
 - ⦿ Windows, Linux, Mac OS X, ...

General Context: Multiprocessors

- Multiprocessor is any computer with several processors
- SIMD
 - ⦿ single instruction, multiple data
 - modern graphics cards
- MIMD
 - ⦿ multiple instructions, multiple data



Lemieux cluster,
Pittsburgh
supercomputing
center

Multiprocessor

- 2—4 multi-core processors on a motherboard
 - ⦿ connected by special short distance interconnect
 - Intel QuickPath



Multiprocessor Memory Types

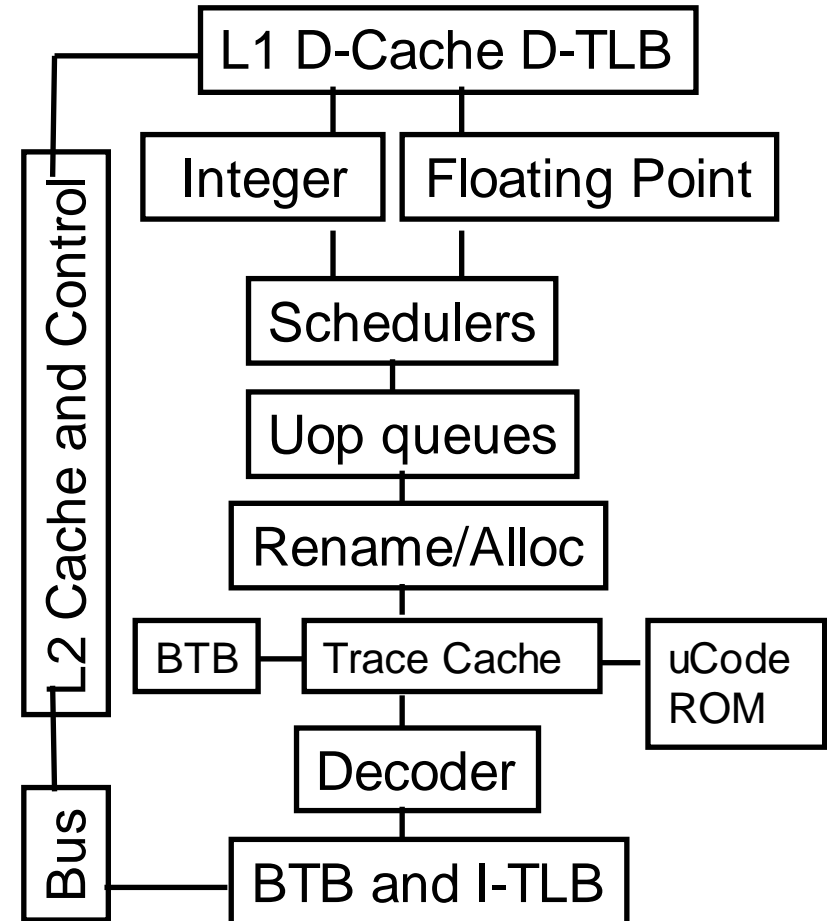
- Shared memory
 - ⦿ one (large) common shared memory for all processors
- Distributed memory
 - ⦿ each processor has its own (small) local memory, and its content is not replicated anywhere else

Multiprocessor Vs. Multi-core Processor

- Multi-core processor
 - ⦿ a special kind of a multiprocessor with all processors on the same chip
- Multi-core processors are MIMD
 - ⦿ Different cores execute different threads (multiple instructions), operating on different parts of memory (multiple data)
- Multi-core is a shared memory multiprocessor
 - ⦿ All cores share the same memory

Simultaneous Multi-Threading (SMT)

- Problem addressed
 - ⦿ processor pipeline can get stalled:
 - Waiting for the result of a long floating-point (or integer) operation
 - Waiting for data to arrive from memory
- Other execution units wait unused

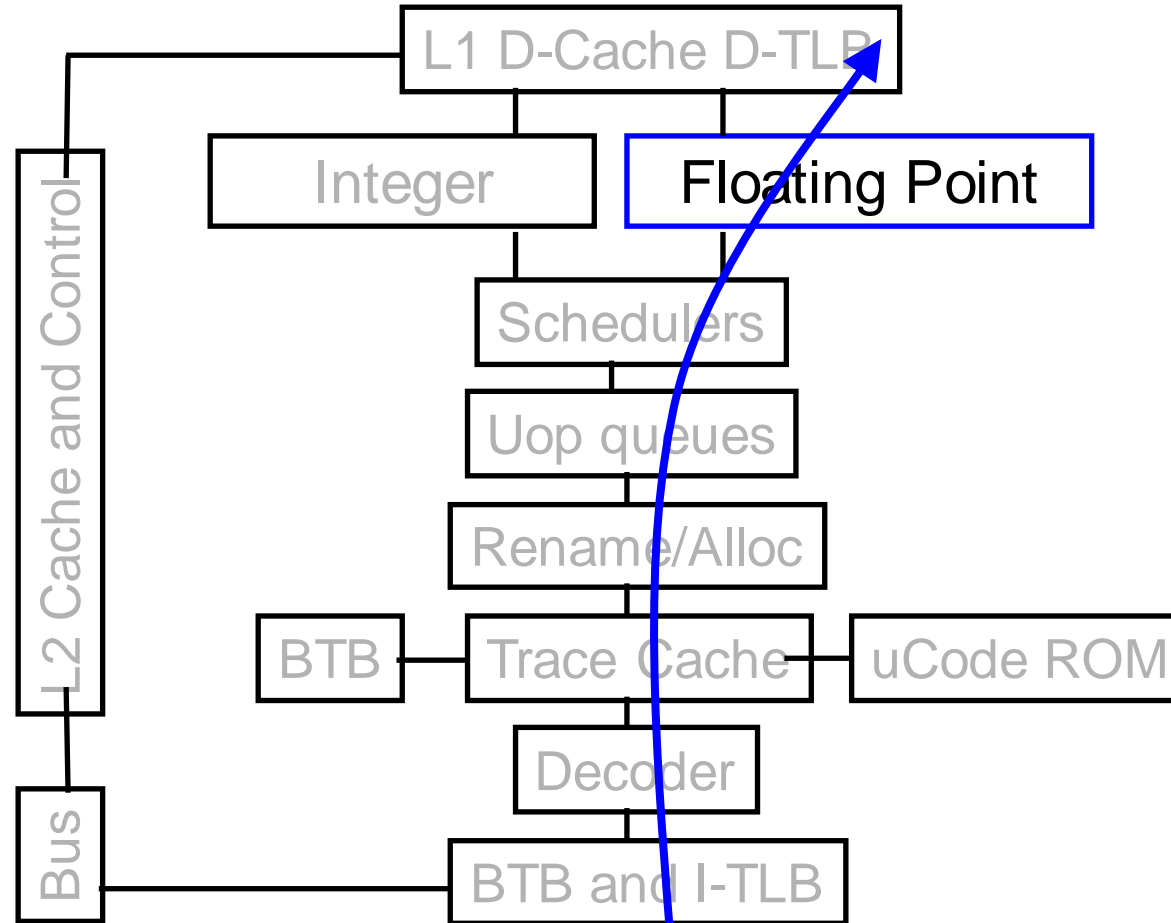


Simultaneous Multi-Threading (SMT)

- Permits multiple independent threads to execute SIMULTANEOUSLY on the SAME core
- Weaving together multiple “threads” on the same core
- Example:
 - ⦿ if one thread is waiting for a floating point operation to complete, another thread can use the integer units
- Another name: **hyper-threading**

Simultaneous Multi-Threading (SMT)

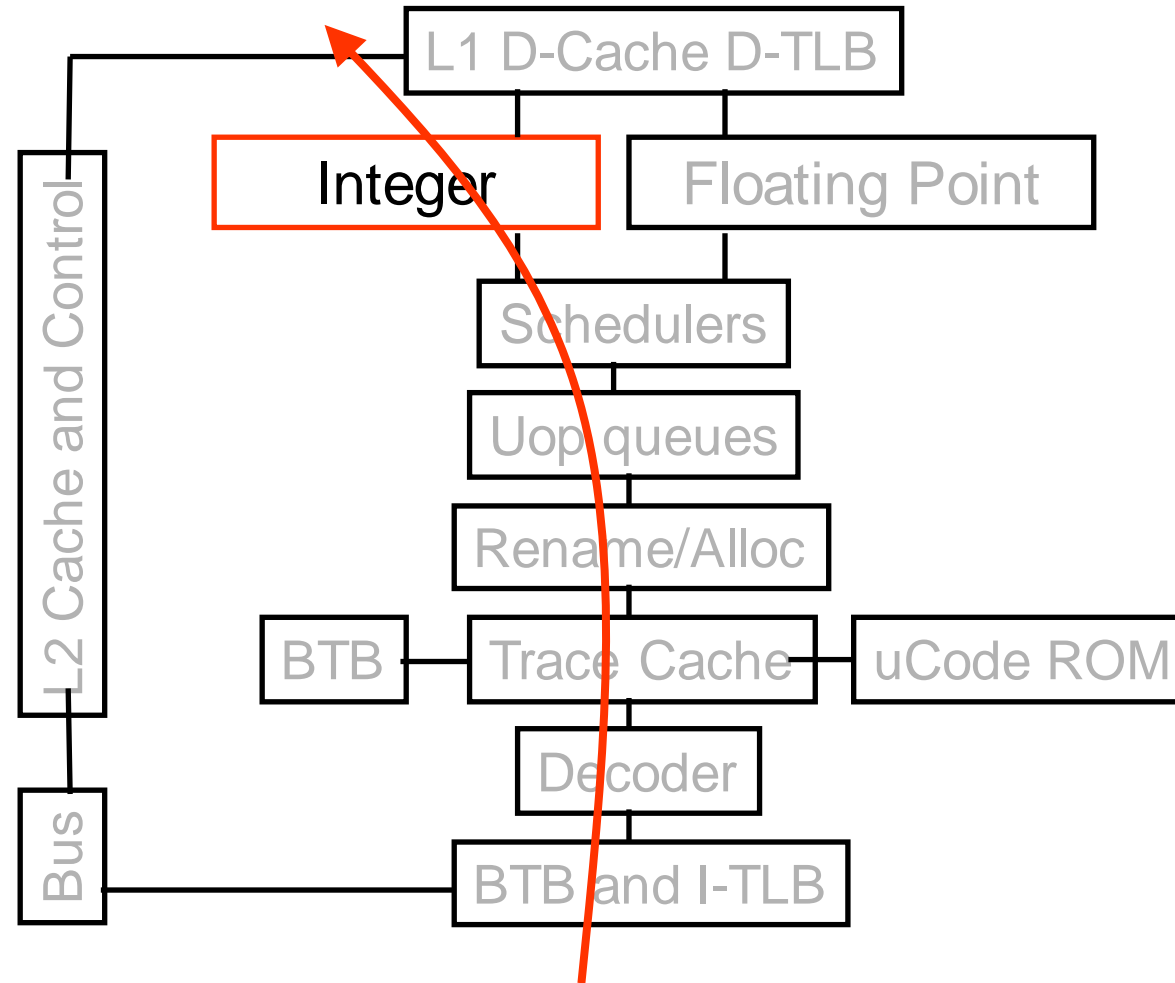
- w/o SMT, only a single thread can run at any given time



Thread 1: floating point

Simultaneous Multi-Threading (SMT)

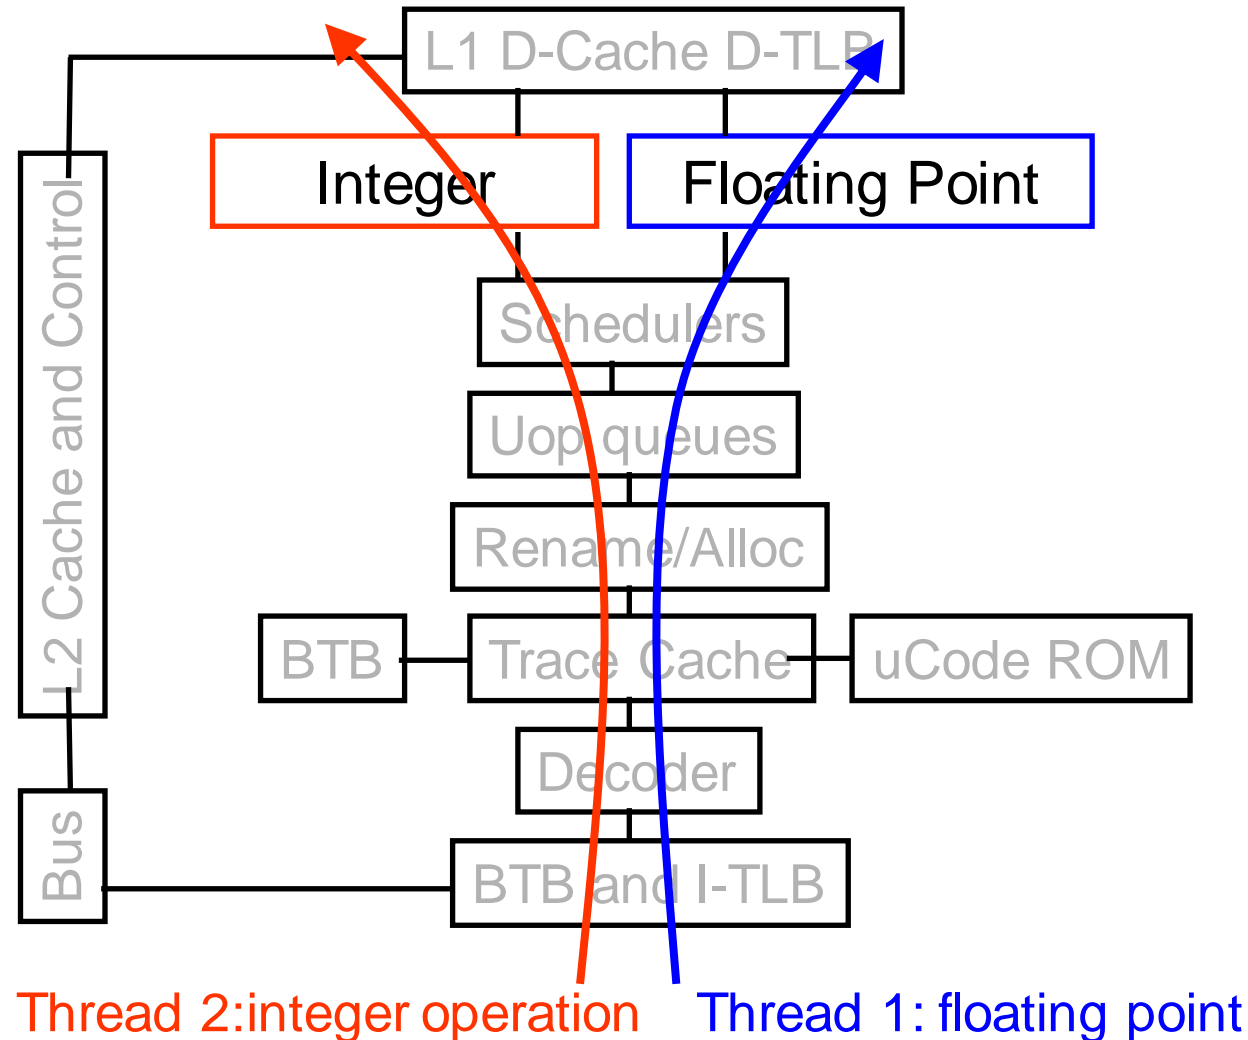
- w/o SMT, only a single thread can run at any given time



Thread 2: integer operation

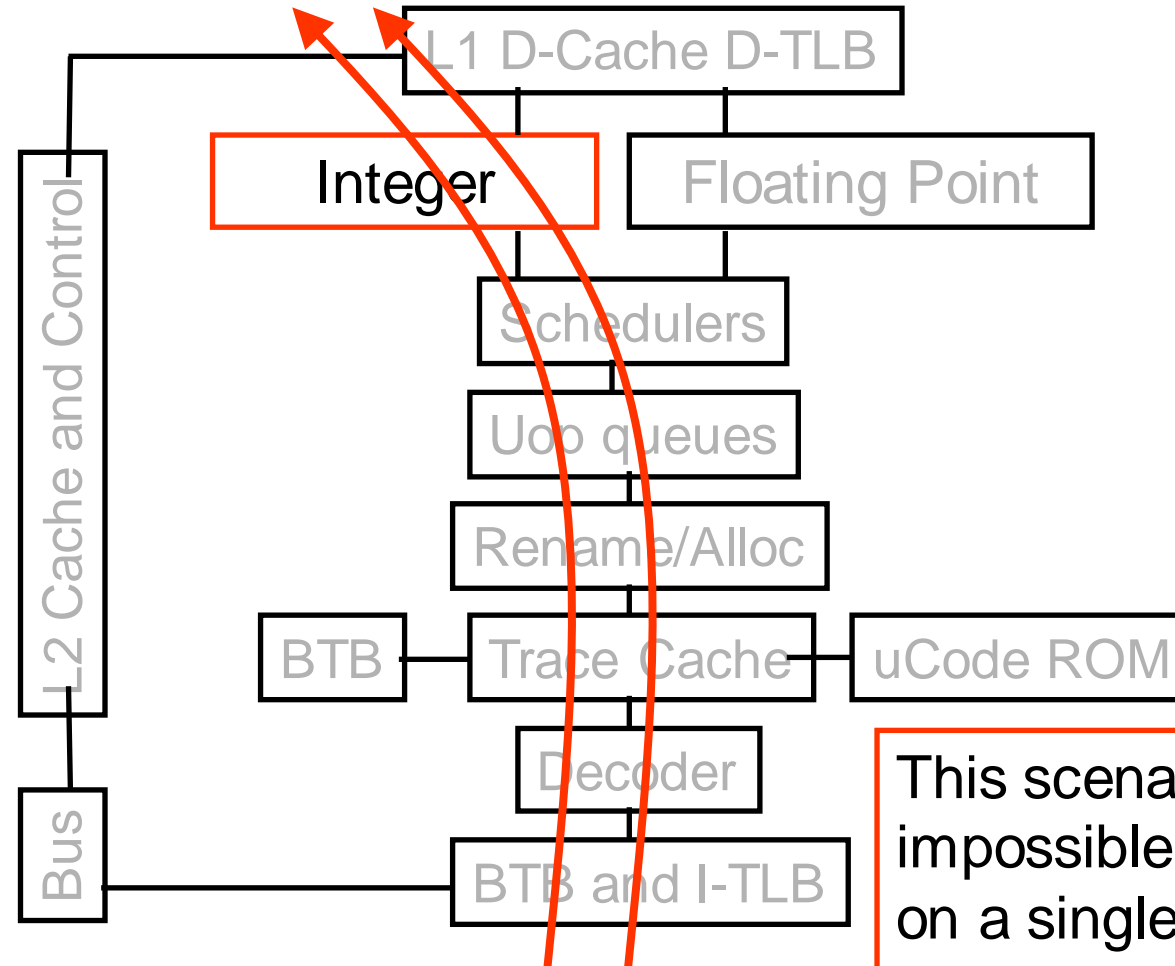
Simultaneous Multi-Threading (SMT)

- SMT processor: both threads can run concurrently



Simultaneous Multi-Threading (SMT)

- But SMT can't simultaneously use the same functional unit



Thread 1 Thread 2
IMPOSSIBLE

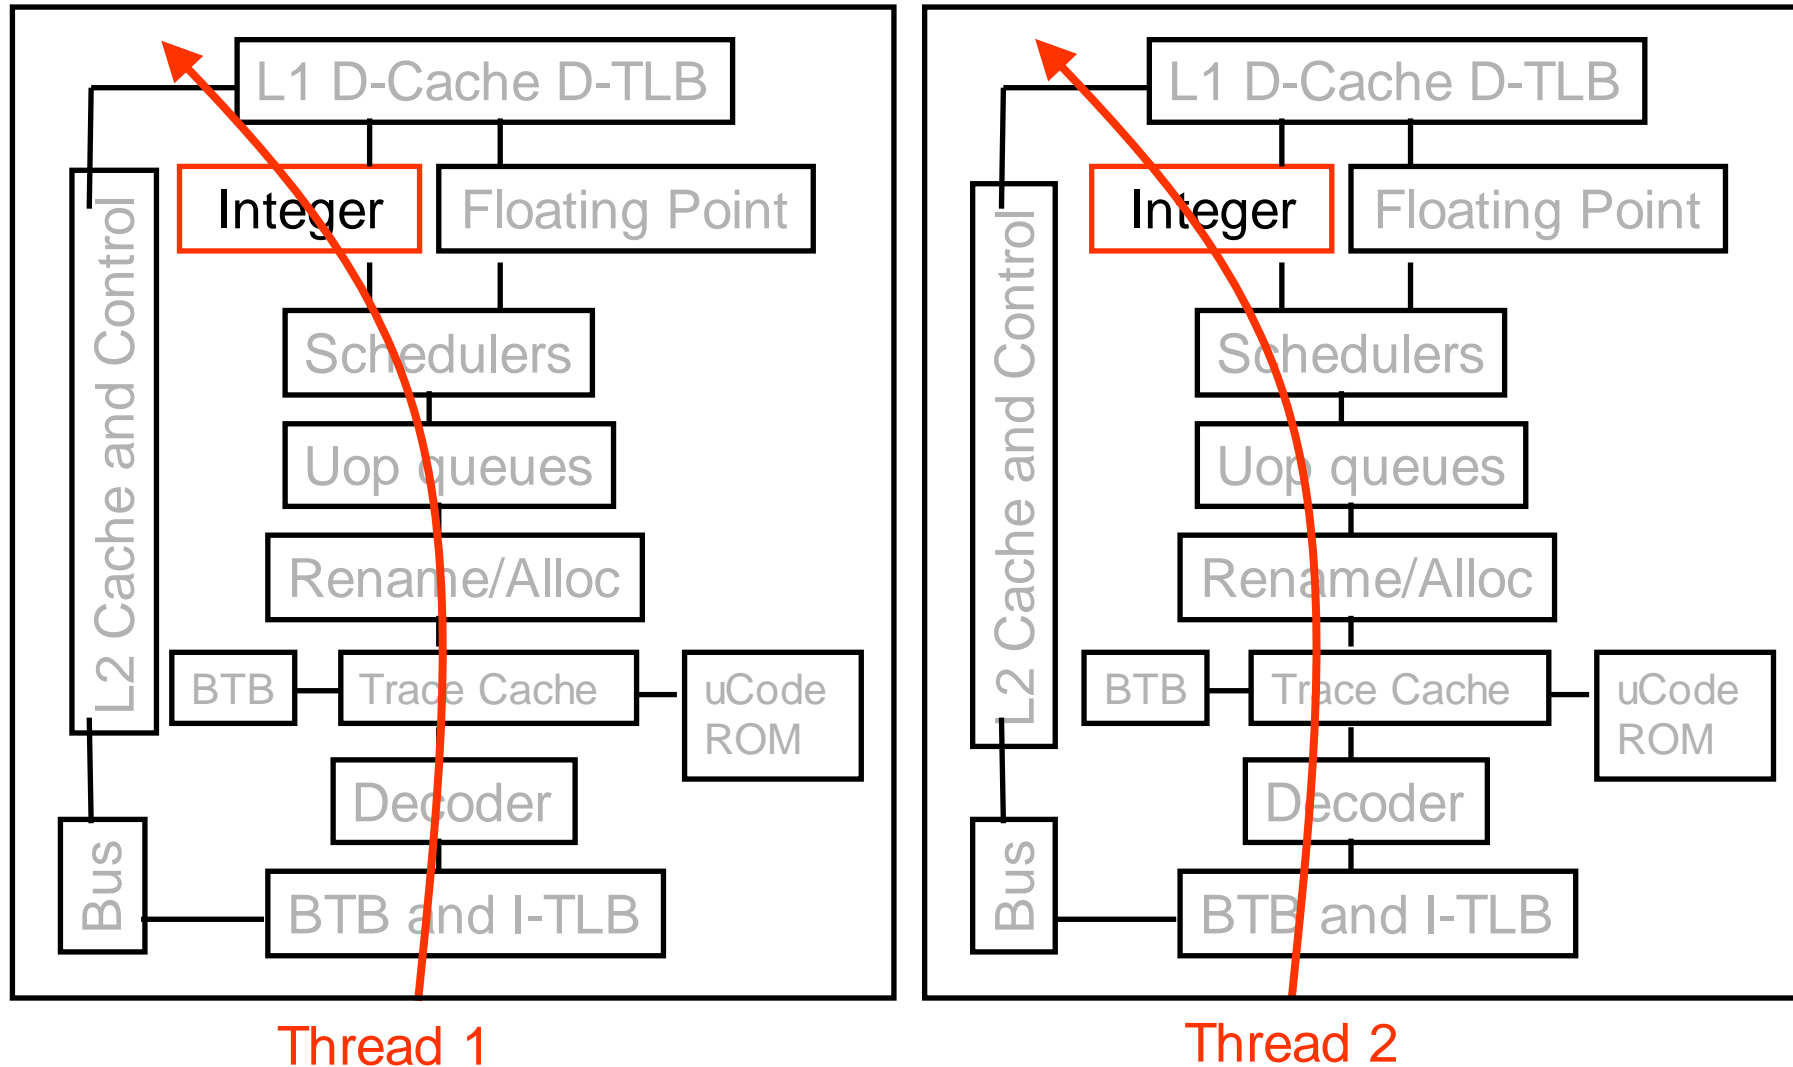
This scenario is impossible with SMT on a single core (assuming a single integer unit)

Simultaneous Multi-Threading (SMT)

- SMT not a “true” parallel processor
 - ⊙ Enables better threading (e.g., up to 30%)
 - ⊙ OS and applications perceive each simultaneous thread as a separate “virtual processor”
 - ⊙ The chip has only a single copy of each resource
 - ⊙ Compared to multi-core, each core has its own copy of resources

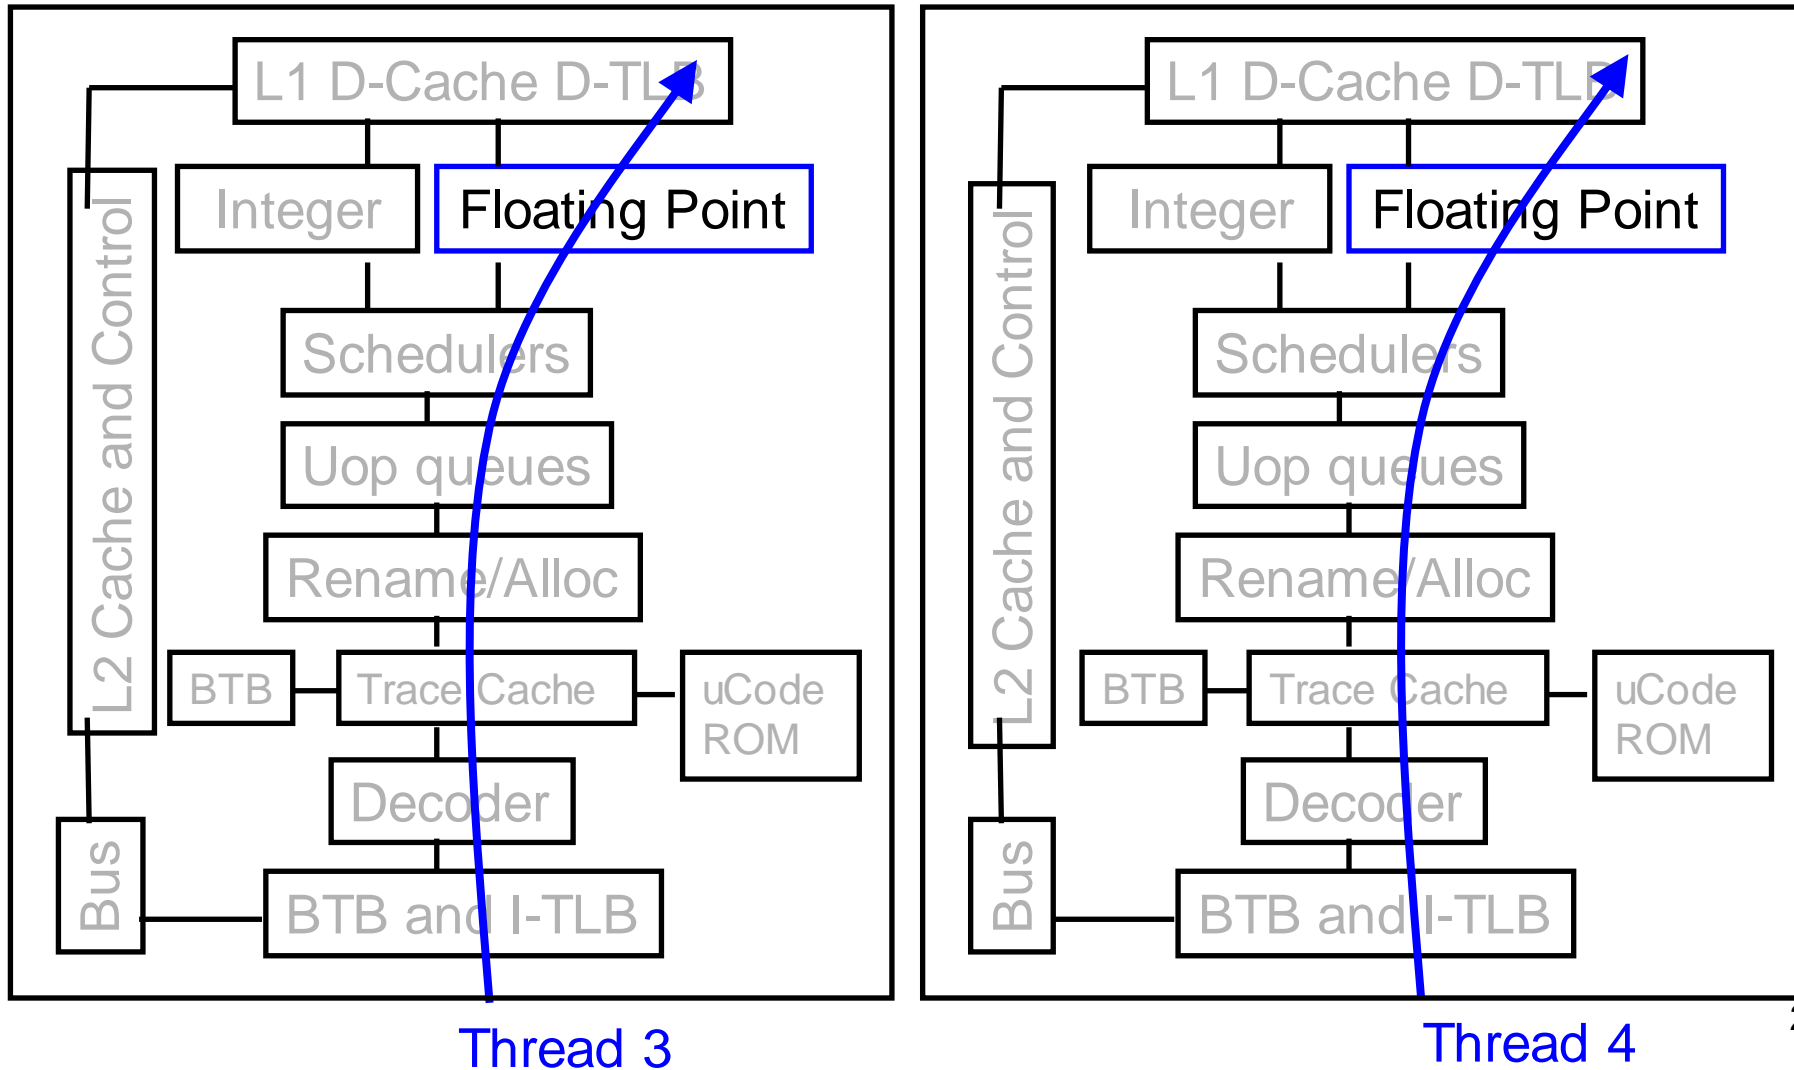
Multi-core

- Threads can run on separate cores



Multi-core

- Threads can run on separate cores

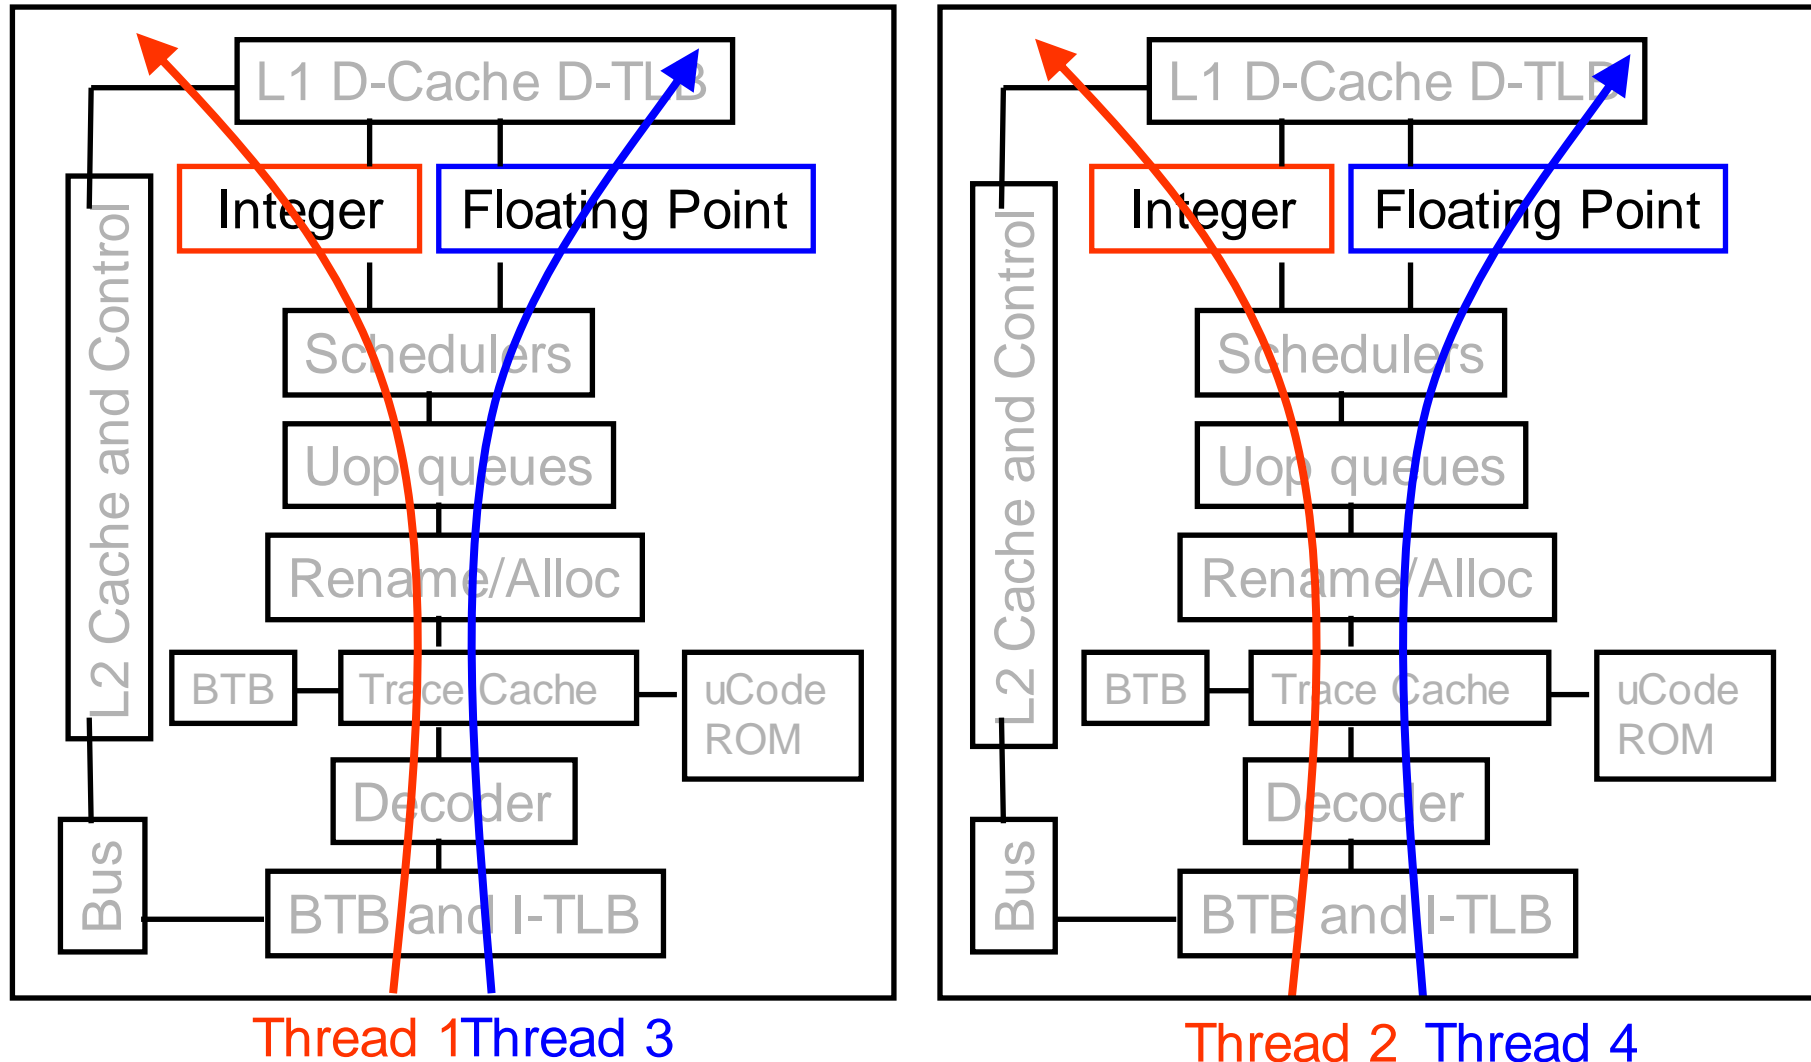


Combining Multi-core and SMT

- Cores can be SMT-enabled (or not)
- different combinations:
 - ⦿ single-core, non-SMT: standard uniprocessor
 - ⦿ single-core, w/ SMT
 - ⦿ multi-core, non-SMT
 - ⦿ multi-core, w/ SMT
- Number of SMT threads
 - ⦿ 2, 4, or sometimes 8 simultaneous threads
- Intel calls them “hyper-threads”

SMT Dual-core

- All four threads can run concurrently



Comparison: Multi-core vs. SMT


- Multi-core:
 - ⊙ Since there are several cores, each is smaller and not as powerful (but also easier to design and manufacture)
 - ⊙ However, great with thread-level parallelism
- SMT
 - ⊙ Can have one large and fast superscalar core
 - ⊙ Great performance on a single thread
 - ⊙ Mostly still only exploits instruction-level parallelism

Memory Hierarchy


- If simultaneous multithreading only:
 - ⊙ all caches shared
- Multi-core chips:
 - ⊙ L1 caches private
 - ⊙ L2 caches private in some architectures and shared in others
- Memory is always shared

Dual-core Intel Xeon Processors

- Each core
 - ⦿ hyper-threaded
- Private L1 caches
- Shared L2 caches




Designs w/ private L2 caches



both L1 and L2 are private

examples: amd opteron,
amd athlon, intel pentium d



a design with L3 caches

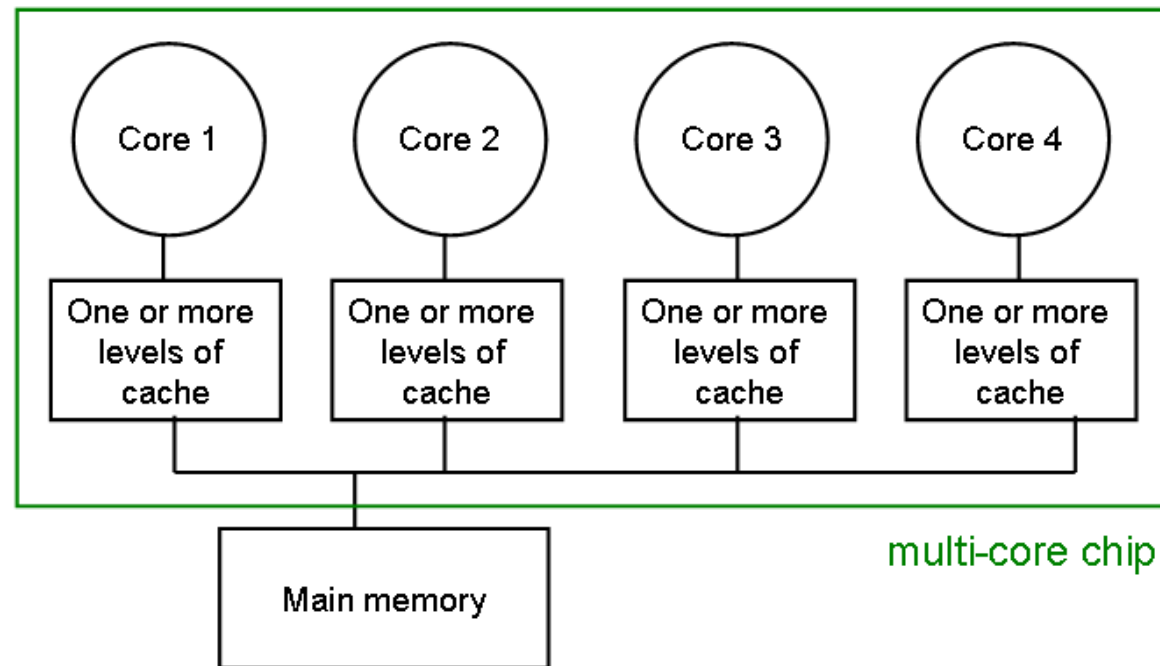
example: intel itanium 2

Private vs shared caches

- Advantages of private:
 - ⊙ They are closer to core, so faster access
 - ⊙ Reduces contention
- Advantages of shared:
 - ⊙ Threads on different cores can share the same cache data
 - ⊙ More cache space available if a single (or a few) high-performance thread runs on the system

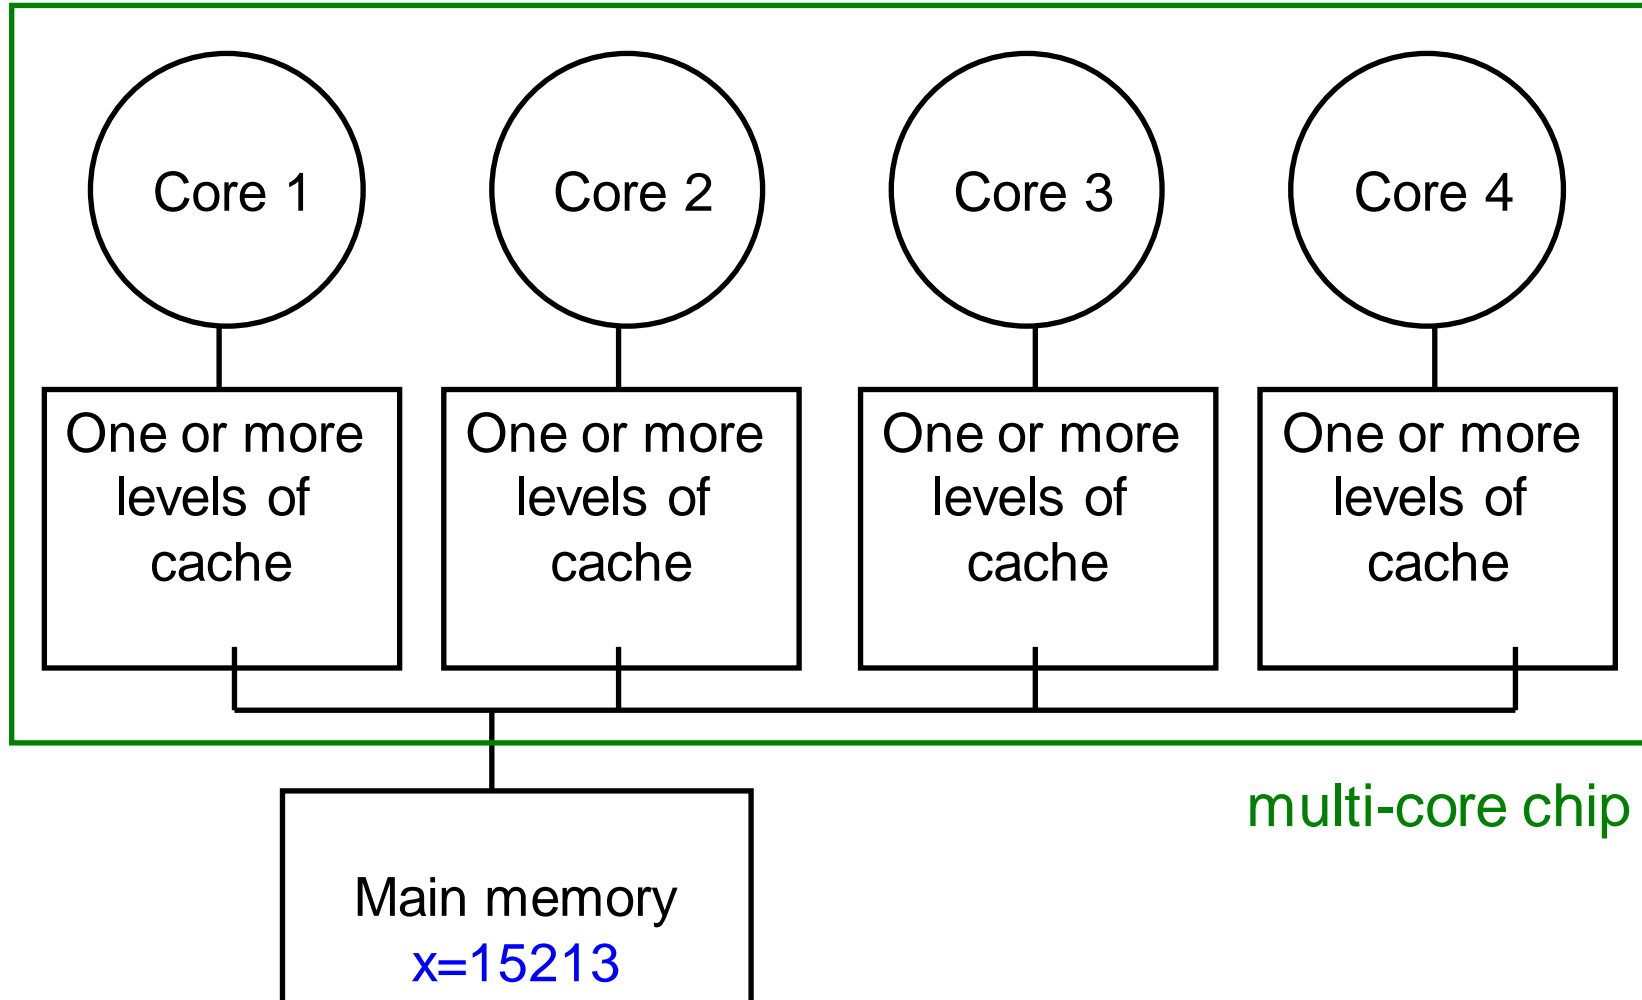
Cache Coherence Problem

- Since we have private caches, how to keep the data consistent across caches?
- Each core should perceive the memory as a monolithic array, shared by all the cores



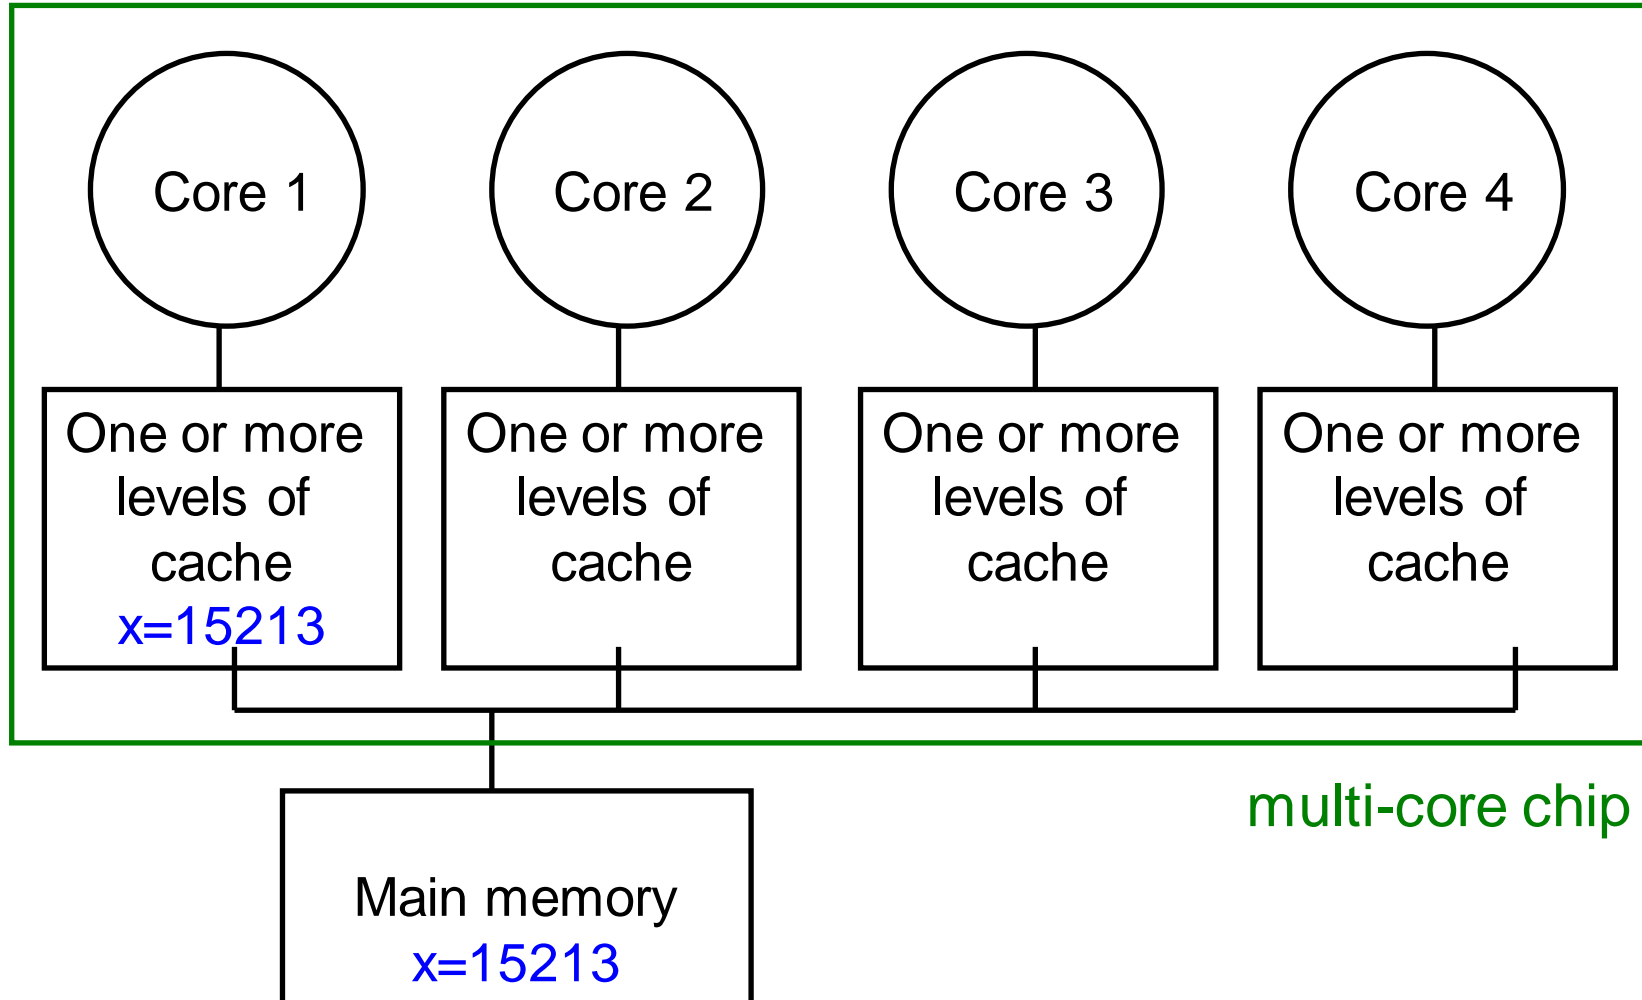
Cache Coherence Problem

- Suppose variable x initially contains 15213



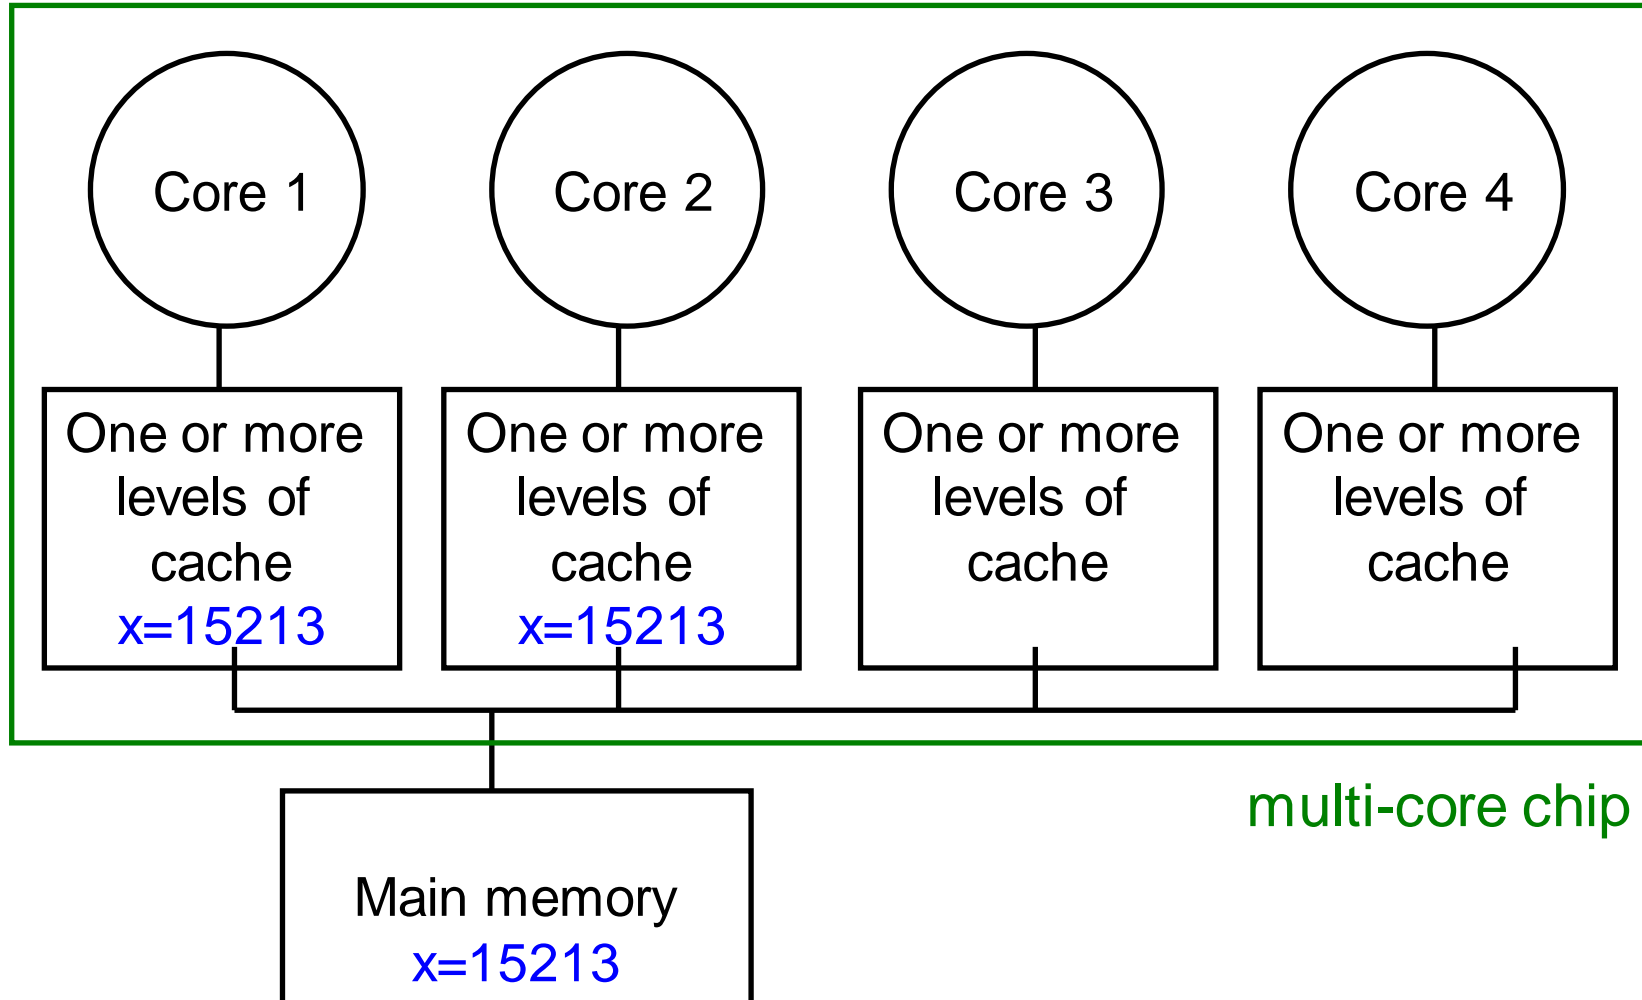
Cache Coherence Problem

- core 1 reads x




Cache Coherence Problem

- core 2 reads x




Cache Coherence Problem

- core 1 writes to x, setting it to 21660



Cache Coherence Problem

- core 2 attempts to read x ... gets a stale copy



Solutions for Cache Coherence Problem

- General problem w/ multiprocessors, not limited just to multi-core
- Many solution algorithms, coherence protocols, etc.
 - ⊙ A simple solution: invalidation-based protocol with snooping
 - ⊙ MSI, MESI (Modified, Exclusive, Shared, Invalid)

Announcement

- Next Lecture: cache coherence